

CSCI 544

Applied Natural Language Processing

Mohammad Rostami
USC Computer Science Department



Logistical Notes

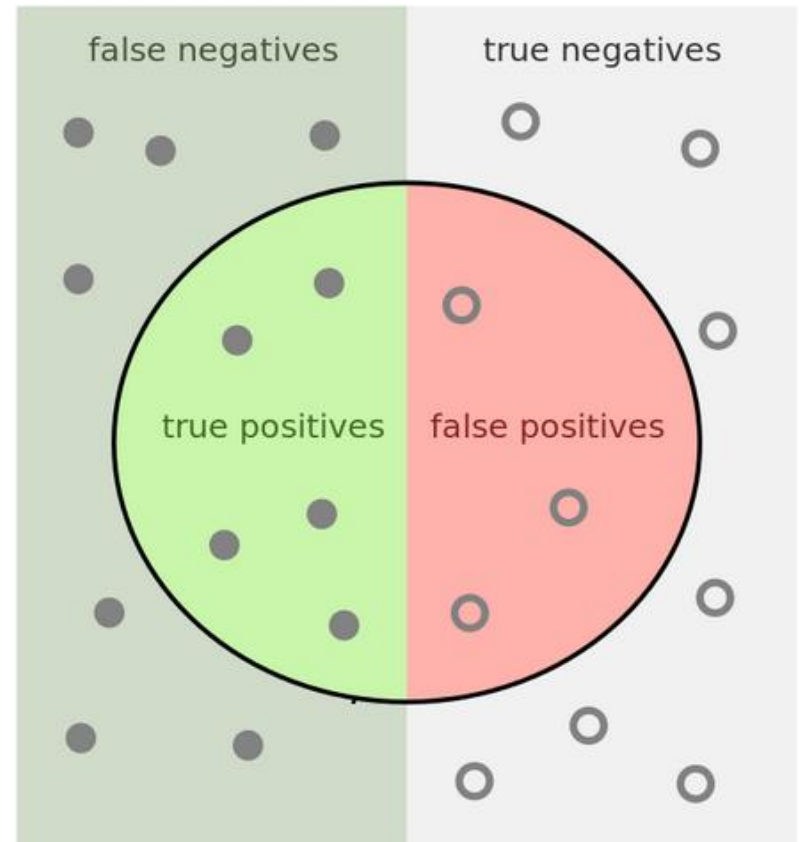
- Quizzes:
 - Quiz Period: 4:05-4:20 and still 10 minutes
 - Quiz 1 average: 84 (108 people scored 100)
 - We will have Quiz 2 next Thursday
- HW1:
 - Please check the homework description for edits done based on feedback from the Slack channel and prepare your report accordingly

Model Evaluation Process

- We use a training dataset for model selection
- A **good** parametric model along with a **suitable** training algorithm guarantees training a model that works well on the training data
- We need to validate that trained models **generalize** well on unseen data instances
- We need a second testing dataset which is fully independent of the training dataset
- We randomly split the annotated dataset into testing and training splits (sometimes, a validation set is generated as well)

Evaluation Metrics

- **Accuracy:** proportion of correctly classified items
 - Accuracy can be dominated by **true negatives** (items correctly classified as not in a class).
 - Sensitive with respect to imbalance
- Precision: $\frac{\text{True Positives}}{\text{True Positives} + \text{False Positive}}$
 - Also called positive predictive value
- Recall: $\frac{\text{True Positives}}{\text{True Positives} + \text{False negative}}$
 - Also called sensitivity
- Precision and recall are not useful metrics when used in isolation?
- We want our model to have good performance with respect to both metrics
- Implemented in sklearn



Evaluation Metrics

- Why having one measure is helpful?
- $F1 = \frac{2 \text{ Precision Recall}}{\text{Precision} + \text{Recall}}$
- F1 is biased towards the lower of precision and recall:
 - harmonic mean < geometric mean < arithmetic mean
 - F1=0 when Precision=0 or Recall=0
- Generalized F score:

$$F_{\beta} = (1 + \beta^2) \cdot \frac{\text{precision} \cdot \text{recall}}{(\beta^2 \cdot \text{precision}) + \text{recall}}$$

Natural Language Representation

- Language processing hierarchy levels:
 - **Documents**
 - Sentences
 - Phrase
 - Words
- Sparsity in the NLP training datasets: natural language has a very huge space:
 - Ex: Average Wikipedia page size is 580 words and English has ~1M words, yet the actual possibilities is far more.
- We need **interpretable** representations or **embeddings** to represent natural language data for model training
 - One-hot representation: too large (15M words) and meaningless
 - Hotel: [0,0,0,0,1,0,0,0,0,0,0,...,0,0,0]
 - Motel:[0,0,0,0,0,0,0,0,0,0,1,0,...,0,0,0]

Similarity of Vectors


- Euclidean distance, i.e., geometric closeness :

- Curse of dimensionality

- Dot product:

$$a \bullet b = ||a|| ||b|| \cos(\theta_{ab})$$

$$= a_1 b_1 + a_2 b_2 + \dots + a_n b_n$$



- Cosine similarity (scale invariant)

$$\cos \theta_{ab} = a \bullet b / ||a|| ||b|| \rightarrow 1 - \cos \theta_{ab} \text{ is a metric}$$

- Invariant with respect to the vector starting point

- **EX:** Hotel: $[0,0,0,0,1,0,0,0,0,0,0,\dots,0,0,0]$, Motel: $[0,0,0,0,0,0,0,0,0,0,1,0,\dots,0,0,0]$

$$\text{Hotel}' * \text{Motel} = 0$$

The Distributional Hypothesis

- Words that occur in the **same contexts** tend to have similar meanings (Zellig Harris, 1954)
 - Example: nice, good
- Word relatedness association (Budanitsky and Hirst, 2006): related words **co-occur** in different contexts
 - Example: cup, coffee
- If semantic similarity and association of words can be encoded into their representations, we may be able to address the challenge of sparsity
 - In the absence of a particular word during training, we can rely on its synonyms that exist in the training dataset: Motel vs Hotel
 - We can draw conclusions:
Lecturers teach in the university-> Professor ____ in the university.

Vector Embedding of Words



- Represent words using dense vectors:
 - Latent Semantic Analysis/Indexing (SC Deerwester et al, 1988)
 - Term weighting-based model
 - Consider occurrences of terms at **document level**.
 - Word2Vec (Mikolov et al, 2013)
 - Prediction-based model.
 - Consider occurrences of terms at **context level**.
 - GloVe (Pennington et al, 2014)
 - Count-based model.
 - Consider occurrences of terms at **context level**.

Word Embedding

- Each word is represented by a vector:
 - The same size is used for all words
 - Relatively low dimensional (~300)
 - Vectors for similar words are similar (measured in dot product)
 - Vector operations can be used for semantic and syntactic

deductions, e.g.,

Queen – Woman + Man = King



- The key idea is to derive the embeddings from the distributions of word context as they appear in a large corpus.


Matrix Factorization

- We can form a matrix of M using the idea of Bag of Words or TF-IDF: the word representations are highly sparse

	Words											Length of the review(in words)
	1 This	2 movie	3 is	4 very	5 scary	6 and	7 long	8 not	9 slow	10 spooky	11 good	
Review 1	1	1	1	1	1	1	1	0	0	0	0	7
Review 2	1	1	2	0	0	1	1	0	1	0	0	8
Review 3	1	1	1	0	0	0	1	0	0	1	1	6

- Singular value decomposition (U, V are orthonormal)


$$M_{m \times n} = U_{m \times n} \Sigma_{n \times n} V_{n \times n}^T$$



Matrix Factorization

- Many singular values are going to be zero or negligible

$$M_{m \times n} \approx U'_{m \times k} \Sigma'_{k \times k} V'^T_{k \times n}$$



$$m_l = u_l \Sigma' V'^T$$
$$u_l = m_l V' \Sigma'^{-1}$$


- We can use rows of U as word embeddings
 - An old idea for dimensionality reduction (it is possible to use other matrix factorization methods, e.g., non-negative matrix factorization)
 - Determining context is heuristic
 - Computationally expensive
 - Hyperparameters: contexts, cell values, optimization
 - computational expensive with $O(mn^2)$ cost for an $n \times m$ matrix
 - Hard to incorporate new words

Word2Vec

- Core idea: find embeddings using a prediction task involving **neighboring words** in a huge real-world corpus.
 - Input data can be sets of **successive word-patterns** from meaningful sentences in the corpus, e.g., “one of the most important”.
 - Try to build **synthetic** prediction tasks using these patterns, e.g., “(one of the __ important, most)”
 - Train a model to solve the prediction task
 - Embeddings are found as a **byproduct** of this process
- More specifically:
 - **We consider a window with the center word w_t and “context words” $w_{t'}$ with a window fixed size, e.g., $(t'=t-5, \dots, t-1, t+1, \dots, t+5)$.**
 - The model is assumed to be a two-layer neural network
 - We train the network to predict all $w_{t'}$ given w_t such that $p(w_{t'} | w_t)$ is maximized
 - We learn embeddings such that the prediction loss is minimized, i.e., if two words occur in close proximity, their representations become similar.

Skip-Gram

- Given a center word, we predict the context words:
 - Vocabulary size: V
 - Input layer: center word in 1-hot form.
 - The row k in $W_{V \times N}$ is the vector embedding of k -th center word.
 - The column k of $W'_{N \times V}$ is context vector of the k -th word.
 - At output layer y_{ij} , $i=1..M$ is computed:
 - We use the context word 1-hot vector to choose its column in $W'_{N \times V}$
 - dot product with h_i for the center word
 - compute the softmax
 - Match** the output one-hot vector
 - After optimization, we will have two vectors for each word. We can set the eventual embedding to be the average of these two vectors




$$h_{N \times 1} = W_{V \times N}^T x_{V \times 1} = v$$

$$\hat{z}_j = h_{N \times 1}^T W'[:, j]_{N \times 1} = v^T u$$

$$\hat{y} = \sigma(W_{N \times V}'^T h_{N \times 1})$$

Continuous Bag of Words

- Given a context word, we predict the center word:
 - Vocabulary size: V
 - Input layer: context words in 1-hot form.
 - The row k in $W_{V \times N}$ is the vector embedding of k -th context word.
 - The column k of $W'_{N \times V}$ is vector of the k -th center word.
 - The output column y_{ij} , $i=1..M$ is computed:
 - We use the center word 1-hot vector to choose its column in $W'_{N \times V}$
 - dot product with h_i for the context word
 - compute the softmax
 - We can set the eventual embedding to be the average of these two vectors
- Skip-gram incorporates non-frequent words better than CBOW



- Ex: Today is really __ day
- Skip-gram: delightful+context
and nice+context

Word2Vec Optimization Problem

- Consider an arbitrary order on vocabulary and let v_t be the word vector for center word t and u_t be the word vector for the context word t :

- Is the word vector unique?

Ex: [...like to **eat** lunch and...]

- We solve for the word vector by maximizing the following likelihood function

$$\begin{aligned} v_t, u_t &= \arg \max \log(\Pi_{t=1}^T P(c|w_t)) = \\ \arg \max \log(\Pi_{t=1}^T P(c|w_t)) &= \\ \arg \max \log(\Pi_{t=1}^T \prod_{\substack{j=-M \\ j \neq 0}}^M P(w_{t-j}|w_t)) &= \quad \leftarrow \text{Important Step} \\ \arg \max \frac{1}{T} \sum_{t=1}^T \sum_{\substack{j=-M \\ j \neq 0}}^M \log(P(w_{t-j}|w_t)) &= \end{aligned}$$

Word2Vec Optimization Problem

- Conditional Probability modeling

$$P(w_{t-j}|w_t) = \frac{e^{u_{t-j}^T v_t}}{\sum_{t=1}^T e^{u_{t-j}^T v_t}}$$

Is it an extension of the logistic function?

Computationally expensive!

- The log-likelihood optimization problem

$$u_o, v_c = \arg \min \frac{1}{V} \sum_{w=1}^V -u_o^T v_c + \log\left(\sum_{w'=1}^V e^{u_o^T v_{c'}}\right)$$

- Can be solved similar to logistic regression objective using numerical optimization techniques, e.g., gradient descent

Word2Vec Optimization Problem

- gradient descent step

$$u_o, v_c = \arg \min \frac{1}{V} \sum_{w=1}^V -u_o^T v_c + \log \left(\sum_{w'=1}^V e^{u_o^T v_c} \right)$$

$$v_c^{i+1} = v_c^i - \eta \nabla f(v_c^i)$$

$$\nabla f(v_c^i) = -u_o + \sum_{w=1}^V \frac{e^{u_o^T v_c}}{\sum_{w'=1}^V e^{u_o^T v_c}} u_o$$

$$\nabla f(v_c^i) = -u_o + \sum_{w=1}^V p(v_o | u_c) u_o = -u_o (1 - E(u_o))$$

- Tutorial: <https://rare-technologies.com/word2vec-tutorial/>

Negative Sampling

- Word2Vec optimization is a highly computationally intensive problem: the **shallow** network has a large number of weights and we will have billions of pairs
- Because we use one-hot vectors, each training pair (c,o) contributes minimally to updating the weights
- Negative sampling: for each positive pair, we randomly generate negative pairs, for which the network output should be 0.

$$u_o, v_c = \arg \min \sigma(u_o^T v_c) + \sum_{k=1}^K E_{v_{w_i} \sim P(w_c)} \log(-\sigma(v_{w_i}^T v_c))$$

- This is an instance of naïve data augmentation

Neural vs Traditional Embeddings

- Comparison is challenging (Levy and Goldberg, NeurIPS 2014):
 - **Hyperparameters**
 - Factorization algorithm
 - Amount of data
- A particular word embedding approach is unlikely to be state-of-the-art for all applications
- Hyperparameters appear to have the largest impact in performance.
- Neural models are less sensitive with respect to hyperparameters, and training data preparation is more straightforward and systematic