


CSCI 544: Applied Natural Language Processing

Sequence Labeling-I

Xuezhe Ma (Max)



USC University of
Southern California

Logistical Points

- **For Team Projects**
 - If you have not joined a team, please write your name on the spreadsheet below all teams.
 - We will assign you randomly

Overview

- **The Sequence Labeling Problem**
 - General Structured Prediction Tasks
 - Part-of-speech Tagging: A case study
 - Generative Models vs. Discriminative Models
 - Maximum Likelihood Estimation (MLE)
- **Hidden Markov Model (HMM)**
 - Basic definitions
 - Parameter estimation
 - The Viterbi algorithm
- **Log-Linear Models**
 - Maximum Entropy Markov Models (MEMMs)
 - Conditional Random Fields (CRFs)

Overview

- **The Sequence Labeling Problem**
 - General Structured Prediction Tasks
 - Part-of-speech Tagging: A case study
 - Generative Models vs. Discriminative Models
 - Maximum Likelihood Estimation (MLE)
- **Hidden Markov Model (HMM)**
 - Basic definitions
 - Parameter estimation
 - The Viterbi algorithm
- **Log-Linear Models**
 - Maximum Entropy Markov Models (MEMMs)
 - Conditional Random Fields (CRFs)

What is Structured Prediction

- Unstructured Prediction
 - Output Y consists of a **single component**



X



{**Cat**, Dog, Finch, Owl}


Y

Structured Prediction

- Y consists of **multiple components** $Y = \{y_1, y_2, \dots, y_n\}$

Max is teaching NLP

X




Structured Prediction

- Y consists of **multiple components** $Y = \{y_1, y_2, \dots, y_n\}$

Max is teaching NLP

X



$Y_1 = < \mathbf{teaching} \rightarrow \mathbf{Max} >$

Structured Prediction

- Y consists of **multiple components** $Y = \{y_1, y_2, \dots, y_n\}$

麦克斯 在 南加大 工作

X



Max is working at USC

↑
 Y_1

↑
 Y_2

↑
 Y_3

↑
 Y_4

↑
 Y_5

Structured Prediction

- Y consists of **multiple components** $Y = \{y_1, y_2, \dots, y_n\}$
- **(Strong) correlations** between output components

麦克斯 在 南加大 工作

X



Max is working at USC

↑
 Y_1

↑
 Y_2

↑
 Y_3

↑
 Y_4

↑
 Y_5

Structured Prediction

- Y consists of **multiple components** $Y = \{y_1, y_2, \dots, y_n\}$
- **(Strong) correlations** between output components
- **Exponential** output space
 - Decoding: $y^* = \operatorname{argmax}_{y \in \mathcal{Y}} p(y | x)$

麦克斯 在 南加大 工作

X



Max is working at USC

↑
 Y_1

↑
 Y_2


↑
 Y_3

↑
 Y_4

↑
 Y_5

What is Sequence Labeling?

A type of structured prediction tasks



Assigning each token of X , e.g. x_i a corresponding label y_i

Why Sequence Labeling?

Part-of-Speech Tagging

PRP VBD DT NN IN DT NN
I saw a girl with a telescope

Named Entity Recognition

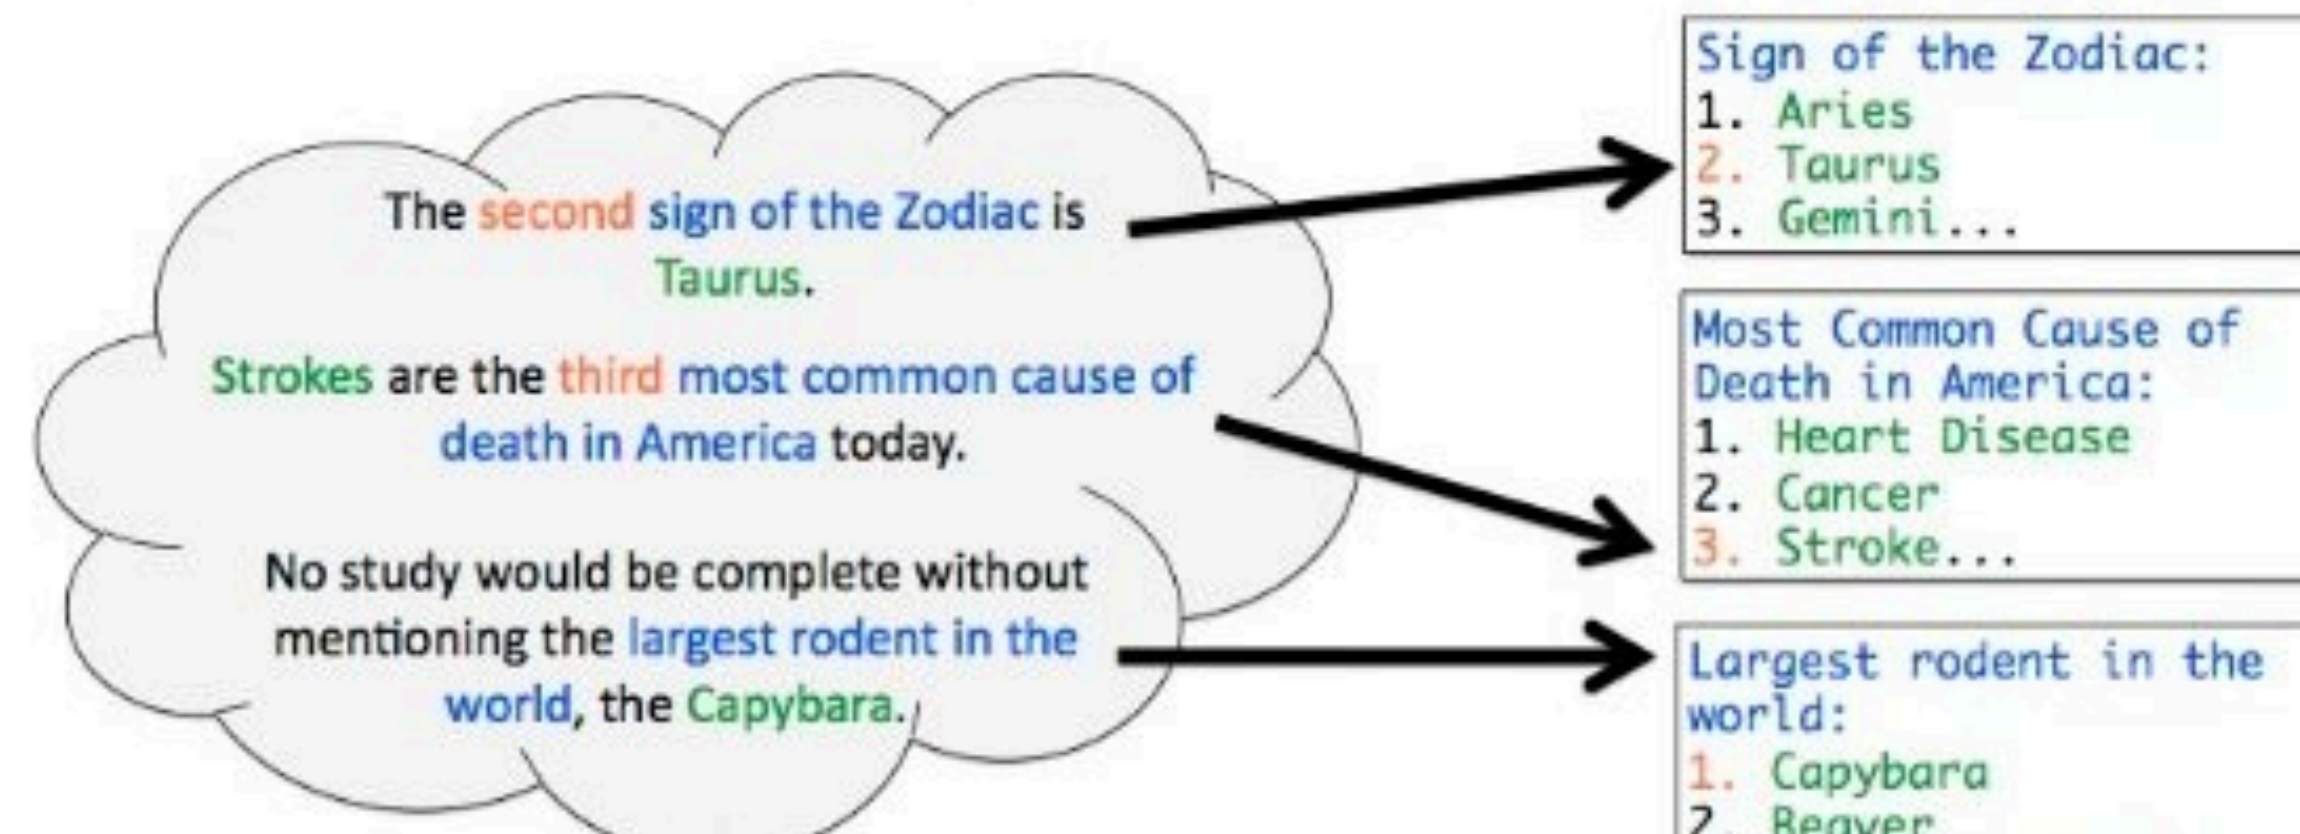
PERSON TITLE ORGANIZATION DATE
Arsène Wenger was named manager of Arsenal in 1996 .

Information Extraction

Unstructured Web Text

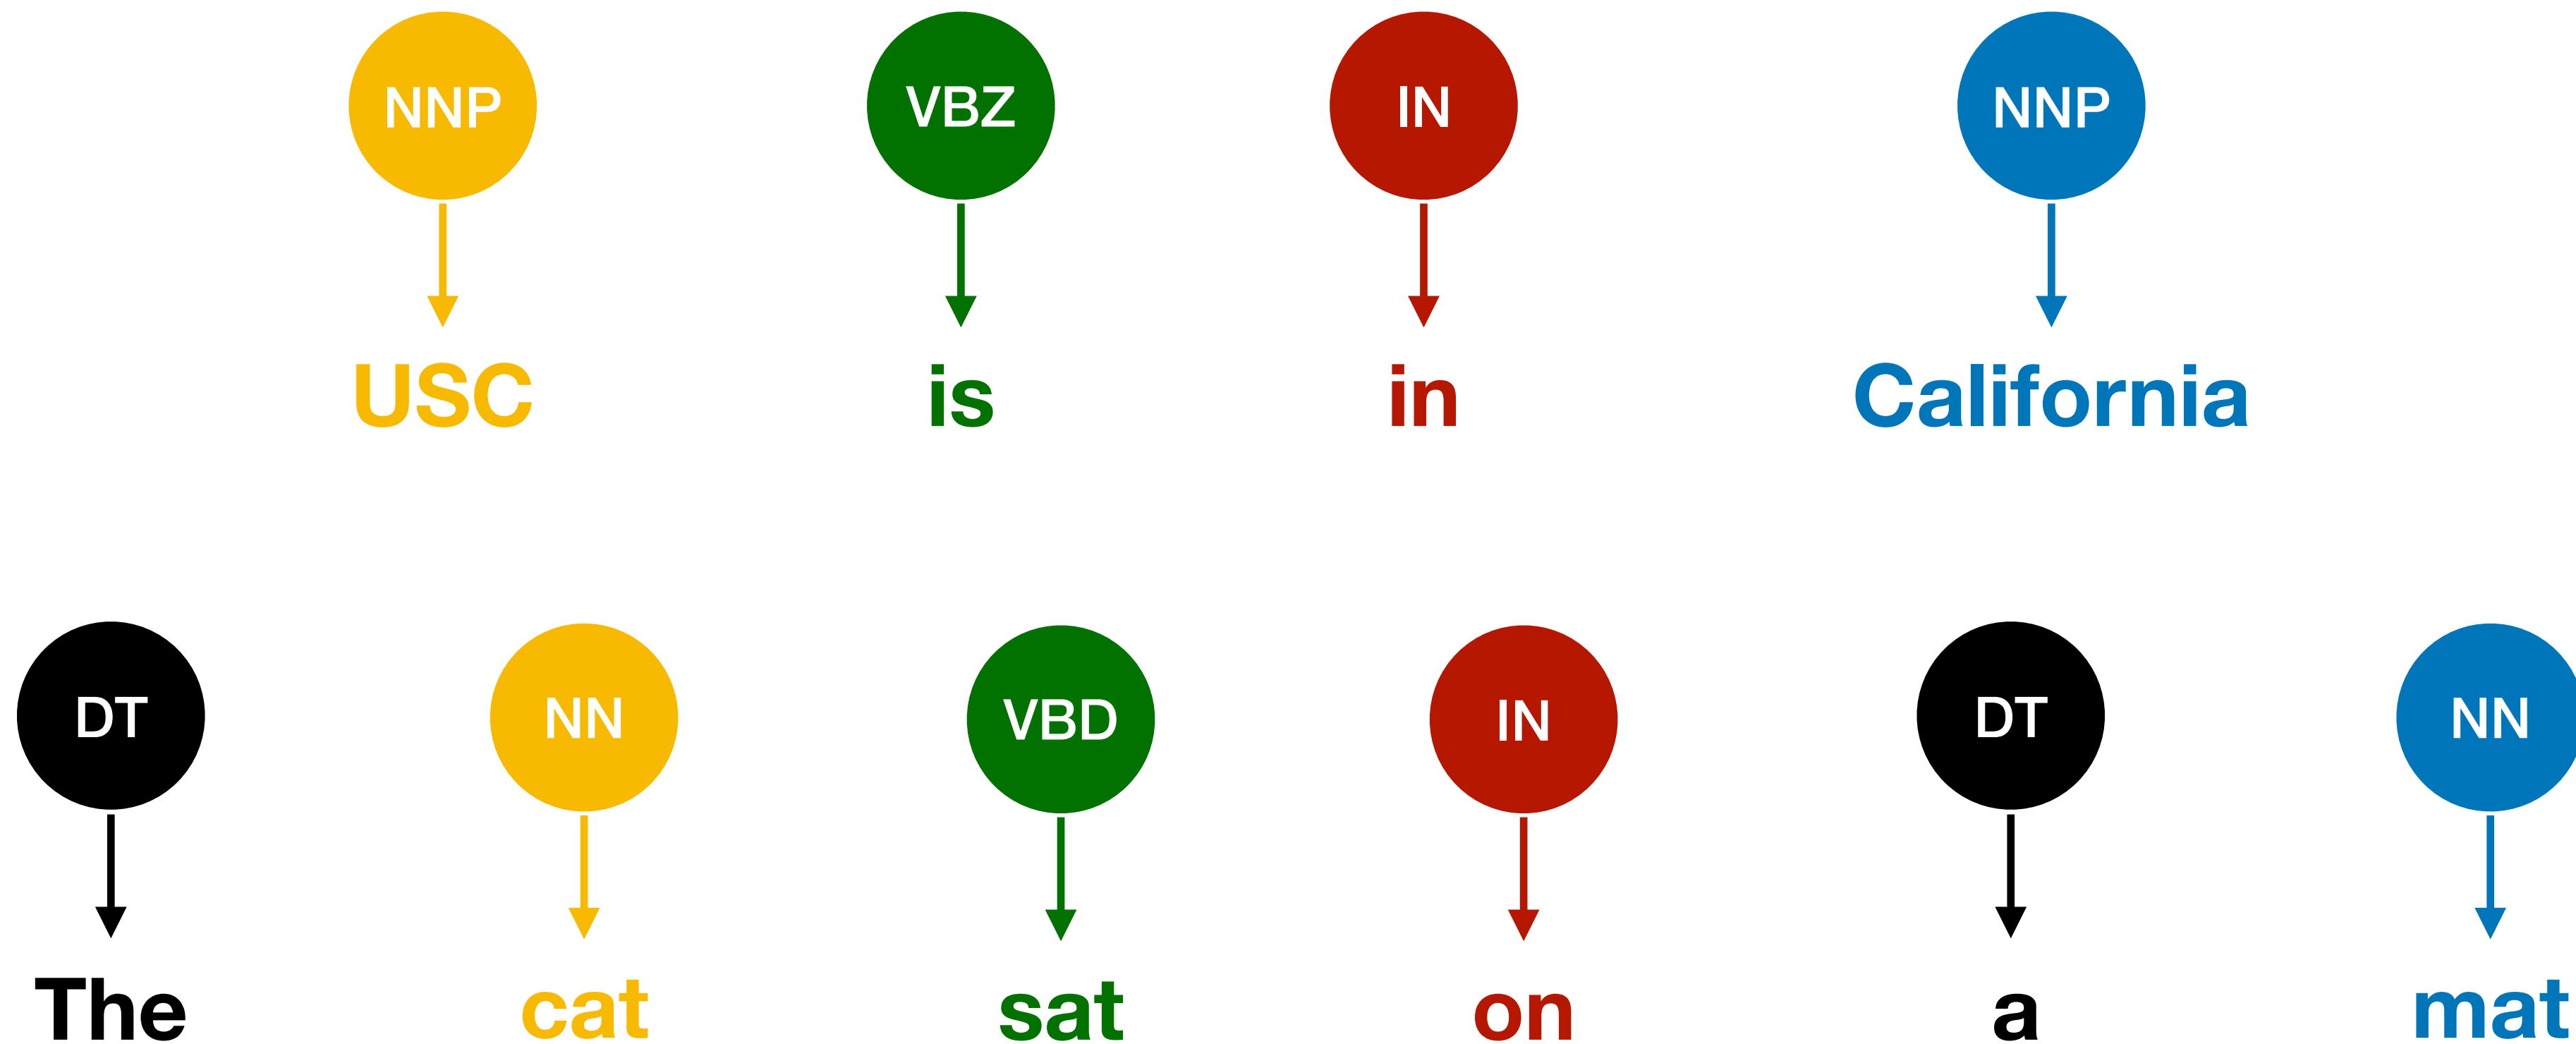


Structured Sequences




What are Part-of-Speech (POS) Tags

- Word classes or syntactic categories
- Reveal useful information about the syntactic role of a word (and its neighbors!)




Part of Speech

- Different words have different syntactic functions
- Can be roughly divided into two classes
 - **Closed class**: fixed membership, function words
 - e.g. prepositions (in, on, of), determiners (a, the)
 - **Open class**: New words get added frequently
 - e.g. nouns (Twitter, Facebook), verbs (google), adjectives and adverbs.



Part of Speech

- How many part of speech tags do you think English has?
 - A. < 10
 - B. 10 - 30
 - C. 30 - 50
 - D. > 50



Penn Tree Bank Tagset

Tag	Description	Example	Tag	Description	Example	Tag	Description	Example
CC	coordinating conjunction	<i>and, but, or</i>	PDT	predeterminer	<i>all, both</i>	VBP	verb non-3sg present	<i>eat</i>
CD	cardinal number	<i>one, two</i>	POS	possessive ending	<i>'s</i>	VBZ	verb 3sg pres	<i>eats</i>
DT	determiner	<i>a, the</i>	PRP	personal pronoun	<i>I, you, he</i>	WDT	wh-determ.	<i>which, that</i>
EX	existential 'there'	<i>there</i>	PRP\$	possess. pronoun	<i>your, one's</i>	WP	wh-pronoun	<i>what, who</i>
FW	foreign word	<i>mea culpa</i>	RB	adverb	<i>quickly</i>	WP\$	wh-possess.	<i>whose</i>
IN	preposition/ subordin-conj	<i>of, in, by</i>	RBR	comparative adverb	<i>faster</i>	WRB	wh-adverb	<i>how, where</i>
JJ	adjective	<i>yellow</i>	RBS	superlatv. adverb	<i>fastest</i>	\$	dollar sign	<i>\$</i>
JJR	comparative adj	<i>bigger</i>	RP	particle	<i>up, off</i>	#	pound sign	<i>#</i>
JJS	superlative adj	<i>wildest</i>	SYM	symbol	<i>+, %, &</i>	“	left quote	<i>' or “</i>
LS	list item marker	<i>1, 2, One</i>	TO	“to”	<i>to</i>	”	right quote	<i>' or ”</i>
MD	modal	<i>can, should</i>	UH	interjection	<i>ah, oops</i>	(left paren	<i>[, (, {, <</i>
NN	sing or mass noun	<i>llama</i>	VB	verb base form	<i>eat</i>)	right paren	<i>],), }, ></i>
NNS	noun, plural	<i>llamas</i>	VBD	verb past tense	<i>ate</i>	,	comma	<i>,</i>
NNP	proper noun, sing.	<i>IBM</i>	VBG	verb gerund	<i>eating</i>	.	sent-end punc	<i>. ! ?</i>
NNPS	proper noun, plu.	<i>Carolinas</i>	VBN	verb past part.	<i>eaten</i>	:	sent-mid punc	<i>: ; ... --</i>

45 tags!

(Marcus et al., 1993)


The Task of Part of Speech Tagging

- Tag each word with its part of speech
- Disambiguation task: each word might have different senses/functions

- The/DT **back/ADJ** door/NN

- On/IN my/PRP\$ **back/NN**

- Win/VB the/DT voters/NNS **back/RP**



Types:		WSJ	Brown
Unambiguous	(1 tag)	44,432 (86%)	45,799 (85%)
Ambiguous	(2+ tags)	7,025 (14%)	8,050 (15%)
Tokens:			
Unambiguous	(1 tag)	577,421 (45%)	384,349 (33%)
Ambiguous	(2+ tags)	711,780 (55%)	786,646 (67%)

Figure 8.2 Tag ambiguity for word types in Brown and WSJ, using Treebank-3 (45-tag) tagging. Punctuation were treated as words, and words were kept in their original case.

A Simple Baseline

- Many words might be easy to disambiguate
- **Most Frequent Class:** Assign each token (word) to the class it occurred most in the training data. (e.g. student/NN)
 - Entirely discarding contextual information
- How accurate do you think this baseline would be at tagging words?
 - A. < 50%
 - B. 50% - 75%
 - C. 75% - 90%
 - D. > 90%

Accurately tags **92.34%** of word tokens on Wall Street Journal (WSJ)

POS Tagging **Not** Solved!

- State of the art: $\sim 97\%$
- Sentence level accuracies
 - Average length of English sentence ~ 14 words
 - $0.92^{14} = 31\%$ vs. $0.97^{14} = 65\%$
- Highly relying on domain information
 - Training data and testing data must be from the same domain
 - $< 70\%$ on data from social media

Some Observations

- The function (or POS) of a word depends on its context
 - The/DT **back/ADJ** door/NN
 - On/IN my/PRP\$ **back/NN**
 - Win/VB the/DT voters/NNS **back/RP**
- Certain POS combinations are extremely unlikely
 - **<JJ, DT>** (“good the”) or **<DT, IN>** (“the in”)
- Better to make predictions on entire sentences instead of individual words

Sequence Labeling Models!

Overview

- The Sequence Labeling Problem
 - General Structured Prediction Tasks
 - Part-of-speech Tagging: A case study
 - Generative Models vs. Discriminative Models
 - Maximum Likelihood Estimation (MLE)
- **Hidden Markov Model (HMM)**
 - Basic definitions
 - Parameter estimation
 - The Viterbi algorithm
- Log-Linear Models
 - Maximum Entropy Markov Models (MEMMs)
 - Conditional Random Fields (CRFs)

Hidden Markov Models

Markov Sequences

- ▶ Consider a sequence of random variables X_1, X_2, \dots, X_m where m is the length of the sequence
- ▶ Each variable X_i can take any value in $\{1, 2, \dots, k\}$
- ▶ How do we model the joint distribution

$$P(X_1 = x_1, X_2 = x_2, \dots, X_m = x_m)$$

?

The Markov Assumption

$$\begin{aligned} & P(X_1 = x_1, X_2 = x_2, \dots, X_m = x_m) \\ &= P(X_1 = x_1) \prod_{j=2}^m P(X_j = x_j | X_1 = x_1, \dots, X_{j-1} = x_{j-1}) \\ &= P(X_1 = x_1) \prod_{j=2}^m P(X_j = x_j | X_{j-1} = x_{j-1}) \end{aligned}$$


Markov assumption

- ▶ The first equality is exact (by the chain rule).
- ▶ The second equality follows from *the Markov assumption*: for all $j = 2 \dots m$,

$$P(X_j = x_j | X_1 = x_1, \dots, X_{j-1} = x_{j-1}) = P(X_j = x_j | X_{j-1} = x_{j-1})$$

Markov Sequences

- A Generative Model for Sequences



Pick x_1 at random from the distribution $P(X_1)$

Pick x_2 at random from the distribution $P(X_2 | X_1 = x_1)$

Pick x_t at random from the distribution $P(X_t | X_{t-1} = x_{t-1})$

Modeling Pairs of Sequences

- In Sequence Labeling, we need to model pairs of sequences

$S = S_1, S_2, \dots, S_n$



$X = X_1, X_2, \dots, X_n$

USC

is

in

California

Hidden Markov Models (HMMs) allow us to *jointly* reason over X and S


Hidden Markov Models

- ▶ We have two sequences of random variables:
 X_1, X_2, \dots, X_m and S_1, S_2, \dots, S_m
- ▶ Intuitively, each X_i corresponds to an “observation” and each S_i corresponds to an underlying “state” that generated the observation. Assume that each S_i is in $\{1, 2, \dots, k\}$, and each X_i is in $\{1, 2, \dots, o\}$
- ▶ How do we model the joint distribution

$$P(X_1 = x_1, \dots, X_m = x_m, S_1 = s_1, \dots, S_m = s_m)$$

?

The HMM Assumptions



1. Markov Assumption on S

$$P(S_j = s_j | S_{j-1} = s_{j-1}, \dots, S_1 = s_1) = P(S_j = s_j | S_{j-1} = s_{j-1})$$


Transition Probabilities

2. Conditional Independence on X given S

$$P(X_1 = x_1, \dots, X_m = x_m | S_1 = s_1, \dots, S_m = s_m) = \prod_{j=1}^m P(X_j = x_j | S_j = s_j)$$

Emission Probabilities

The HMM Assumptions



1. Markov Assumption on S

$$P(S_3 = \text{IN} \mid S_2 = \text{VBZ}, S_1 = \text{NNP}) = P(S_3 = \text{IN} \mid S_2 = \text{VBZ})$$

2. Conditional Independence on X given S

$$P(\text{USC is in CA} \mid \text{NNP VBZ IN NNP}) = P(\text{USC} \mid \text{NNP})P(\text{is} \mid \text{VBZ})P(\text{in} \mid \text{IN})P(\text{CA} \mid \text{NNP})$$

Which assumption do you think is stronger?

Second assumption is stronger..why? Becoz its more powerful than the markov assumption

Joint Distribution of Sequence Pairs in HMMs

$$P(X_1 = x_1, \dots, X_m = x_m, S_1 = s_1, \dots, S_m = s_m)$$

$$= P(X_1 = x_1, \dots, X_m = x_m \mid S_1 = s_1, \dots, S_m = s_m)$$

Output Independence

$$\times P(S_1 = s_1, \dots, S_m = s_m)$$

Markov Assumption

$$= \prod_{j=1}^m P(X_j = x_j \mid S_j = s_j)$$

How to model $P(X_j = x_j \mid S_j = s_j)$

and $P(S_j = s_j \mid S_{j-1} = s_{j-1})$?

$$\times P(S_1 = s_1) \prod_{j=1}^m P(S_j = s_j \mid S_{j-1} = s_{j-1})$$

Homogeneous HMMs

- In a *homogeneous* HMM, we make an additional assumption:

$$P(S_j = s_j | S_{j-1} = s_{j-1}) = t(s_j | s_{j-1})$$

$$P(X_j = x_j | S_j = s_j) = e(x_j | s_j)$$

- Idea behind this assumption: the transition and emission probabilities do not depend on the position in the Markov chain (do not depend on the index j)

The Model Form for Homogeneous HMMs


- ▶ The model takes the following form:

$$p(x_1 \dots x_m, s_1 \dots s_m; \underline{\theta}) = t(s_1) \prod_{j=2}^m t(s_j | s_{j-1}) \prod_{j=1}^m e(x_j | s_j)$$

- ▶ Parameters in the model:

1. Initial state parameters $t(s)$ for $s \in \{1, 2, \dots, k\}$
2. Transition parameters $t(s' | s)$ for $s, s' \in \{1, 2, \dots, k\}$
3. Emission parameters $e(x | s)$ for $s \in \{1, 2, \dots, k\}$ and $x \in \{1, 2, \dots, o\}$

Example: Sequence Probability



What is the joint probability $P(\text{the cat, DT NN})$?


- A) $(0.8 * 0.8) * (0.9 * 0.5)$
- B) $(0.2 * 0.8) * (0.9 * 0.5)$
- C) $(0.3 * 0.7) * (0.5 * 0.5)$

Dummy start state

		s_{t+1}	
		DT	NN
s_t	\emptyset	0.8	0.2
	DT	0.2	0.8
	NN	0.3	0.7


		o_t	
		the	cat
s_t	DT	0.9	0.1
	NN	0.5	0.5

HMMs are Generative Models



1. Pick s_1 at random from the distribution $t(s)$. Pick x_1 from the distribution $e(x|s_1)$
2. For $j = 2 \dots m$:
 - ▶ Choose s_j at random from the distribution $t(s|s_{j-1})$
 - ▶ Choose x_j at random from the distribution $e(x|s_j)$

HMMs are Generative Models



1. Pick s_1 at random from the distribution $t(s)$. Pick x_1 from the distribution $e(x|s_1)$
2. For $j = 2 \dots m$:
 - ▶ Choose s_j at random from the distribution $t(s|s_{j-1})$
 - ▶ Choose x_j at random from the distribution $e(x|s_j)$

Learning a Hidden Markov Model

Parameter Estimation

- Assuming we have fully observed data $\{X_i, S_i\}_{i=1}^N$, e.g. WSJ

Training set:

1 Pierre/**NNP** Vinken/**NNP** ,/, 61/**CD** years/**NNS** old/**JJ** ,/
join/**VB** the/**DT** board/**NN** as/**IN** a/**DT** nonexecutive/**JJ** di
Nov./**NNP** 29/**CD** ./.

2 Mr./**NNP** Vinken/**NNP** is/**VBZ** chairman/**NN** of/**IN** Elsev
N.V./**NNP** ,/, the/**DT** Dutch/**NNP** publishing/**VBG** group/

3 Rudolph/**NNP** Agnew/**NNP** ,/, 55/**CD** years/**NNS** old/**JJ**
chairman/**NN** of/**IN** Consolidated/**NNP** Gold/**NNP** Fields/**N**
./, was/**VBD** named/**VBN** a/**DT** nonexecutive/**JJ** director/
this/**DT** British/**JJ** industrial/**JJ** conglomerate/**NN** ./.

...

38,219 It/**PRP** is/**VBZ** also/**RB** pulling/**VBG** 20/**CD** peopl
of/**IN** Puerto/**NNP** Rico/**NNP** ,/, who/**WP** were/**VBD** help
Hurricane/**NNP** Hugo/**NNP** victims/**NNS** ,/, and/**CC** sendin
them/**PRP** to/**TO** San/**NNP** Francisco/**NNP** instead/**RB** ./

Maximum Likelihood Estimate:

$$\max_{t(\cdot|\cdot), e(\cdot|\cdot)} \prod_{i=1}^N P(X_i, S_i)$$

$$t(s' | s) = \frac{\text{count}(s \rightarrow s')}{\text{count}(s)}$$

$$e(x | s) = \frac{\text{count}(s \rightarrow x)}{\text{count}(s)}$$

Learning Example

1. the/**DT** cat/**NN** sat/**VBD** on/**IN** the/**DT** mat/**NN**
2. Princeton/**NNP** is/**VBZ** in/**IN** New/**NNP** Jersey/**NNP**
3. the/**DT** old/**NN** man/**VB** the/**DT** boats/**NNS**

$$t(\mathbf{NN} | \mathbf{DT}) = \frac{3}{4}$$
$$e(\mathbf{cat} | \mathbf{NN}) = \frac{1}{3}$$

Maximum Likelihood Estimate:

$$\max_{t(\cdot|\cdot), e(\cdot|\cdot)} \prod_{i=1}^N P(X_i, S_i)$$

$$t(s' | s) = \frac{\text{count}(s \rightarrow s')}{\text{count}(s)}$$

$$e(x | s) = \frac{\text{count}(s \rightarrow x)}{\text{count}(s)}$$


Decoding with HMMs

Decoding with HMMs


- Goal: for a given input sequence x_1, \dots, x_m , find

$$\arg \max_{s_1, \dots, s_m} p(x_1 \dots x_m, s_1 \dots s_m; \underline{\theta})$$

- This is the most likely state sequence $s_1 \dots s_m$ for the given input sequence $x_1 \dots x_m$




Decoding with HMMs



$$S^* = \arg \max_{s_1, \dots, s_m} p(x_1, \dots, x_m, s_1, \dots, s_m) = t(s_1) \prod_{j=2}^m t(s_j | s_{j-1}) \prod_{j=1}^m e(x_j | s_j)$$

How can we maximize this over all state sequences?

Greedy Decoding




Decode/reveal one state at a time

$$s_1^* = \arg \max_{s_1} t(s_1) e(x_1 | s_1)$$

$$S^* = \arg \max_{s_1, \dots, s_m} p(x_1, \dots, x_m, s_1, \dots, s_m) = t(s_1) \prod_{j=2}^m t(s_j | s_{j-1}) \prod_{j=1}^m e(x_j | s_j)$$

Greedy Decoding

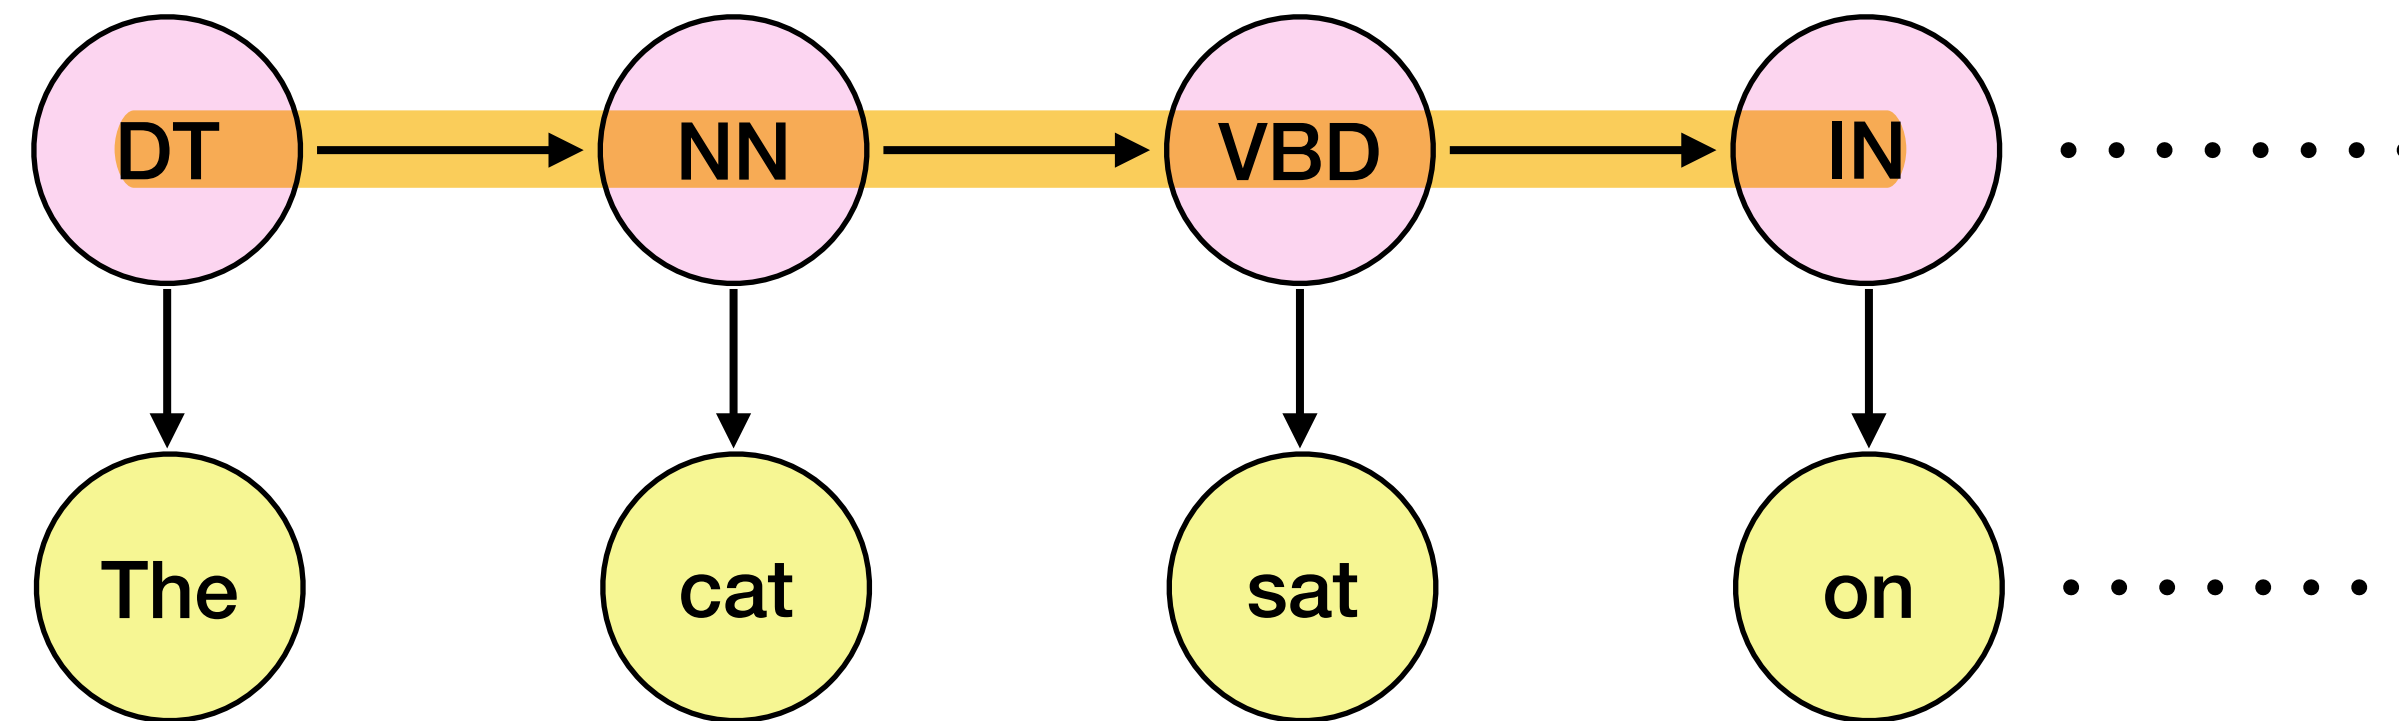


Decode/reveal one state at a time

$$s_2^* = \arg \max_{s_2} t(s_2 | s_1^*) e(x_2 | s_2)$$

$$S^* = \arg \max_{s_1, \dots, s_m} p(x_1, \dots, x_m, s_1, \dots, s_m) = t(s_1) \prod_{j=2}^m t(s_j | s_{j-1}) \prod_{j=1}^m e(x_j | s_j)$$

Greedy Decoding



$$s_j^* = \arg \max_{s_j} t(s_j | s_{j-1}^*) e(x_j | s_j), \quad \forall j$$

TIME COMPLEXITY OF THE ALGO: $O(KN)$ where k is the number of the number of the tags

- Local decisions
- Not guaranteed to produce the overall optimal sequence

Viterbi Decoding

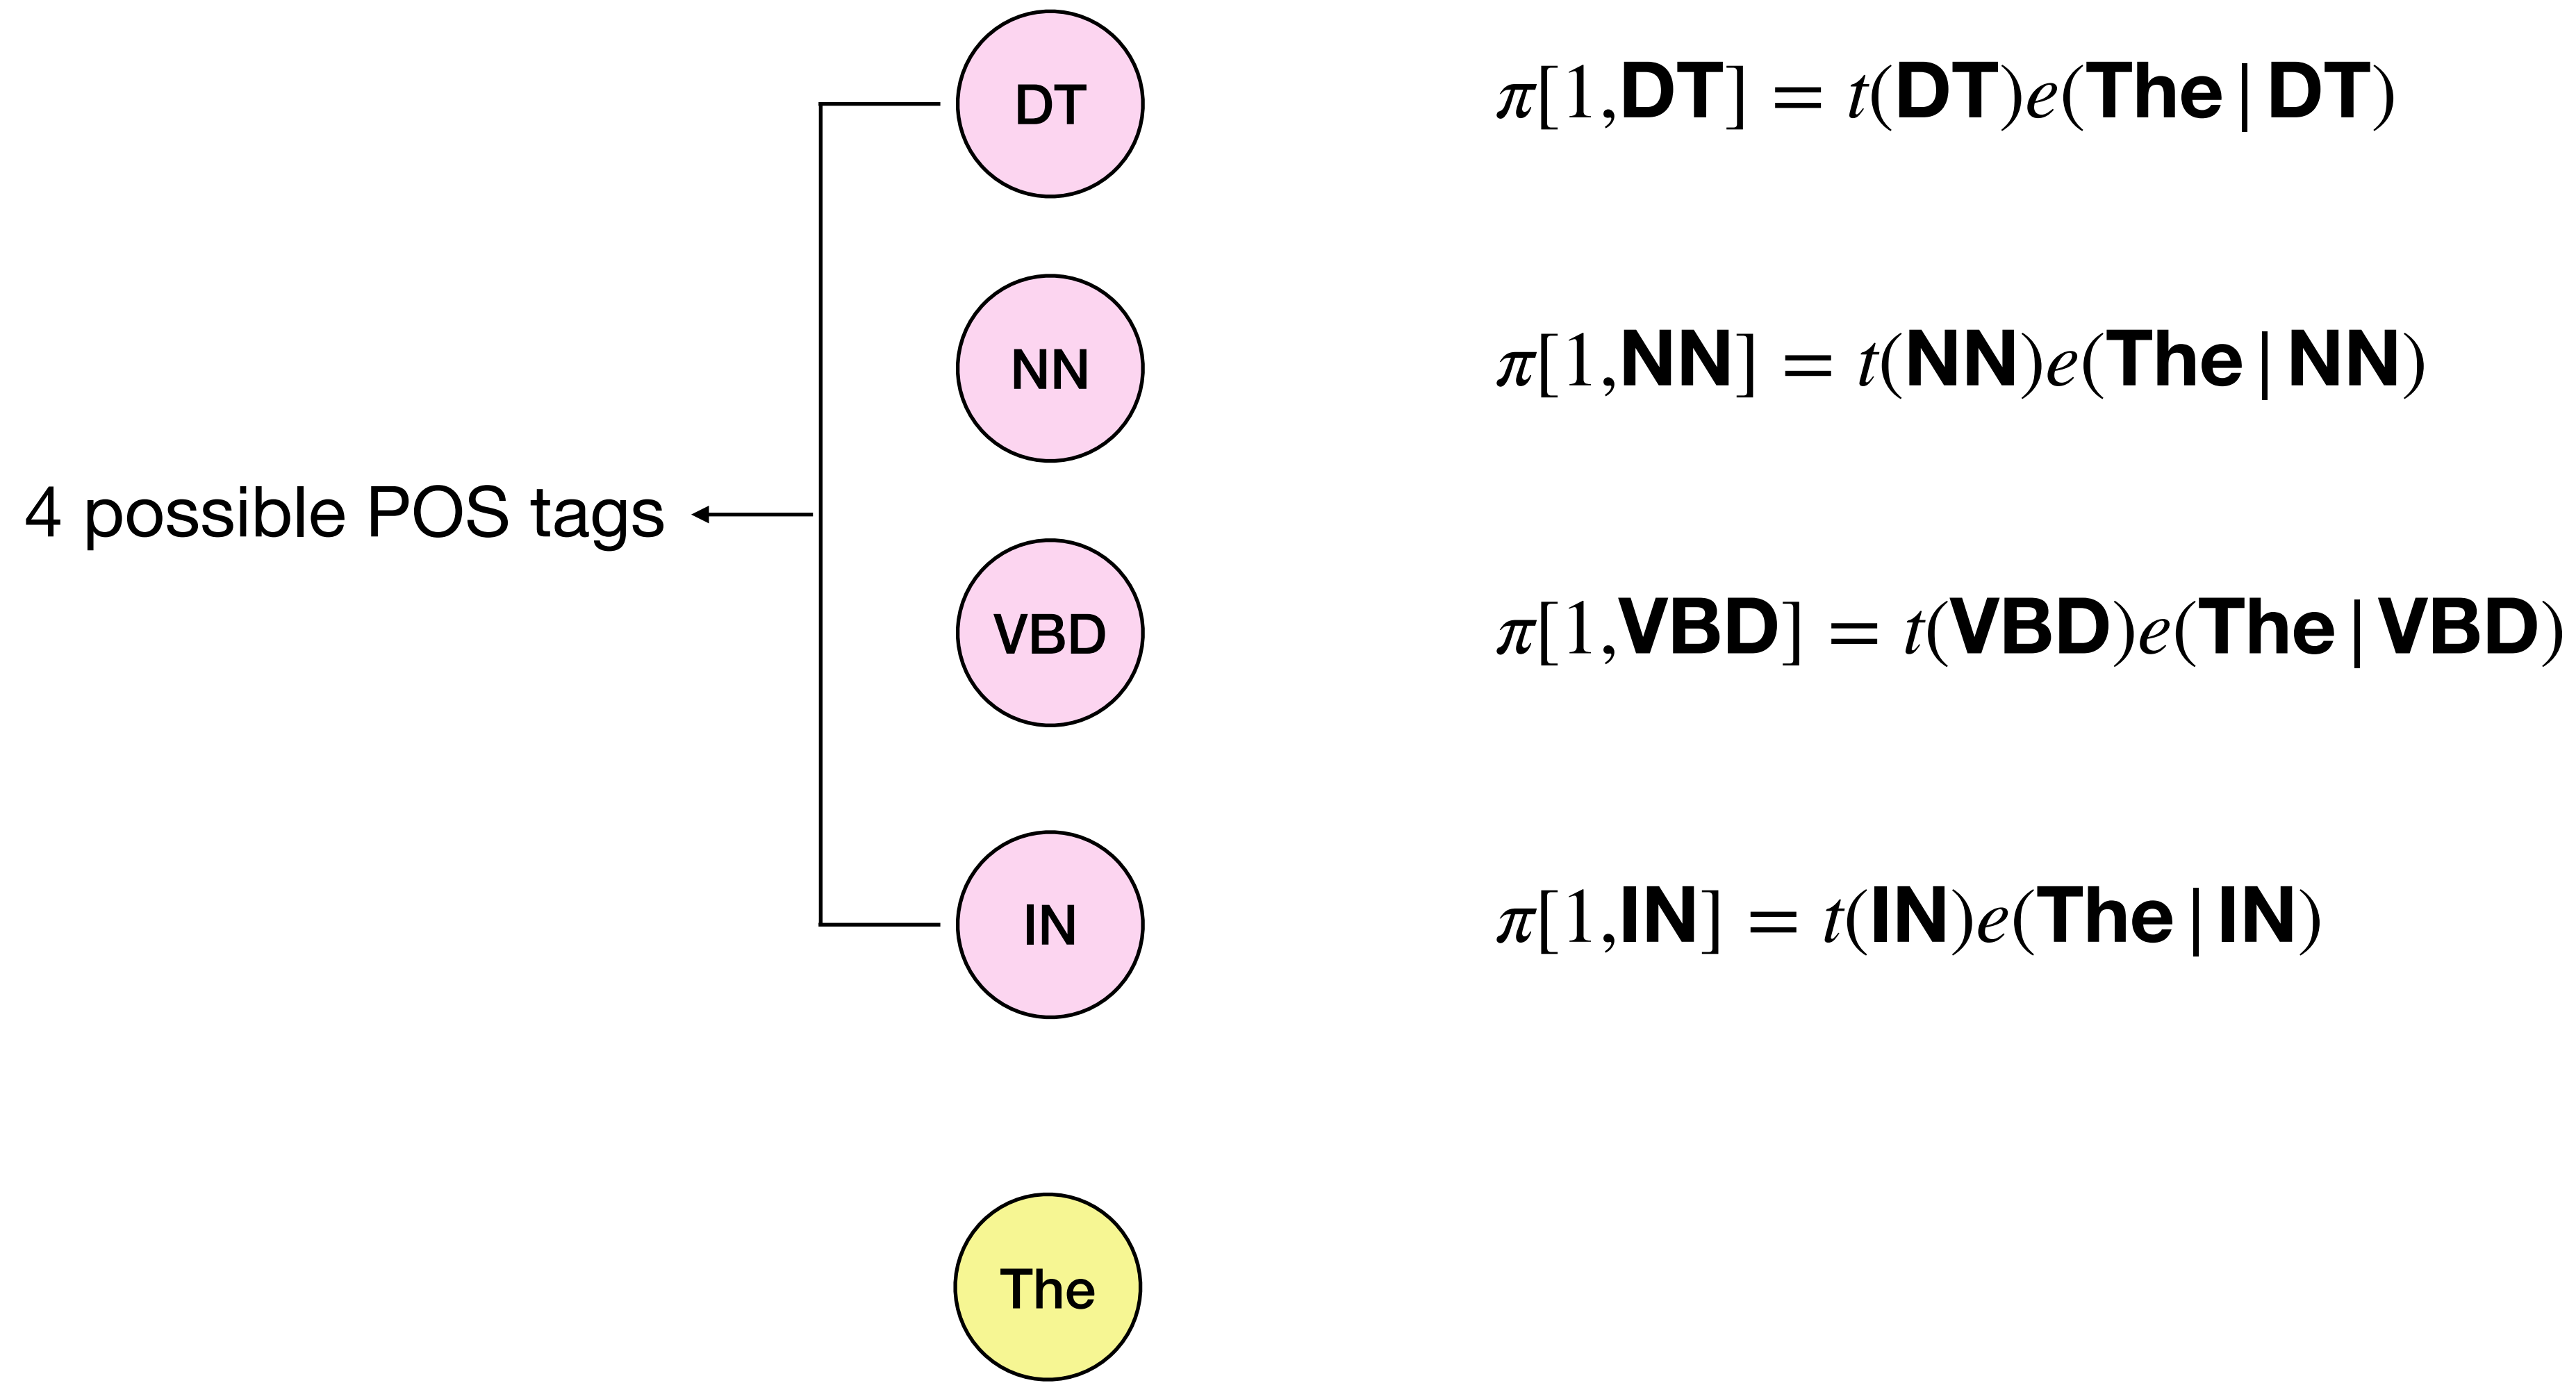
- ▶ The *Viterbi algorithm* is a dynamic programming algorithm.
Basic data structure:

$$\pi[j, s]$$


will be a table entry that stores the maximum probability for any state sequence ending in state s at position j . More formally: $\pi[1, s] = t(s)e(x_1|s)$, and for $j > 1$,

$$\pi[j, s] = \max_{s_1 \dots s_{j-1}} \left[t(s_1)e(x_1|s_1) \left(\prod_{k=2}^{j-1} t(s_k|s_{k-1})e(x_k|s_k) \right) \boxed{t(s|s_{j-1})e(x_j|s)} \right]$$

Viterbi Decoding

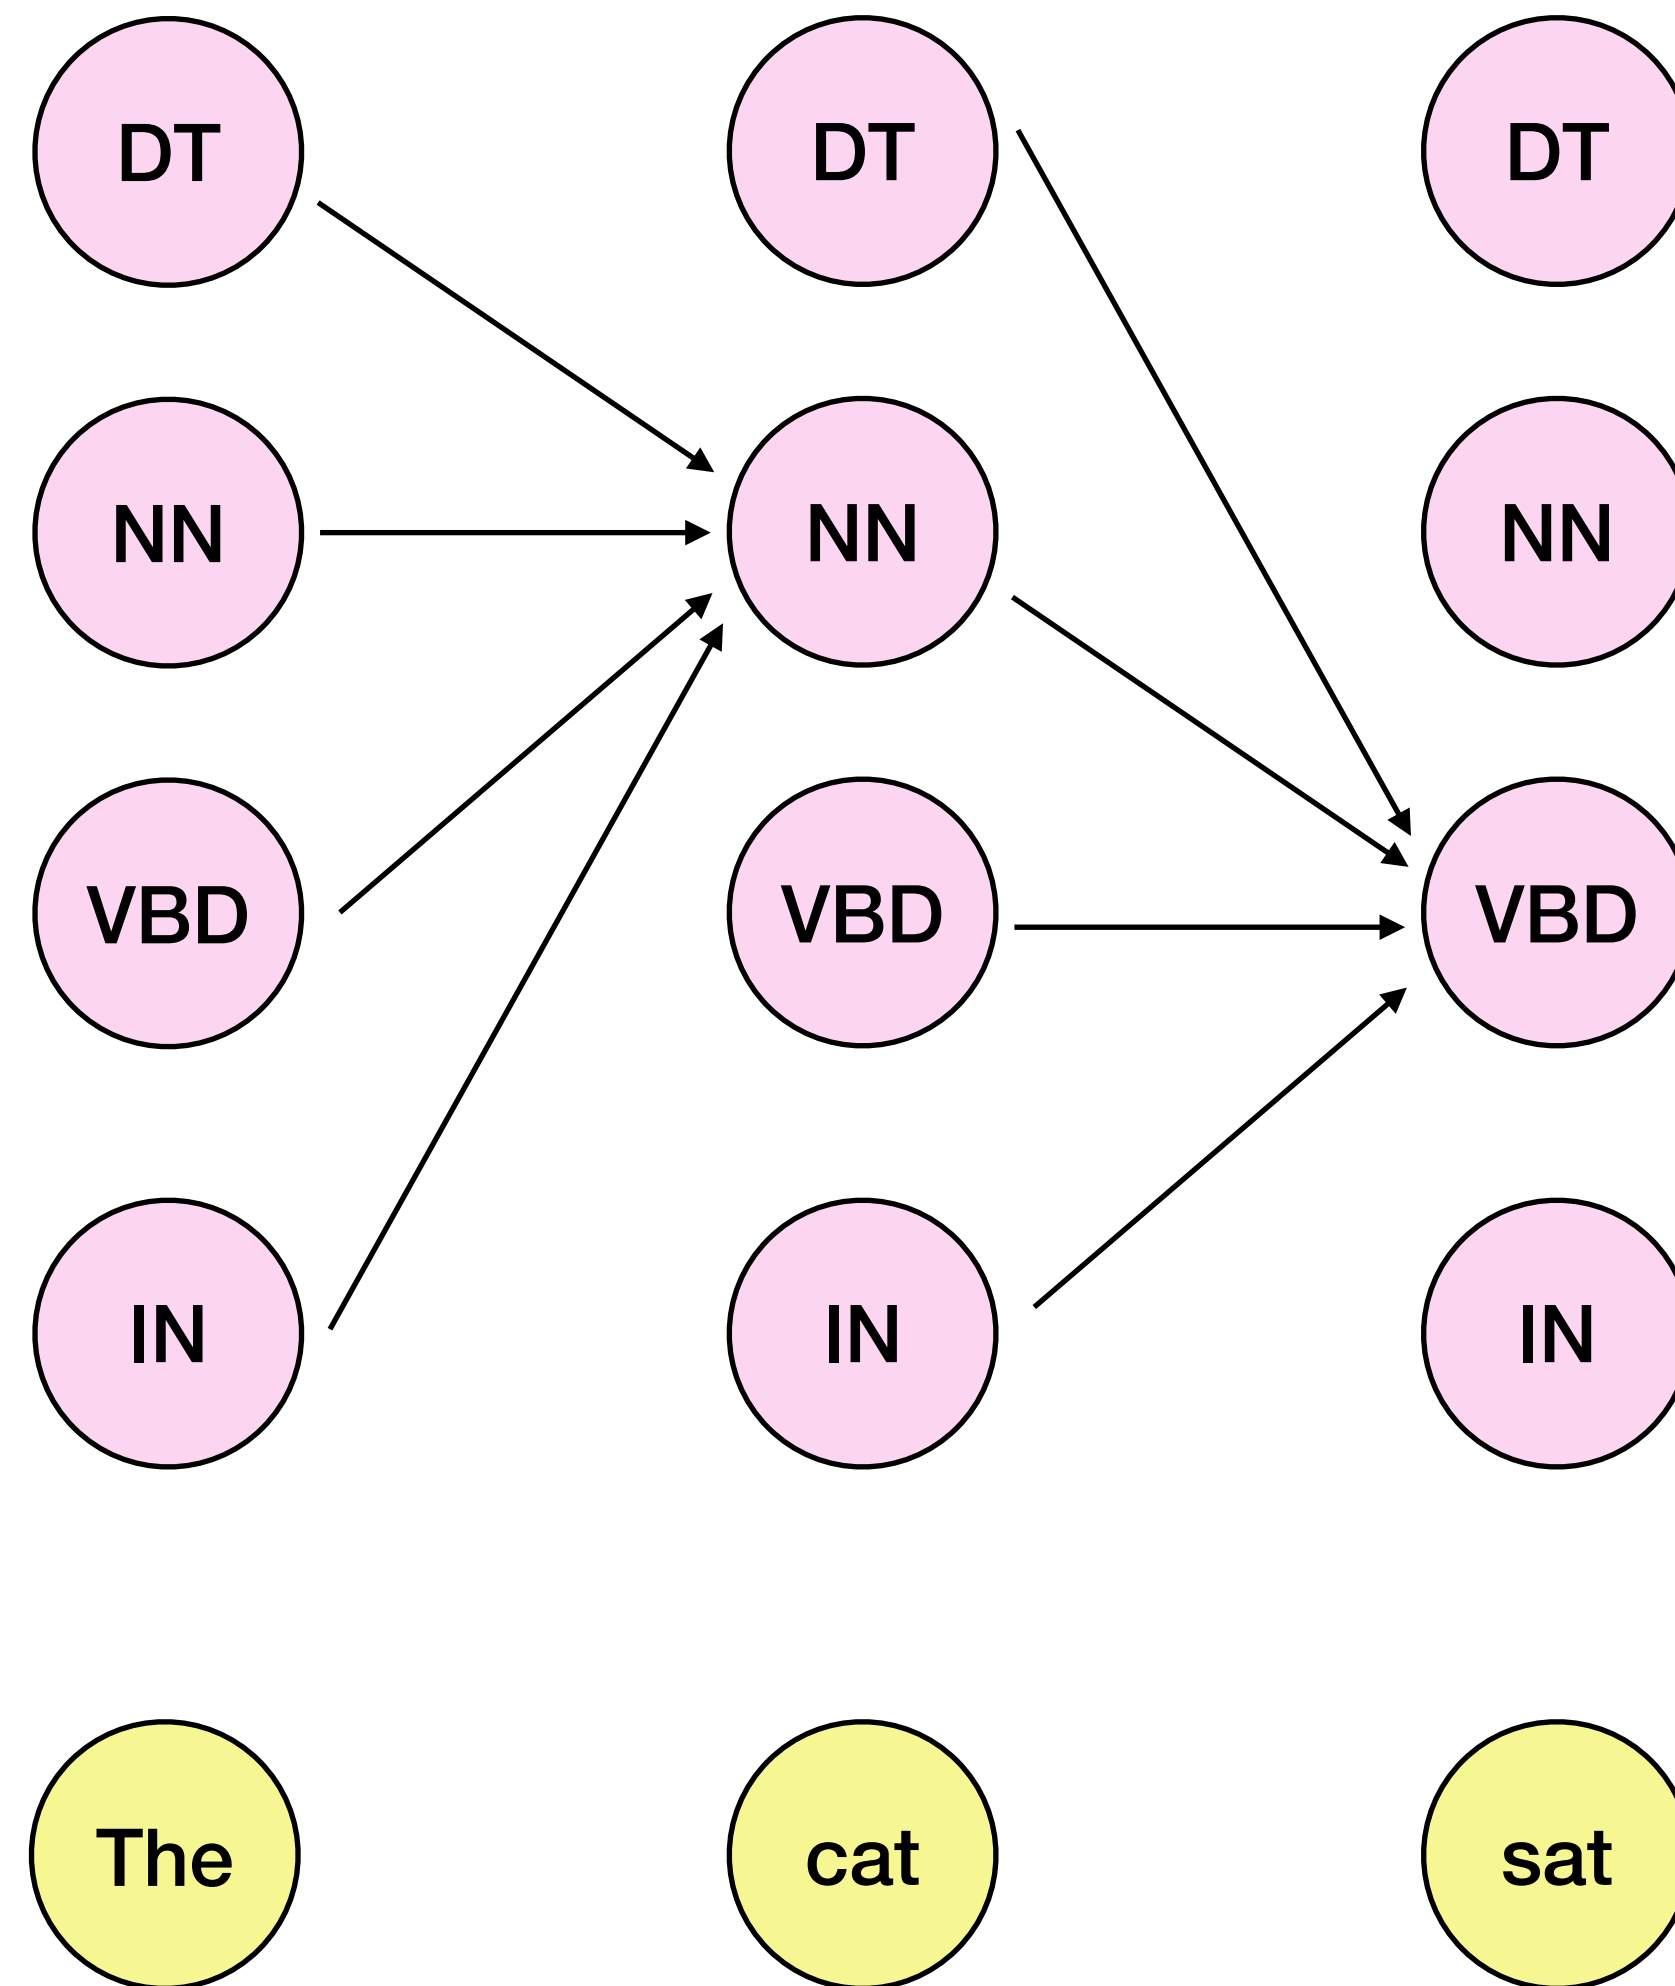


Viterbi Decoding



$$\pi[2, \mathbf{NN}] = \max_{s_1} \boxed{\pi[1, s_1]} t(\mathbf{NN} | s_1) e(\mathbf{cat} | \mathbf{NN})$$

Viterbi Decoding



$$\pi[2, \mathbf{NN}] = \max_{s_1} \pi[1, s_1] t(\mathbf{NN} | s_1) e(\mathbf{cat} | \mathbf{NN})$$

$$\pi[3, \mathbf{VBD}] = \max_{s_2} \pi[2, s_2] t(\mathbf{VBD} | s_2) e(\mathbf{sat} | \mathbf{VBD})$$

Viterbi Decoding

- ▶ Initialization: for $s = 1 \dots k$

$$\pi[1, s] = t(s)e(x_1|s)$$

- ▶ For $j = 2 \dots m$, $s = 1 \dots k$:

$$\pi[j, s] = \max_{s' \in \{1 \dots k\}} [\pi[j-1, s'] \times t(s|s') \times e(x_j|s)]$$

- ▶ We then have

$$\max_{s_1 \dots s_m} p(x_1 \dots x_m, s_1 \dots s_m; \underline{\theta}) = \max_s \pi[m, s]$$

- ▶ The algorithm runs in $O(mk^2)$ time

Pros and Cons

- **HMMs are simple to train**
 - Just need to compile counts from the training corpus
- **Performs relatively well**
 - > 96% on POS tagging (92.3% of most frequent class)
 - > 90% on Named Entity Recognition
- **Main difficulty is modeling $e(word | tag)$**
 - Words are very complex
 - Unknown words

Reading Materials

- **Notes from Michael Collins:**
 - Sequence Labeling and HMMs
 - EM Algorithm