


CSCI 544: Applied Natural Language Processing

Sequence Labeling-2

Xuezhe Ma (Max)



USC University of
Southern California

Recap

- **The Sequence Labeling Problem**
 - General Structured Prediction Tasks
 - Part-of-speech Tagging: A case study
- **Hidden Markov Model (HMM)**
 - Basic definitions
 - Parameter estimation: Maximum Likelihood Estimation (MLE)
 - The Viterbi algorithm

Recap: Structured Prediction

- Y consists of **multiple components** $Y = \{y_1, y_2, \dots, y_n\}$
- **(Strong) correlations** between output components
- **Exponential** output space
 - Decoding: $y^* = \operatorname{argmax}_{y \in \mathcal{Y}} p(y | x)$

麦克斯 在 南加大 工作

X



Max is working at USC

\uparrow
 Y_1

\uparrow
 Y_2


\uparrow
 Y_3

\uparrow
 Y_4

\uparrow
 Y_5


Recap: Sequence Labeling

A type of structured prediction tasks



Assigning each token of X , e.g. x_i a corresponding label y_i


Recap: Hidden Markov Models



- Set of states $S \in \{1, 2, \dots, k\}$ and set of observations (words) X
- Transition probabilities $P(s_{j+1} | s_j) = t(s_{j+1} | s_j)$
- Emission probabilities $P(x_j | s_j) = e(x_j | s_j)$
- Decoding with Viterbi algorithm

$$p(x_1, \dots, x_m, s_1, \dots, s_m) = t(s_1) \prod_{j=2}^m t(s_j | s_{j-1}) \prod_{j=1}^m e(x_j | s_j)$$

Recap: Hidden Markov Models



- A generative model with **strong assumptions**
 - Is generative model necessary?
- Simple to train
 - Just need to compile counts from the training corpus
- Performs relatively well
 - 96% on POS tagging (**92.3% of most frequent class**)
- Features only on **word type** and **tag**

$$p(x_1, \dots, x_m, s_1, \dots, s_m) = t(s_1) \prod_{j=2}^m t(s_j | s_{j-1}) \prod_{j=1}^m e(x_j | s_j)$$

Overview

- The Sequence Labeling Problem
 - General Structured Prediction Tasks
 - Part-of-speech Tagging: A case study
 - Maximum Likelihood Estimation (MLE)
- Hidden Markov Model (HMM)
 - Basic definitions
 - Parameter estimation
 - The Viterbi algorithm
- **Log-Linear Models**
 - Generative Models vs. Discriminative Models
 - General Form of Log-Linear Models
 - Maximum Entropy Markov Models (MEMMs)
 - Conditional Random Fields (CRFs)
 - Parameter estimation

Generative Models vs Discriminative Models

Generative vs Discriminative

- **Generative Models**

- Modeling the joint distribution: $P(X, S)$

- **Discriminative Models**

- Modeling $P(S | X)$ directly?

Generative

Classification

Naive Bayes: $P(y)P(x | y)$

Discriminative

Logistic Regression: $P(y | x)$

Sequence Labeling

HMM:

$P(s_1, \dots, s_n)P(x_1, \dots, x_n | s_1, \dots, s_n)$

MEMM/CRF:

$P(s_1, \dots, s_n | x_1, \dots, x_n)$

Log-Linear Models

Log-Linear Models

- The General Problem

- ▶ We have some **input domain** \mathcal{X}
- ▶ Have a finite **label set** \mathcal{Y}
- ▶ Aim is to provide a **conditional probability** $p(y \mid x)$ for any x, y where $x \in \mathcal{X}, y \in \mathcal{Y}$

Log-Linear Models

- ▶ We have some input domain \mathcal{X} , and a finite label set \mathcal{Y} . Aim is to provide a conditional probability $p(y \mid x)$ for any $x \in \mathcal{X}$ and $y \in \mathcal{Y}$.
- ▶ A feature is a function $f : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$
(Often binary features or indicator functions $f_k : \mathcal{X} \times \mathcal{Y} \rightarrow \{0, 1\}$).
- ▶ Say we have m features f_k for $k = 1 \dots m$
 \Rightarrow A feature vector $f(x, y) \in \mathbb{R}^m$ for any $x \in \mathcal{X}$ and $y \in \mathcal{Y}$.
- ▶ We also have a **parameter vector** $v \in \mathbb{R}^m$
- ▶ We define

$$p(y \mid x; v) = \frac{e^{v \cdot f(x, y)}}{\sum_{y' \in \mathcal{Y}} e^{v \cdot f(x, y')}}$$


Why the name?

$$p(y \mid x; v) = \frac{e^{v \cdot f(x, y)}}{\sum_{y' \in \mathcal{Y}} e^{v \cdot f(x, y')}}$$

$$\log p(y \mid x; v) = \underbrace{v \cdot f(x, y)}_{\text{Linear term}} - \underbrace{\log \sum_{y' \in \mathcal{Y}} e^{v \cdot f(x, y')}}_{\text{Normalization term}}$$

Features in Log-linear Models

$S = S_1, S_2, \dots, S_n$



$X = X_1, X_2, \dots, X_n$

USC

is

in

California

Theoretically, we can use any features in X and S : $f(X, S)$

- The current word: **<is>**
- The surrounding words: **<USC>**, **<in>**, ...
- The current POS tag: **<VBZ>**
- The surrounding tags: **<NNP>**, **<IN>**, ...

How to design these features into numerical vectors?

Binary Feature Vectors

$$f_1 = \begin{cases} 1, & \text{if } x_i = \text{is}, s_i = \text{VBZ} \\ 0, & \text{otherwise} \end{cases}$$

$$f_2 = \begin{cases} 1, & \text{if } x_{i-1} = \text{USC}, x_i = \text{is}, s_{i-1} = \text{NNP}, \text{ and } s_i = \text{VBZ} \\ 0, & \text{otherwise} \end{cases}$$

$$f_3 = \begin{cases} 1, & \text{if } x_{i-1} = \text{USC}, x_i = \text{is}, x_{i+1} = \text{in}, s_{i-1} = \text{NNP}, \text{ and } s_i = \text{VBZ} \\ 0, & \text{otherwise} \end{cases}$$

Task-Specific Features


- Spelling features for prefixes/suffixes

$$f_4 = \begin{cases} 1, & \text{if } x_i \text{ ends in } \textit{ing}, s_{i-1} = \text{NNP}, \text{ and } s_i = \text{VBZ} \\ 0, & \text{otherwise} \end{cases}$$

$$f_5 = \begin{cases} 1, & \text{if } x_i \text{ starts with } \textit{pre}, s_{i-1} = \text{NNP}, \text{ and } s_i = \text{VBZ} \\ 0, & \text{otherwise} \end{cases}$$

Features in Log-linear Models

$$S = S_1, S_2, \dots, S_n$$



$$X = X_1, X_2, \dots, X_n$$

USC

is

in

California

Theoretically, we can use any features in X and S : $f(X, S)$

- The current word: **<is>**
- The surrounding words: **<USC>**, **<in>**, ...
- The current POS tag: **<VBZ>**
- The surrounding tags: **<NNP>**, **<IN>**, ...

How to design these features into numerical vectors?

Can we design any features in practice?

Feature Sparsity

Number of Features

$$f_1 = \begin{cases} 1, & \text{if } x_i = \text{is}, s_i = \text{VBZ} \\ 0, & \text{otherwise} \end{cases}$$

VK

$$f_2 = \begin{cases} 1, & \text{if } x_{i-1} = \text{USC}, x_i = \text{is}, s_{i-1} = \text{NNP}, \text{ and } s_i = \text{VBZ} \\ 0, & \text{otherwise} \end{cases}$$

V^2K^2

$$f_3 = \begin{cases} 1, & \text{if } x_{i-1} = \text{USC}, x_i = \text{is}, x_{i+1} = \text{in}, s_{i-1} = \text{NNP}, \text{ and } s_i = \text{VBZ} \\ 0, & \text{otherwise} \end{cases}$$

V^3K^2

Features vs. Independence

- We need independence assumptions to compute the nominator
- Stronger assumptions lead to less flexible features

$$p(y \mid x; v) = \frac{e^{v \cdot f(x, y)}}{\sum_{y' \in \mathcal{Y}} e^{v \cdot f(x, y')}} \quad \boxed{\phantom{\sum_{y' \in \mathcal{Y}} e^{v \cdot f(x, y')}}}$$

Overview

- The Sequence Labeling Problem
 - General Structured Prediction Tasks
 - Part-of-speech Tagging: A case study
 - Generative Models vs. Discriminative Models
 - Maximum Likelihood Estimation (MLE)
- Hidden Markov Model (HMM)
 - Basic definitions
 - Parameter estimation
 - The Viterbi algorithm
- **Log-Linear Models**
 - Generative Models vs. Discriminative Models
 - General Form of Log-Linear Models
 - **Maximum Entropy Markov Models (MEMMs)**
 - Conditional Random Fields (CRFs)
 - Parameter estimation

Maximum-Entropy Markov Models (MEMMs)

- Goal: modeling the distribution

$$p(s_1, \dots, s_m \mid x_1, \dots, x_m)$$

Independence Assumptions in MEMMs

- Markov Assumption on S

$$p(s_1, \dots, s_m | x_1, \dots, x_m) = \prod_{j=1}^m p(s_j | s_1, \dots, s_{j-1}, x_1, \dots, x_m)$$

chain rule (no assumptions)

$$= \prod_{j=1}^m p(s_j | s_{j-1}, x_1, \dots, x_m)$$

Markov assumption

Using Log-Linear Models

- We then model each term using a log-linear model:

$$p(s_j | s_{j-1}, x_1, \dots, x_m) = \frac{\exp(v \cdot f(x_1, \dots, x_m, i, s_{j-1}, s_j))}{\sum_{s'_j \in \mathcal{S}} \exp(v \cdot f(x_1, \dots, x_m, i, s_{j-1}, s'_j))}$$

- Here $f(x_1, \dots, x_m, j, s, s')$ is the feature vector:
 - x_1, \dots, x_m is the sequence of words to be tagged
 - j is the position to be tagged (any value from $1, \dots, m$)
 - s is the previous state
 - s' is the new state

Using Log-Linear Models

- We then model each term using a log-linear model:

$$p(s_j | s_{j-1}, x_1, \dots, x_m) = \frac{\exp(v \cdot f(x_1, \dots, x_m, \overset{\cdot}{i}, s_{j-1}, s_j))}{\sum_{s'_j \in \mathcal{S}} \exp(v \cdot f(x_1, \dots, x_m, \underset{==j}{i}, s_{j-1}, s'_j))}$$

Trackable

- Here $f(x_1, \dots, x_m, j, s, s')$ is the feature vector:

- x_1, \dots, x_m is the sequence of words to be tagged
- j is the position to be tagged (any value from $1, \dots, m$)
- s is the previous state
- s' is the new state

The whole sequence of X
Only two successive tags

Features in MEMMs

$$S = S_1, S_2, \dots, S_n$$

$$X = X_1, X_2, \dots, X_n$$



USC

is

in

California

What are the most important features $f(x_1, \dots, x_m, j, s, s')$?

- The current word: **<is>**
- The current tag: **<VBZ>**
- The surrounding words: **<USC>**, **<in>**
- The previous POS tag: **<NNP>**
- Prefix or suffix features
- ...

Decoding with MEMMs: Viterbi Algorithm

- ▶ Goal: for a given input sequence x_1, \dots, x_m , find

$$\arg \max_{s_1, \dots, s_m} p(s_1 \dots s_m | x_1 \dots x_m)$$

- ▶ We can use the *Viterbi* algorithm again (see last lecture on HMMs). Basic data structure:

$$\pi[j, s]$$

will be a table entry that stores the maximum probability for any state sequence ending in state s at position j . More formally:

$$\pi[j, s] = \max_{s_1 \dots s_{j-1}} \left(p(s | s_{j-1}, x_1 \dots x_m) \prod_{k=1}^{j-1} p(s_k | s_{k-1}, x_1 \dots x_m) \right)$$

Decoding with MEMMs: Viterbi Algorithm

- Initialization: for $s \in \mathcal{S}$

$$\pi[1, s] = p(s|s_0, x_1 \dots x_m)$$

where s_0 is a special “initial” state.

- For $j = 2 \dots m$, $s = 1 \dots k$:

$$\pi[j, s] = \max_{s' \in \mathcal{S}} [\pi[j - 1, s'] \times p(s|s', x_1 \dots x_m)]$$

- We then have

$$\max_{s_1 \dots s_m} p(s_1 \dots s_m | x_1 \dots x_m) = \max_s \pi[m, s]$$

Model Performance

	POS Tagging	NER
HMM	96.4%	75.3
MEMM	96.9%	85.9

HMMs vs. MEMMs

- In MEMMs, each state transition has probability

$$p(s_j | s_{j-1}, x_1, \dots, x_m) = \frac{\exp(v \cdot f(x_1, \dots, x_m, i, s_{j-1}, s_j))}{\sum_{s' \in S} \exp(v \cdot f(x_1, \dots, x_m, i, s_{j-1}, s'_j))}$$

- In HMMs, each state transition has probability

$$p(s_j | s_{j-1})p(x_j | s_j)$$

- Feature vectors f allows much richer representations in MEMMs:

- Sensitivity to *any* word in the input sequence x_1, \dots, x_m , not just x_j
- Sensitivity to spelling features (prefixes, suffixes etc.) of current or surrounding words

- Parameter estimation in MEMMs is more expensive than in HMMs (but is still not prohibitive for most tasks)

Can we relax the Markov assumption in MEMMs but keep the same features?

Overview

- The Sequence Labeling Problem
 - General Structured Prediction Tasks
 - Part-of-speech Tagging: A case study
 - Generative Models vs. Discriminative Models
 - Maximum Likelihood Estimation (MLE)
- Hidden Markov Model (HMM)
 - Basic definitions
 - Parameter estimation
 - The Viterbi algorithm
- **Log-Linear Models**
 - Maximum Entropy Markov Models (MEMMs)
 - **Conditional Random Fields (CRFs)**
 - Parameter estimation

Conditional Random Fields (CRFs)

- Goal: modeling the distribution

$$p(s_1, \dots, s_m \mid x_1, \dots, x_m)$$

- In MEMMs we had

$$p(s_1, \dots, s_m \mid x_1, \dots, x_m) = \prod_{j=1}^m p(s_j \mid s_1, \dots, s_{j-1}, x_1, \dots, x_m)$$

chain rule (no assumptions)

$$= \prod_{j=1}^m p(s_j \mid s_{j-1}, x_1, \dots, x_m)$$

Markov assumption

- Using log-linear model

$$p(s_j \mid s_{j-1}, x_1, \dots, x_m) = \frac{\exp(v \cdot f(x_1, \dots, x_m, i, s_{j-1}, s_j))}{\sum_{s' \in \mathcal{S}} \exp(v \cdot f(x_1, \dots, x_m, i, s_{j-1}, s'_j))}$$


Can we build a *giant* log-linear model?

$$p(s_1, \dots, s_m \mid x_1, \dots, x_m)$$


Conditional Random Fields (CRFs)

- Globally Normalized Model


$$p(s_1, \dots, s_m | x_1, \dots, x_m) = \frac{\prod_{j=1}^m \exp(v \cdot f(x_1, \dots, x_m, i, s_{j-1}, s_j))}{\sum_{s'_1, \dots, s'_m \in \mathcal{S}} \prod_{j=1}^m \exp(v \cdot f(x_1, \dots, x_m, i, s'_{j-1}, s'_j))}$$



HMM




MEMM



CRF

Independence Assumptions in CRFs


$$p(s_1, \dots, s_m | x_1, \dots, x_m) = \frac{\prod_{j=1}^m \exp(v \cdot f(x_1, \dots, x_m, i, s_{j-1}, s_j))}{\sum_{s'_1, \dots, s'_m \in \mathcal{S}} \prod_{j=1}^m \exp(v \cdot f(x_1, \dots, x_m, i, s'_{j-1}, s'_j))}$$



CRF

Independence Assumptions in CRFs

$$p(s_1, \dots, s_m | x_1, \dots, x_m) = \frac{\prod_{j=1}^m \exp(v \cdot f(x_1, \dots, x_m, i, s_{j-1}, s_j))}{\sum_{s'_1, \dots, s'_m \in \mathcal{S}} \prod_{j=1}^m \exp(v \cdot f(x_1, \dots, x_m, i, s'_{j-1}, s'_j))}$$



CRF

weaker than MEMMs!

Decoding with CRFs

- Viterbi Algorithm still applicable!

$$p(s_1, \dots, s_m | x_1, \dots, x_m) = \frac{\prod_{j=1}^m \exp(v \cdot f(x_1, \dots, x_m, i, s_{j-1}, s_j))}{\sum_{s'_1, \dots, s'_m \in \mathbb{S}} \prod_{j=1}^m \exp(v \cdot f(x_1, \dots, x_m, i, s'_{j-1}, s'_j))}$$

Makes no effects on decoding!

Computation of the Global Nominator

- How to compute the global nominator?
 - Dynamic programming similar to the Viterbi algorithm
 - Replacing the maximum operation in decoding with sum operation

$$p(s_1, \dots, s_m | x_1, \dots, x_m) = \frac{\prod_{j=1}^m \exp(v \cdot f(x_1, \dots, x_m, i, s_{j-1}, s_j))}{\sum_{s'_1, \dots, s'_m \in \mathcal{S}} \prod_{j=1}^m \exp(v \cdot f(x_1, \dots, x_m, i, s'_{j-1}, s'_j))} \quad \text{Partition Function}$$

$$\pi[j, s] = \sum_{s_1, \dots, s_{j-1}} \left[\prod_{k=1}^{j-1} \exp(v \cdot f(x_1, \dots, x_m, k, s_{k-1}, s_k)) \right] \exp(v \cdot f(x_1, \dots, x_m, k, s_{j-1}, s))$$

Overview

- The Sequence Labeling Problem
 - General Structured Prediction Tasks
 - Part-of-speech Tagging: A case study
 - Generative Models vs. Discriminative Models
 - Maximum Likelihood Estimation (MLE)
- Hidden Markov Model (HMM)
 - Basic definitions
 - Parameter estimation
 - The Viterbi algorithm
- **Log-Linear Models**
 - Maximum Entropy Markov Models (MEMMs)
 - Conditional Random Fields (CRFs)
 - **Parameter estimation**

Maximum Likelihood Estimation

- **Need to maximize:**

$$\begin{aligned}\max_v L(v) &= \sum_{i=1}^N \log P(S_i | X_i; v) \\ &= \sum_{i=1}^N v \cdot f(X_i, S_i) - \sum_{i=1}^N \log \sum_{s' \in \mathbb{S}} e^{v \cdot f(X_i, S')}\end{aligned}$$

- **Calculating gradients:**

$$\frac{\partial L(v)}{\partial v_k} = \underbrace{\sum_{i=1}^N f_k(X_i, S_i)}_{\text{Empirical counts}} - \underbrace{\sum_{i=1}^N \sum_{s' \in \mathbb{S}} f_k(X_i, S') p(S' | X_i; v)}_{\text{Expected counts}}$$

See Collin's notes for derivations

Model Performance

	POS Tagging	NER
HMM	96.4%	75.3
MEMM	96.9%	85.9
CRF	97.3%	88.7

Reading Materials

- Notes from Michael Collins:
 - Log-linear Models
 - MEMMs and CRFs