


CSCI 544: Applied Natural Language Processing

Self-attention & Transformer

Xuezhe Ma (Max)



USC University of
Southern California


Logistic Points

- **Midterm:**

- In-person unless approved to participate remotely
- Remote students: sign up in this Google Sheet <https://docs.google.com/spreadsheets/d/1BjmjBs6tceW5RHLj1-nskg9m79limq2BzJFVxDYRXcE/edit#gid=0>
- We will email questions to students only in this list
- Camera and microphone should be open with no virtual background (please find a quite location)
- Please follow the discussions on the Slack channel, which are sometimes useful


Recap: Seq2seq Generation

- **Sequence-to-Sequence (Seq2seq) Generation**
 - Input: $X = \{x_1, x_2, \dots, x_L\}, x_i \in \mathcal{X}$
 - Output: $Y = \{y_1, y_2, \dots, y_T\}, y_i \in \mathcal{Y}$
 - Model: $p_\theta(Y|X)$
- **Difference from Sequence Labeling**
 - The length of Y can be different from the length of X
 - The size of \mathcal{Y} is often much larger



Recap: Autoregressive Seq2seq Generation

- Autoregressive Factorization:

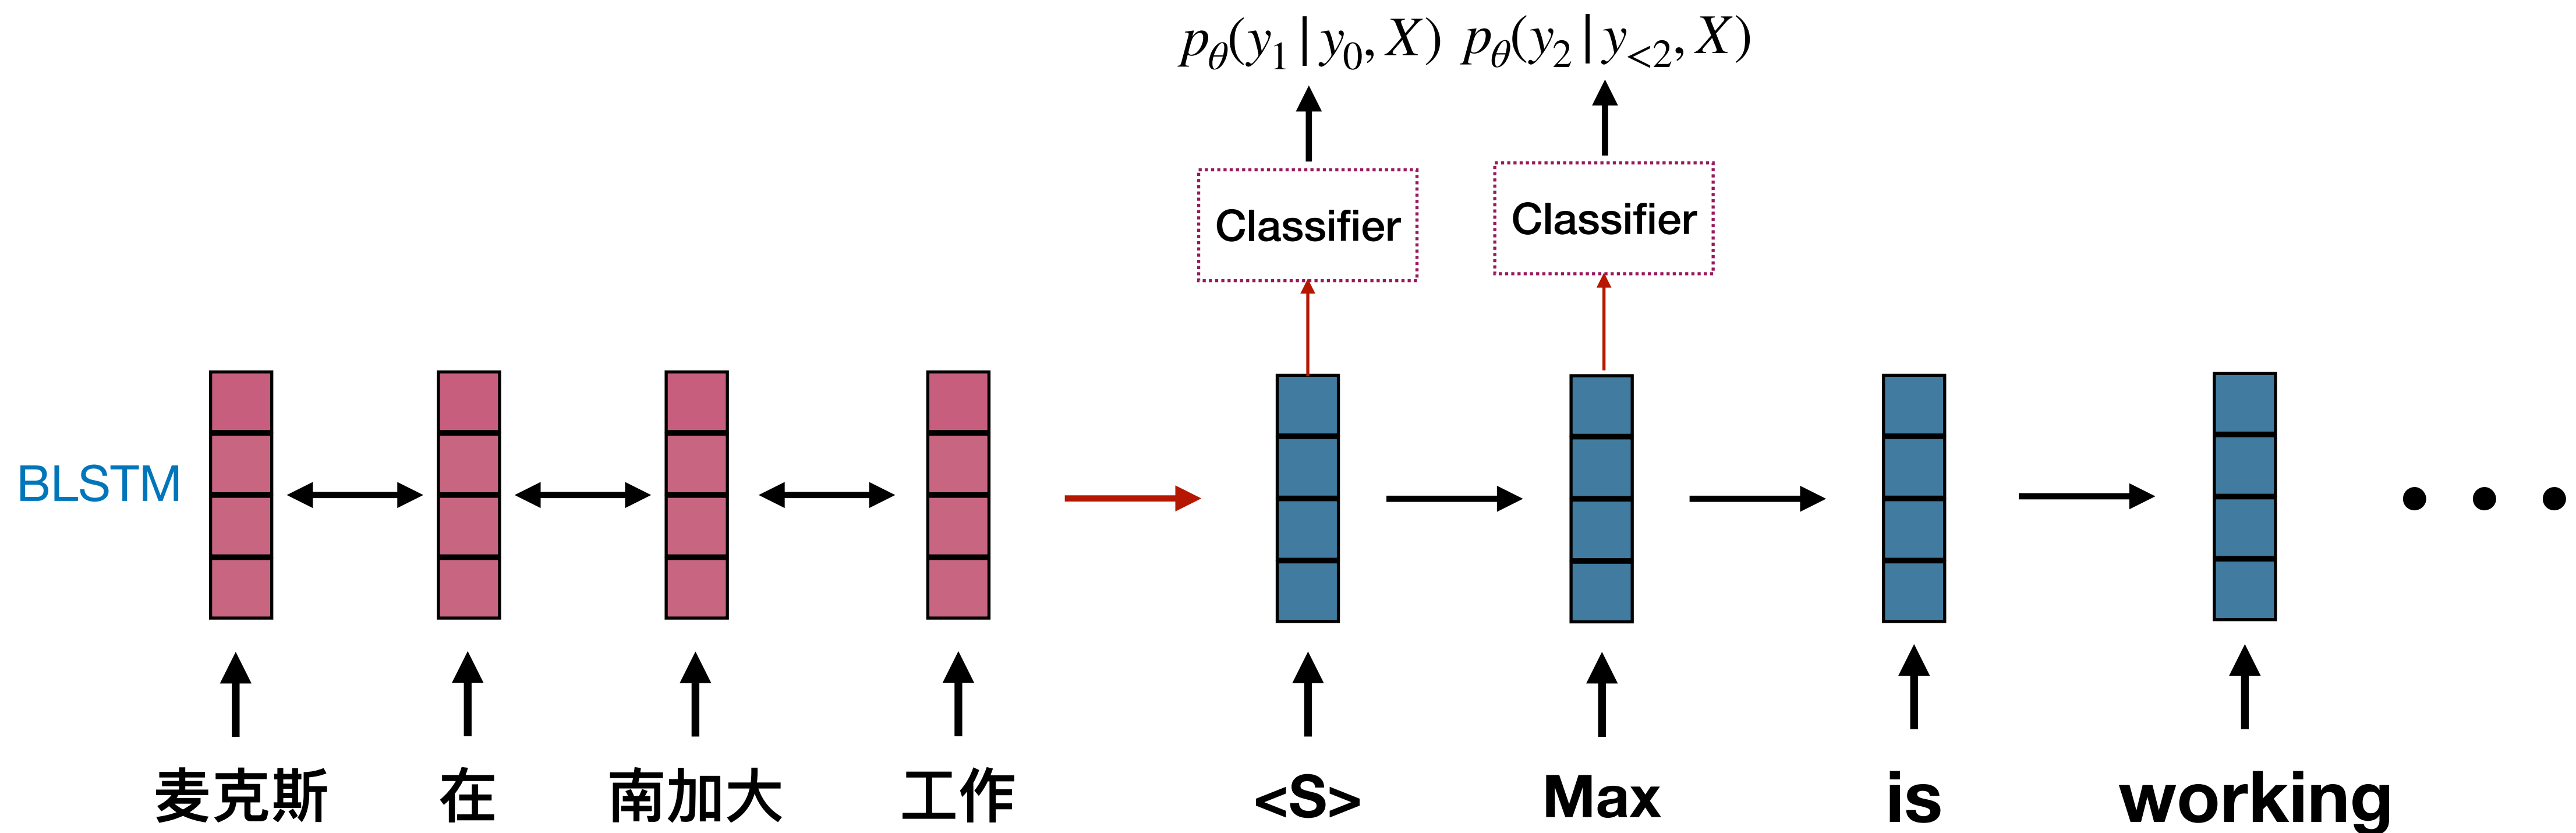
$$p_{\theta}(Y|X) = \prod_{t=1}^T p_{\theta}(y_t | y_{<t}, X)$$


- Autoregressive factorization is just chain-rule (HMMs, MEMMs)
- Autoregressive factorization does **NOT** assume any independence
- With autoregressive factorization, we need to model each $p_{\theta}(y_t | y_{<t}, X)$


Recap: Encoder-Decoder Architecture

- **Two Components:**


- **Encoder:** Convert input sequence into a sequence of vectors
- **Decoder:** Convert encoding into a sequence in the output space



Recap: MT Progress



Recap: Attention Mechanism



Recap: Types of Attention

- Dot-product attention (assumes equal dimensions for c and h)

$$\text{sim}(c_j, h_t) = c_j^T h_t$$

- Multiplicative attention

$$\text{sim}(c_j, h_t) = c_j^T W h_t, \text{ where } W \text{ is learnable weight matrix}$$

- Additive attention

$$\text{sim}(c_j, h_t) = v^T \tanh(W_c c_j + W_h h_t)$$


where W_c and W_h are learnable weight matrices and v is a learnable weight vector

Google's Neural Machine Translation System: Bridging the Gap between Human and Machine Translation

Table 10: Mean of side-by-side scores on production data

	PBMT	GNMT	Human	Relative Improvement
English → Spanish	4.885	5.428	5.504	87%
English → French	4.932	5.295	5.496	64%
English → Chinese	4.035	4.594	4.987	58%
Spanish → English	4.872	5.187	5.372	63%
French → English	5.046	5.343	5.404	83%
Chinese → English	3.694	4.263	4.636	60%

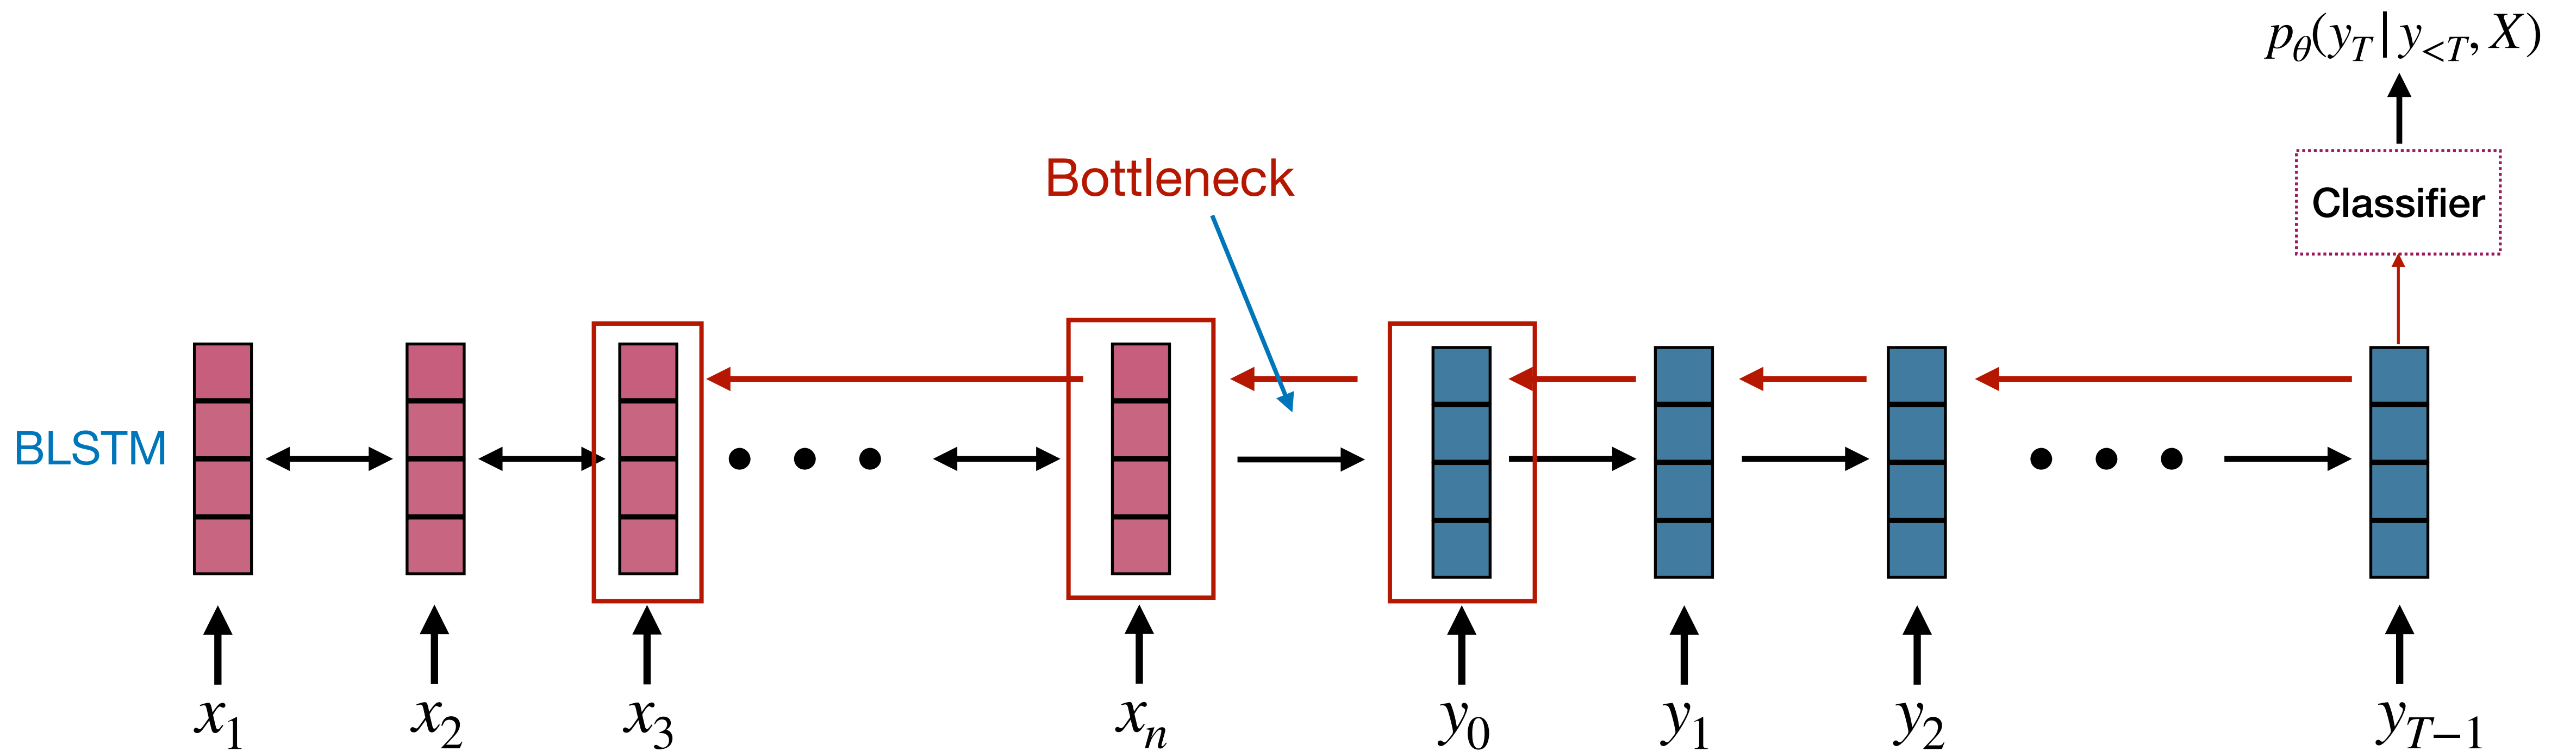
MT Progress



Self-Attention & Transformer

Revisit: Motivation of Attention Mechanism


- A single encoding vector needs to capture **all the information** about source sentence
- Longer sequences can lead to **vanishing gradients**



Issues with RNNs

- **One vector** to memorize all historical information


Inherent Markov Property



Issues with RNNs

- **One vector** to memorize all historical information
- Hard to capture long-distance information: **vanishing gradients**

Inherent Markov Property



Attention is the key to solve the problem!

This Lecture

- Do we really need RNNs to model the arbitrary context?
- **Maybe attention is all you need!**

Attention Is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Jakob Uszkoreit*
Google Research
usz@google.com

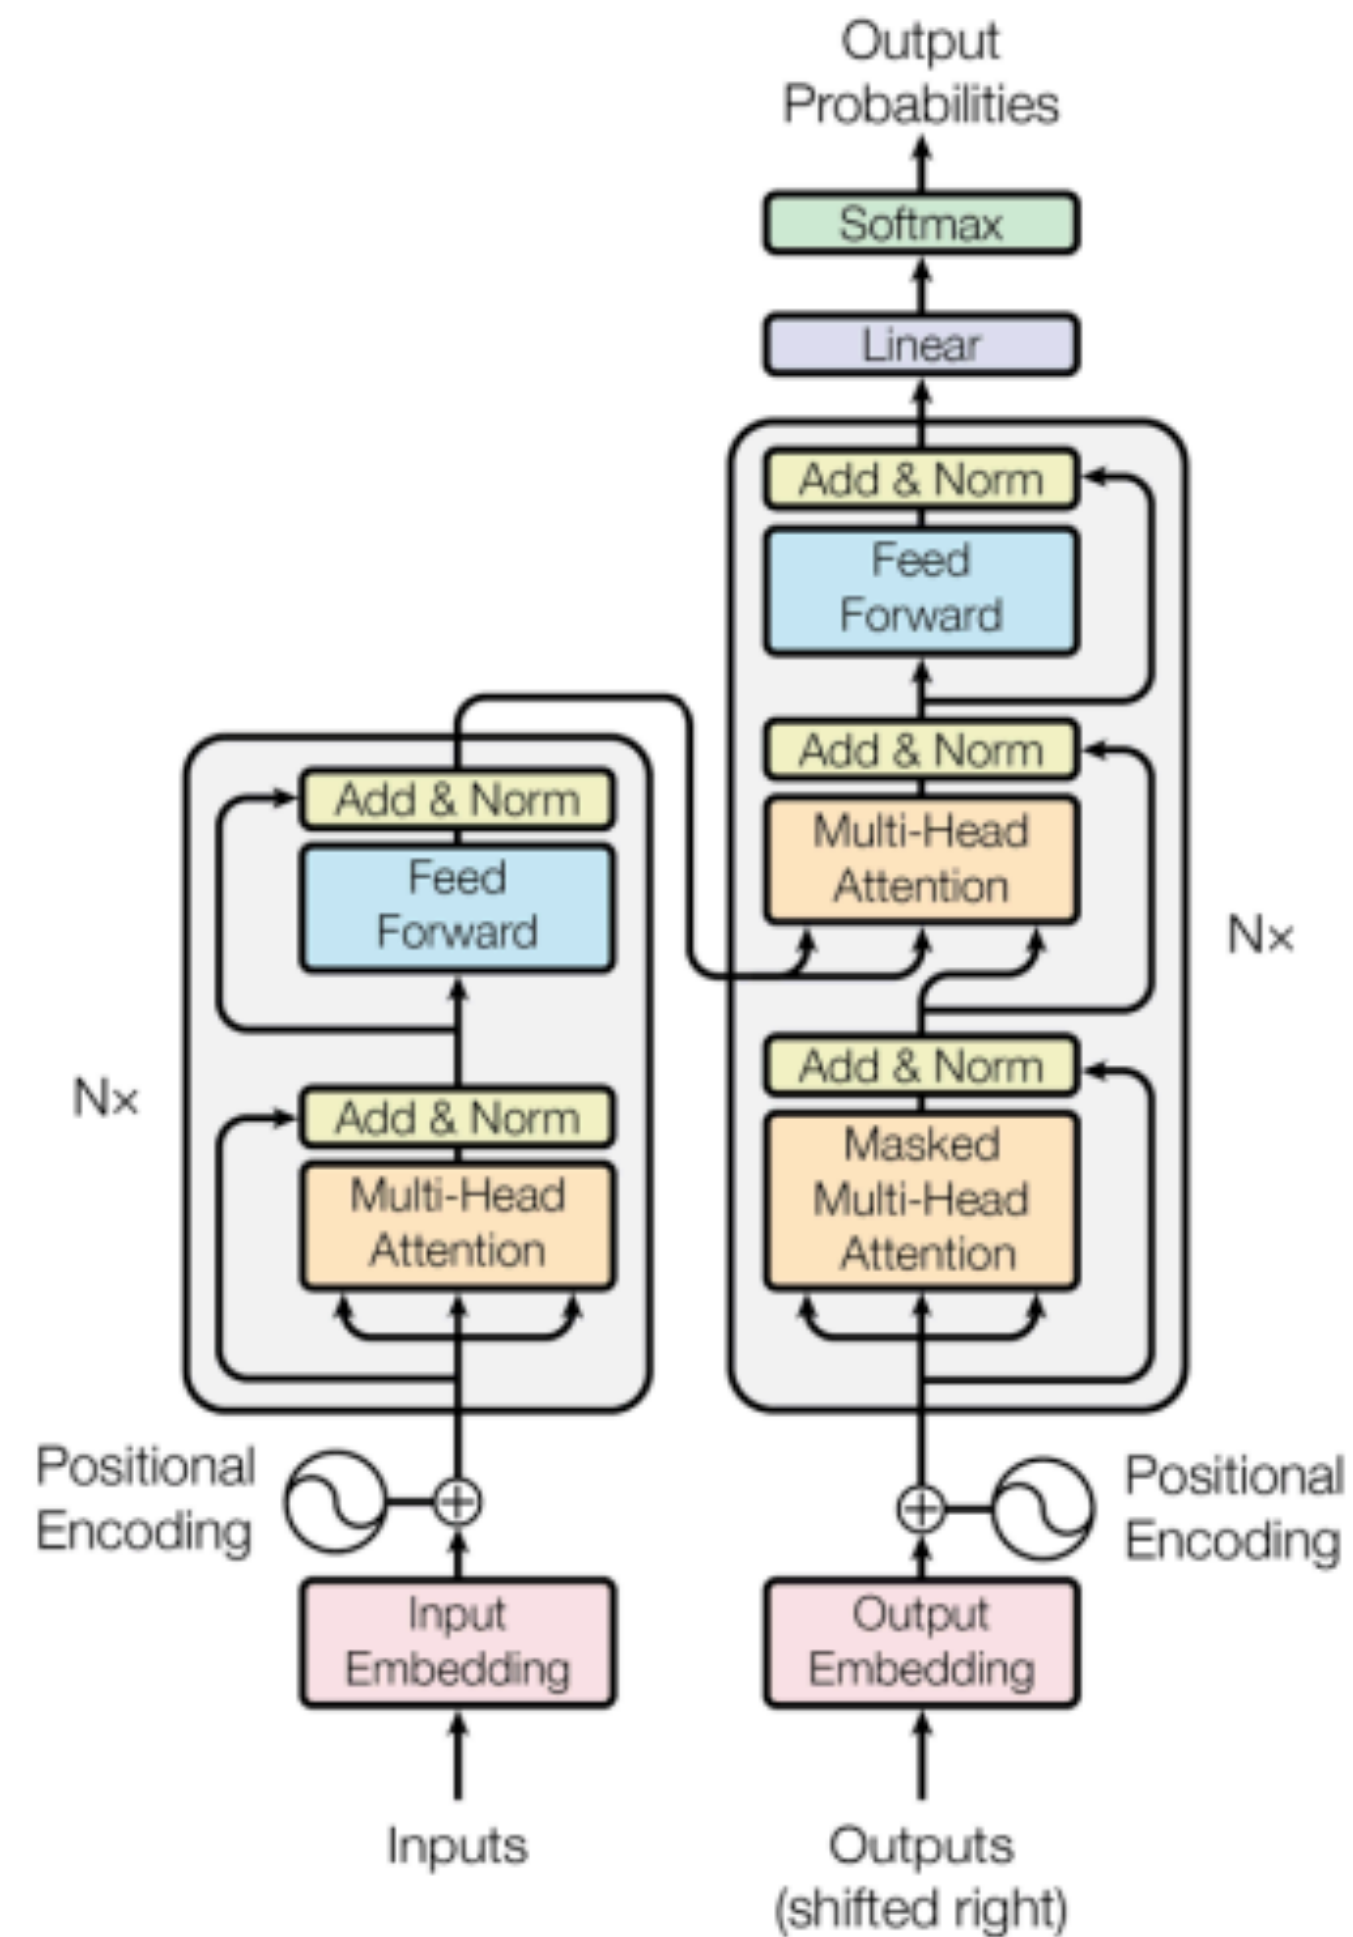
Llion Jones*
Google Research
llion@google.com

Aidan N. Gomez* †
University of Toronto
aidan@cs.toronto.edu

Łukasz Kaiser*
Google Brain
lukaszkaizer@google.com

Illia Polosukhin* ‡
illia.polosukhin@gmail.com

Transformers



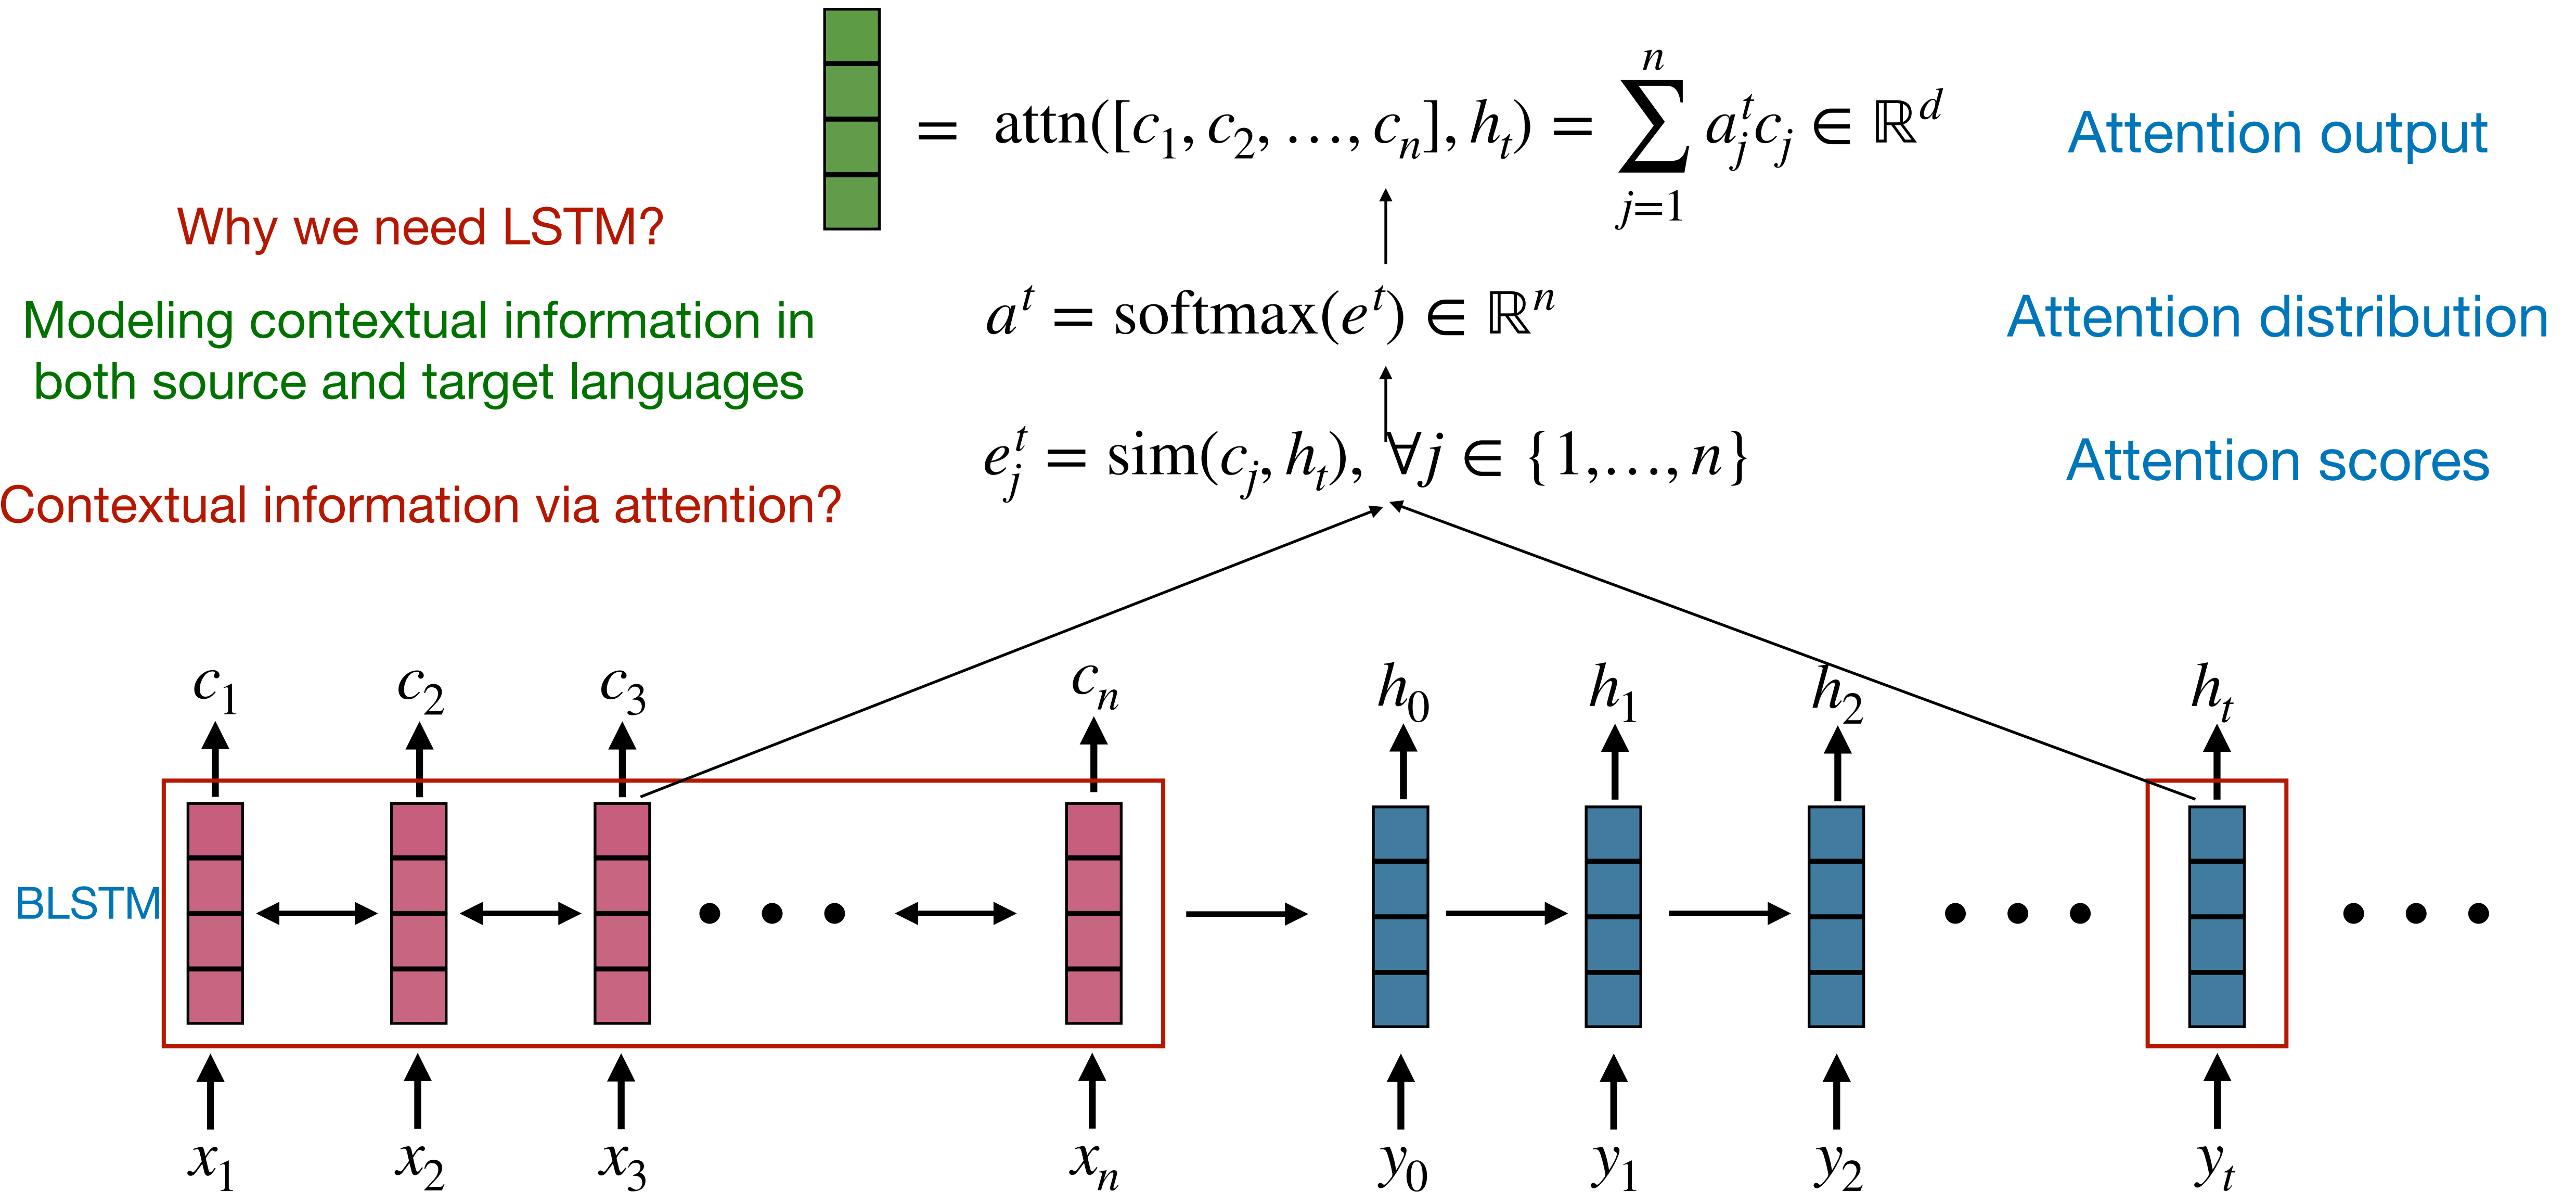
- Consists of an encoder and a decoder
- Originally proposed for neural machine translation and later adapted for almost all the NLP tasks
 - For example, BERT only uses the **encoder** of the Transformer architecture (**next lecture**)
- Both encoder and decoder consist of N layers
 - Each encoder layer has two sub-layers
 - Each decoder layer has three sublayers
 - Key innovation: **multi-head self-attention**

Recap: Attention Mechanism

Why we need LSTM?


Modeling contextual information in both source and target languages

Contextual information via attention?




Self-Attention

- **Self-attention: attention within on single sequence**
 - Contexts and queries are drawn from the same source
- **Contextual information via self-attention**



Self-Attention


- **Self-attention: attention within on single sequence**
 - Contexts and queries are drawn from the same source
- **Contextual information via self-attention**



Self-Attention

- **Self-attention: attention within on single sequence**
 - Contexts and queries are drawn from the same source
- **Contextual information via self-attention**

- Capturing long-distance dependencies
- No gradient vanishing



Self-attention in equations

- A sequence of input vectors $x_1, \dots, x_n \in \mathbb{R}^d$
- First, construct a set of **queries**, **keys** and **values**:

$$q_i = W_Q x_i, k_i = W_K x_i, v_i = W_V x_i$$

- Second, for each q_i , compute attention scores and attention distributions:

$$a_{i,j} = \text{softmax}\left(\frac{q_i^T k_j}{\sqrt{d}}\right) \quad \text{aka. "scaled dot product"}$$

- Finally, compute the weighted sum:

$$y_i = \sum_{j=1}^n a_{i,j} v_j$$

Why *Scaled* Dot Product?

- **Softmax is sensitive to scale**

If $[x_1, x_2] = [0.1, 0.5]$, $\alpha = 10$

$$\text{softmax}([x_1, x_2]) = \left[\frac{e^{x_1}}{e^{x_1} + e^{x_2}}, \frac{e^{x_2}}{e^{x_1} + e^{x_2}} \right]$$

[0.4013, 0.5987]

$$\text{softmax}([\alpha x_1, \alpha x_2]) = \left[\frac{e^{\alpha x_1}}{e^{\alpha x_1} + e^{\alpha x_2}}, \frac{e^{\alpha x_2}}{e^{\alpha x_1} + e^{\alpha x_2}} \right]$$

[0.0180, 0.9820]

Self-attention in equations

- A sequence of input vectors $x_1, \dots, x_n \in \mathbb{R}^d$
- First, construct a set of **queries**, **keys** and **values**:

$$q_i = W_Q x_i, k_i = W_K x_i, v_i = W_V x_i$$


- Second, for each q_i , compute attention scores and attention distributions:

$$a_{i,j} = \text{softmax}\left(\frac{q_i^T k_j}{\sqrt{d}}\right) \quad \text{aka. "scaled dot product"}$$


- Finally, compute the weighted sum:

$$y_i = \sum_{j=1}^n a_{i,j} v_j$$

Self-attention: Illustration



Self-attention: Illustration



<http://jalammar.github.io/illustrated-transformer/>

Self-attention: matrix notations



Self-attention: matrix notations

$$\text{attn}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right)V$$

The diagram illustrates the self-attention mechanism using matrix notations and dimensions. It shows the calculation of the attention matrix Z as follows:

$$\text{softmax}\left(\frac{Q \times K^T}{\sqrt{d_k}}\right) V = Z$$

Where:

- Q (Query matrix) is a 2x3 matrix (pink).
- K^T (Key matrix, transposed) is a 3x2 matrix (orange).
- V (Value matrix) is a 2x3 matrix (blue).
- Z (Attention matrix) is a 2x3 matrix (pink).
- d_k is the dimension of the key matrix.

Self-attention: matrix notations



hardmaru
@hardmaru

The most important formula in deep learning after 2018

Self-Attention

What is self-attention? Self-attention calculates a weighted average of feature representations with the weight proportional to a similarity score between pairs of representations. Formally, an input sequence of n tokens of dimensions d , $X \in \mathbf{R}^{n \times d}$, is projected using three matrices $W_Q \in \mathbf{R}^{d \times d_q}$, $W_K \in \mathbf{R}^{d \times d_k}$, and $W_V \in \mathbf{R}^{d \times d_v}$ to extract feature representations Q , K , and V , referred to as query, key, and value respectively with $d_k = d_q$. The outputs Q , K , V are computed as

$$Q = XW_Q, \quad K = XW_K, \quad V = XW_V. \quad (1)$$

So, self-attention can be written as,

$$S = D(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_q}} \right) V, \quad (2)$$

where softmax denotes a *row-wise* softmax normalization function. Thus, each element in S depends on all other elements in the same row.

9:08 PM · Feb 9, 2021 · Twitter Web App

580 Retweets 38 Quote Tweets 3,407 Likes

Attention is *General*

- Given a set of **key** and **value** vectors, and a **query** vector, attention is a technique to compute a weighted sum of the **value** vectors, dependent on the **query and keys**
 - We sometimes say that the **query** attends to the **values** via **keys**
 - In the NMT vase, each decoder hidden state (**query**) attends to all the encoder hidden states (**keys and values**)
- **Intuition**
 - The weighted sum is a *selective summary* of the information contained in the **values**, where the **query** and **keys** determines which **values** to focus on
 - Attention is a way to obtain a *fixed-size representation* of an arbitrary set of representations (the **values**), dependent on some other representation (the **query**)

$$\text{attn}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right)V$$

Multi-head Attention


- Problem with self-attention?

$$y_i = \sum_{j=1}^n a_{i,j} v_j$$

one set of attention weights a_i

- It is better to use multiple attention weights instead of one!
 - Each attention can focus on different positions
- How to do this? Splits queries, keys, values to multiple heads!

Multi-head Attention: Head Split

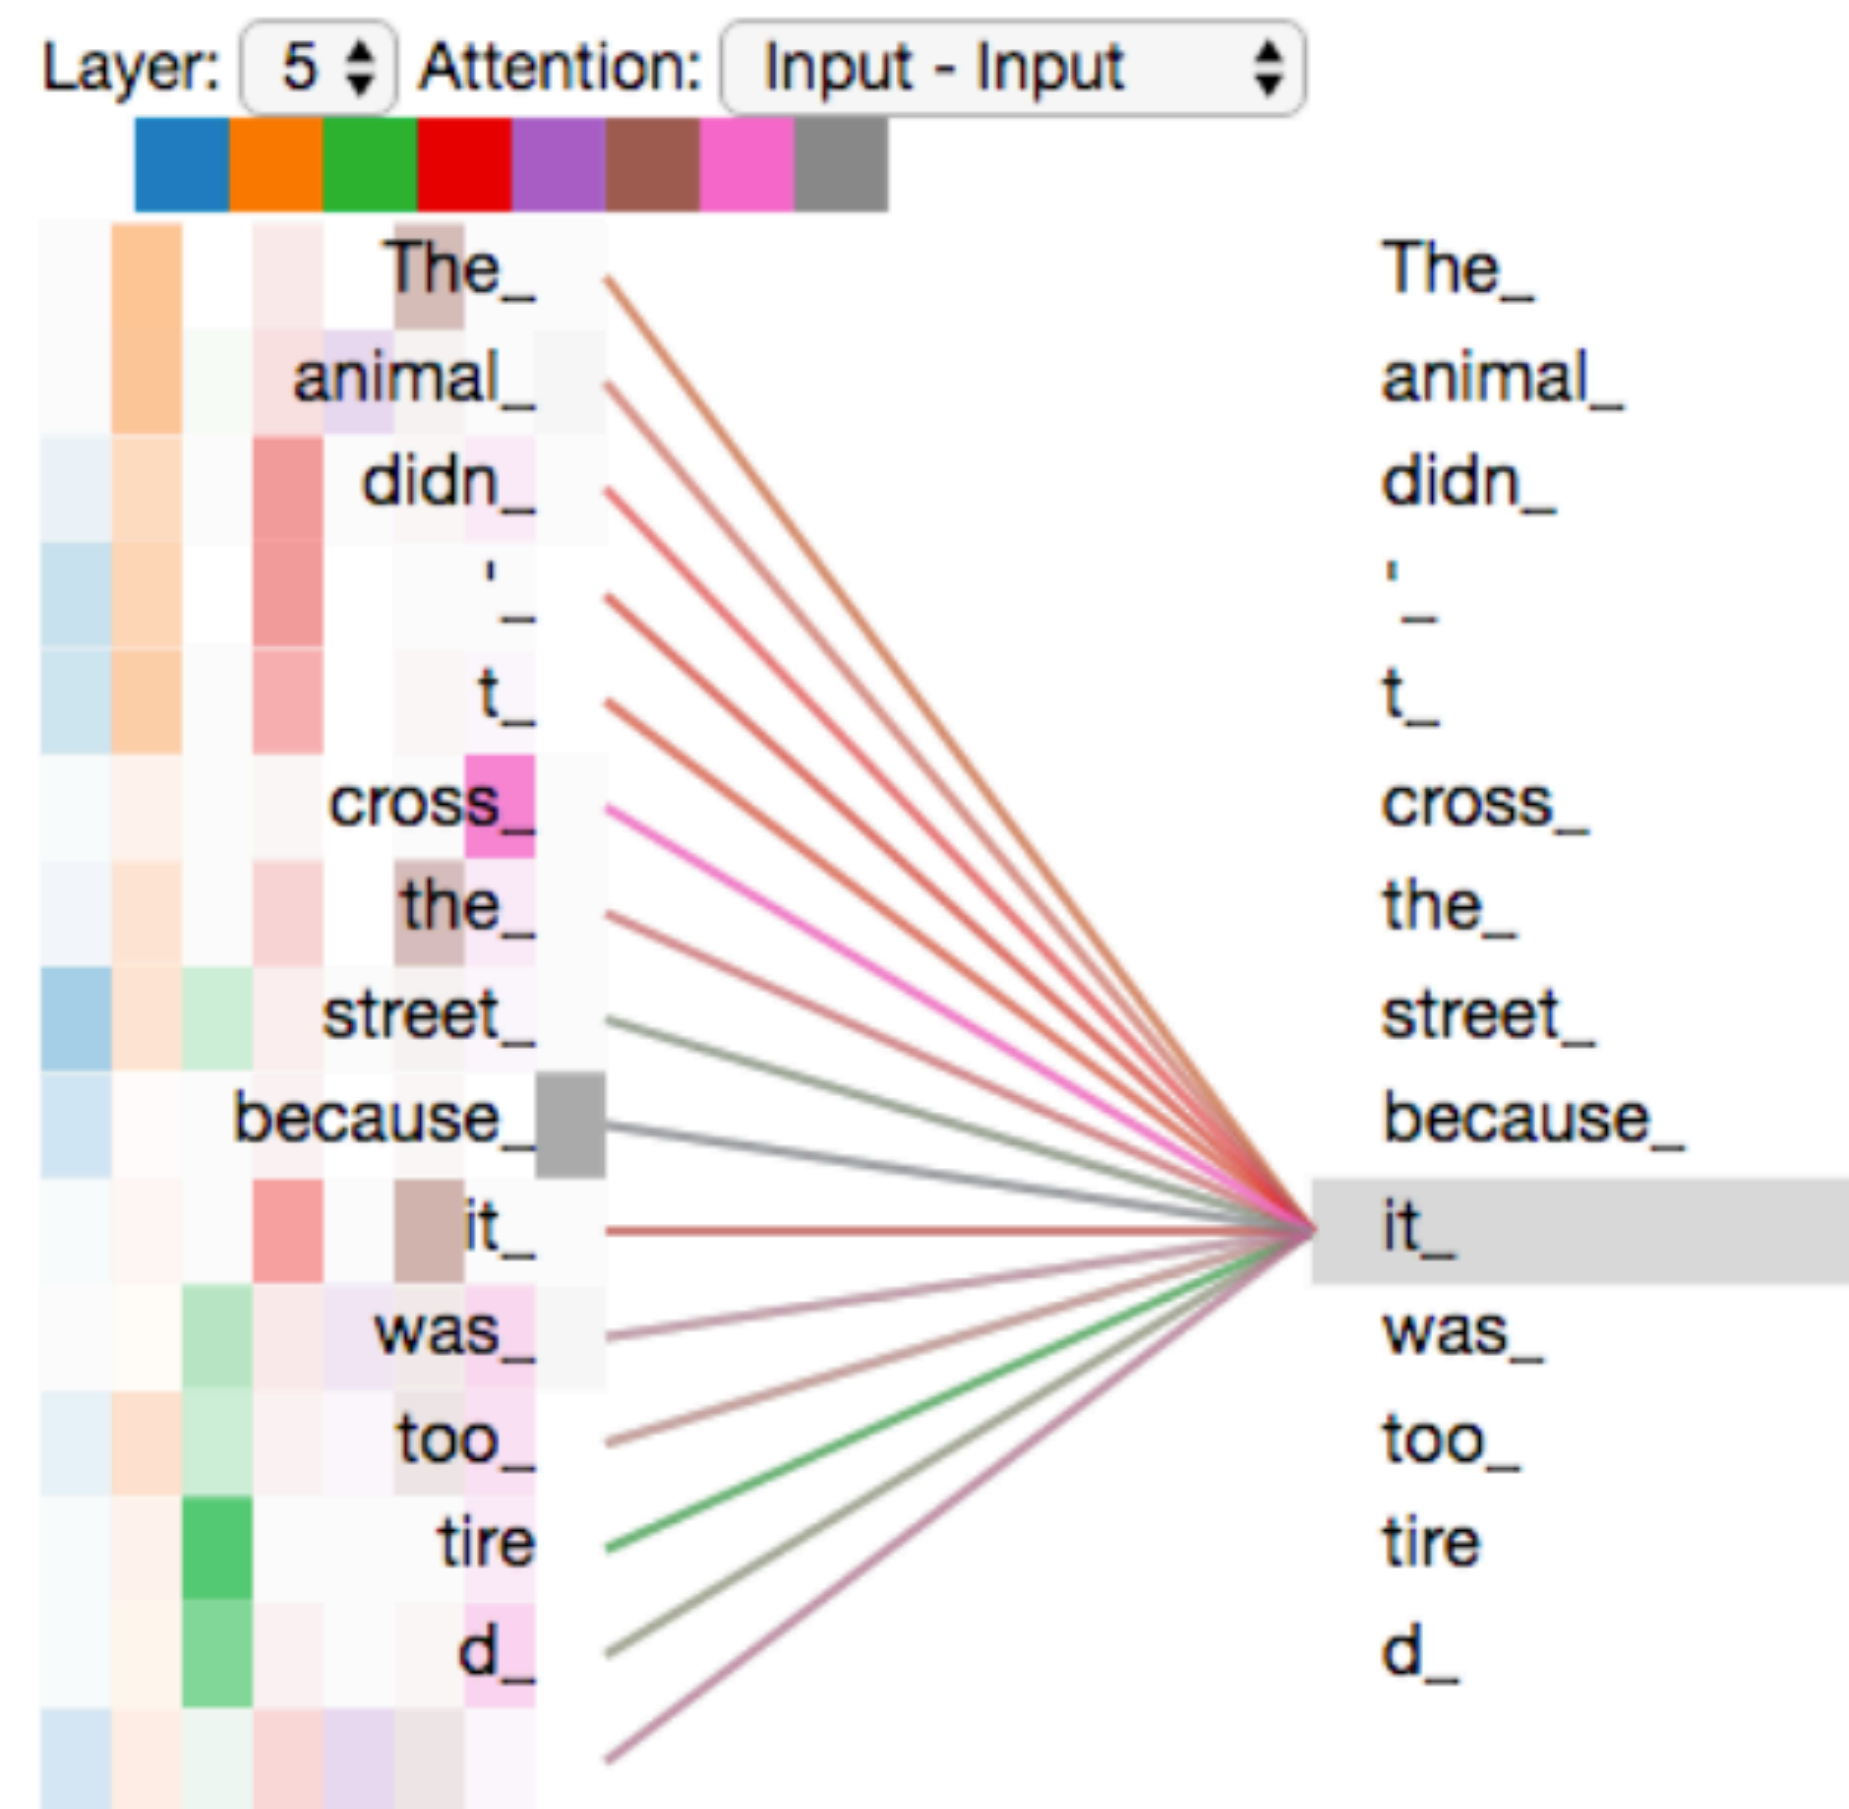
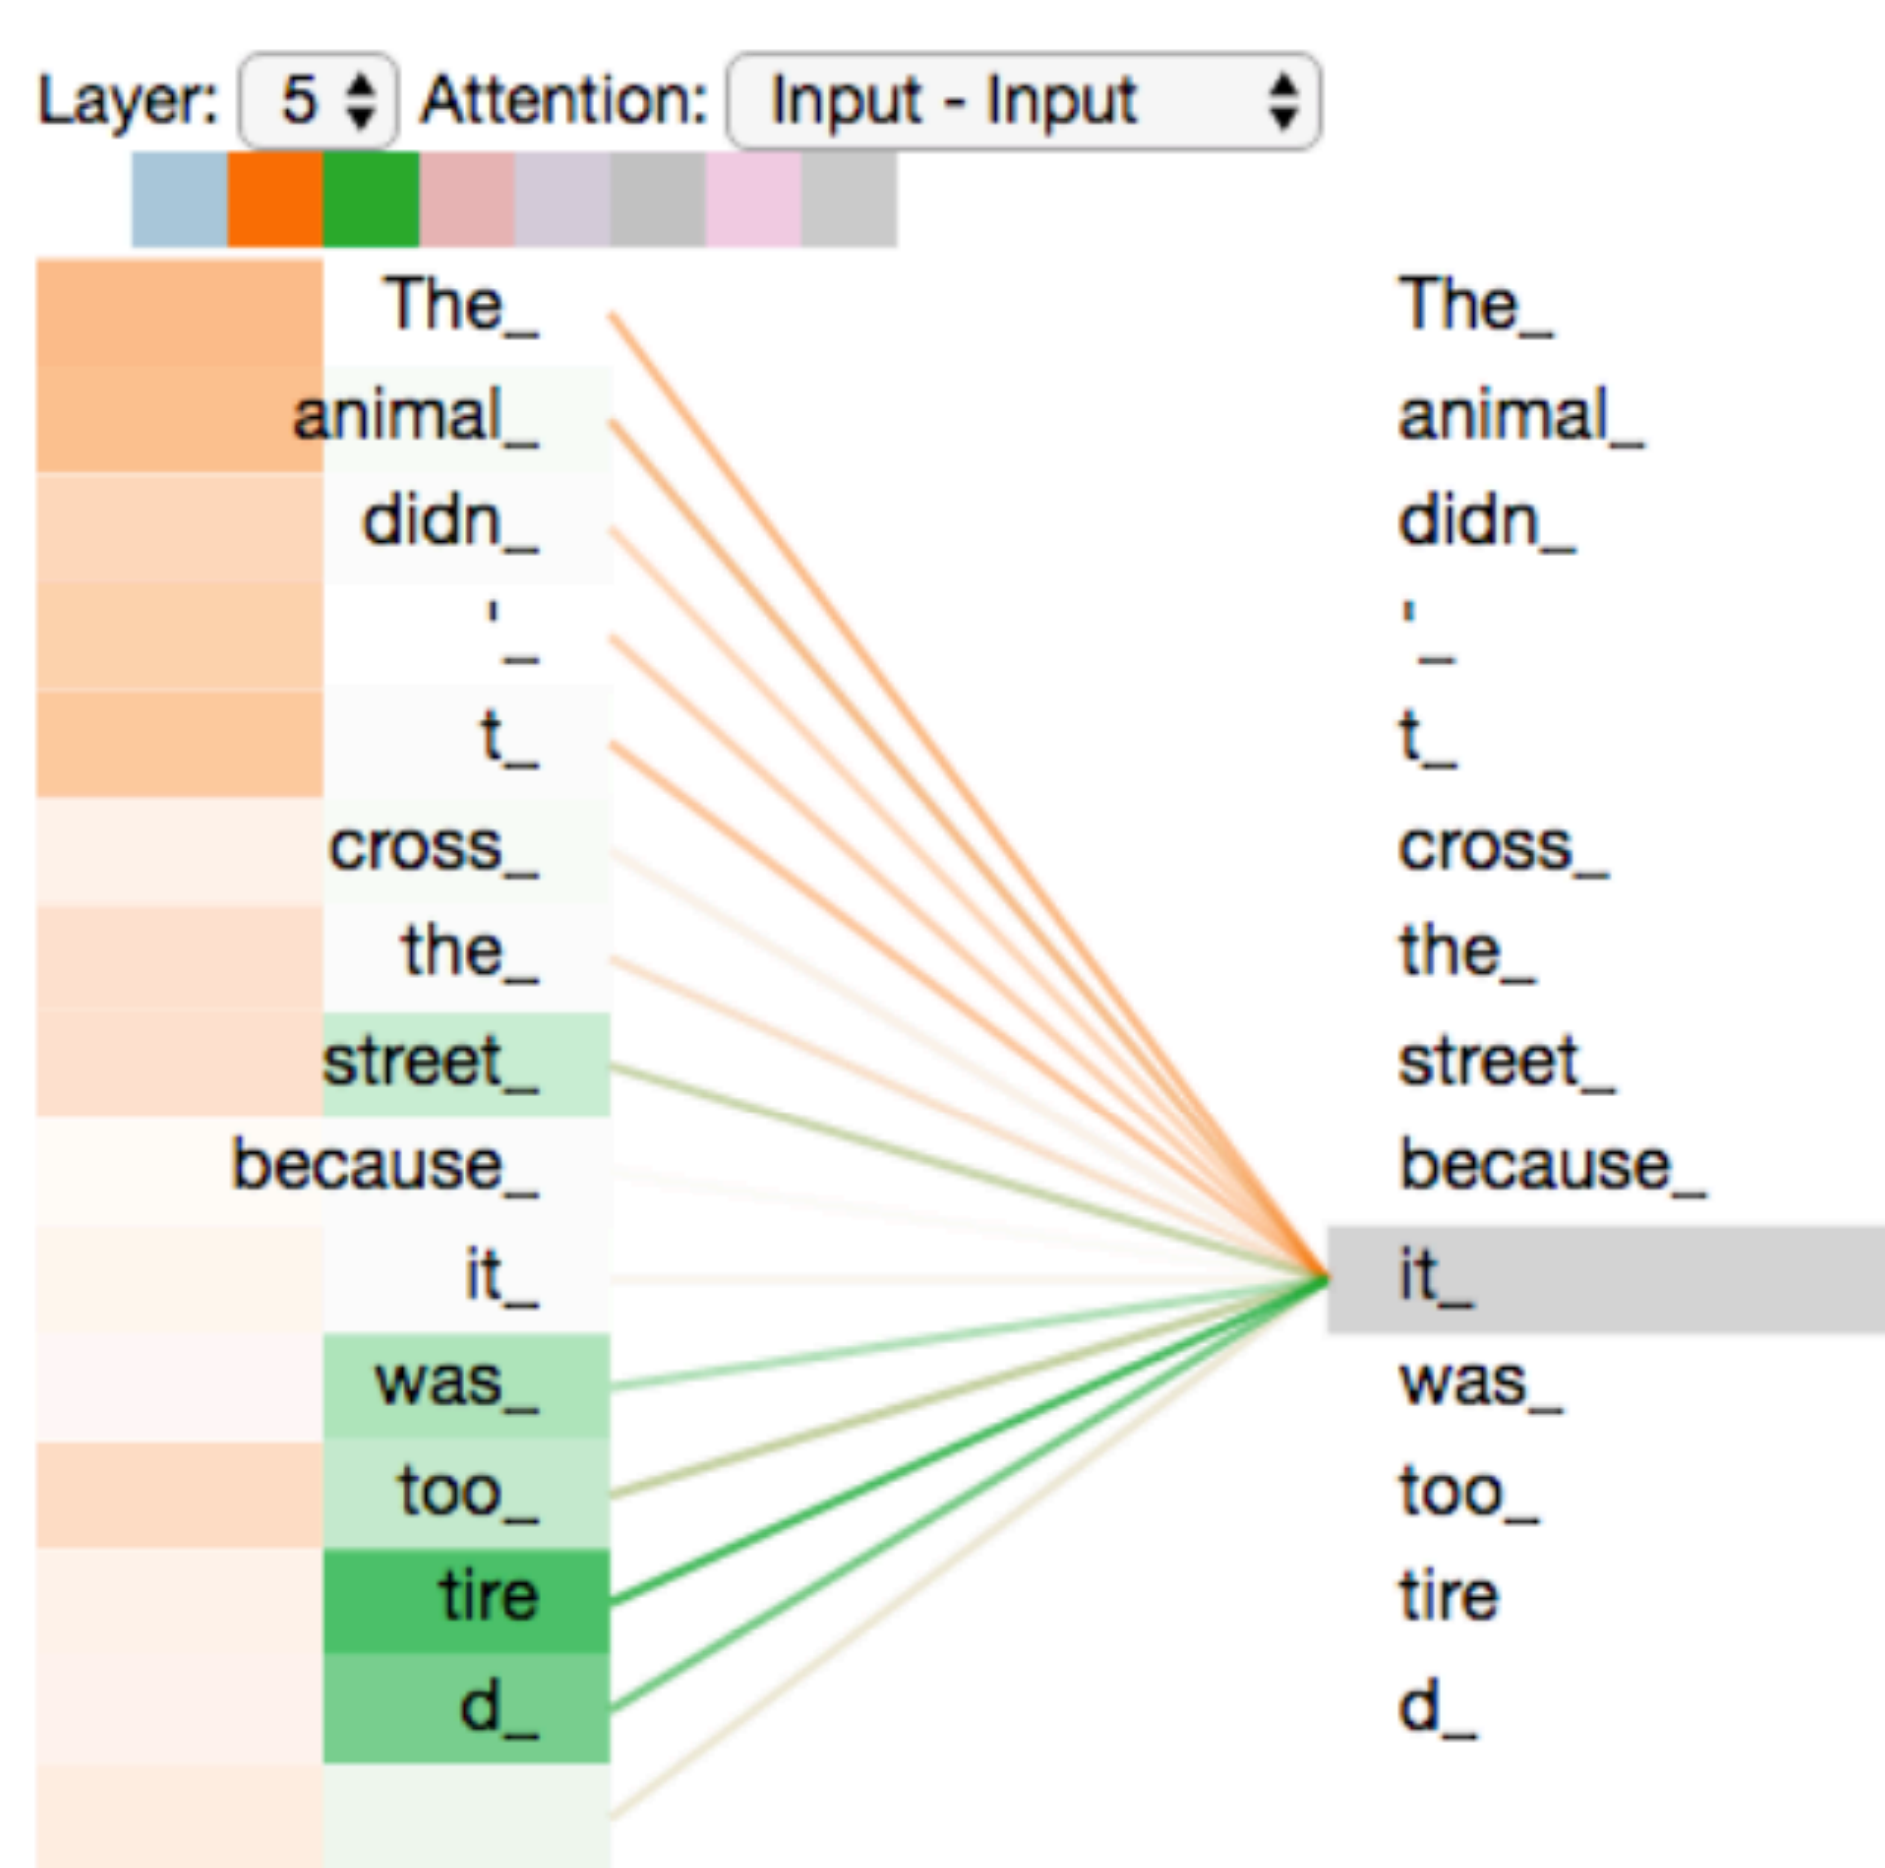


$$h_1 = \text{attn}(Q_1, K_1, V_1) = \text{softmax}\left(\frac{Q_1 K_1^T}{\sqrt{d/2}}\right) V_1$$

$$h_2 = \text{attn}(Q_2, K_2, V_2) = \text{softmax}\left(\frac{Q_2 K_2^T}{\sqrt{d/2}}\right) V_2$$

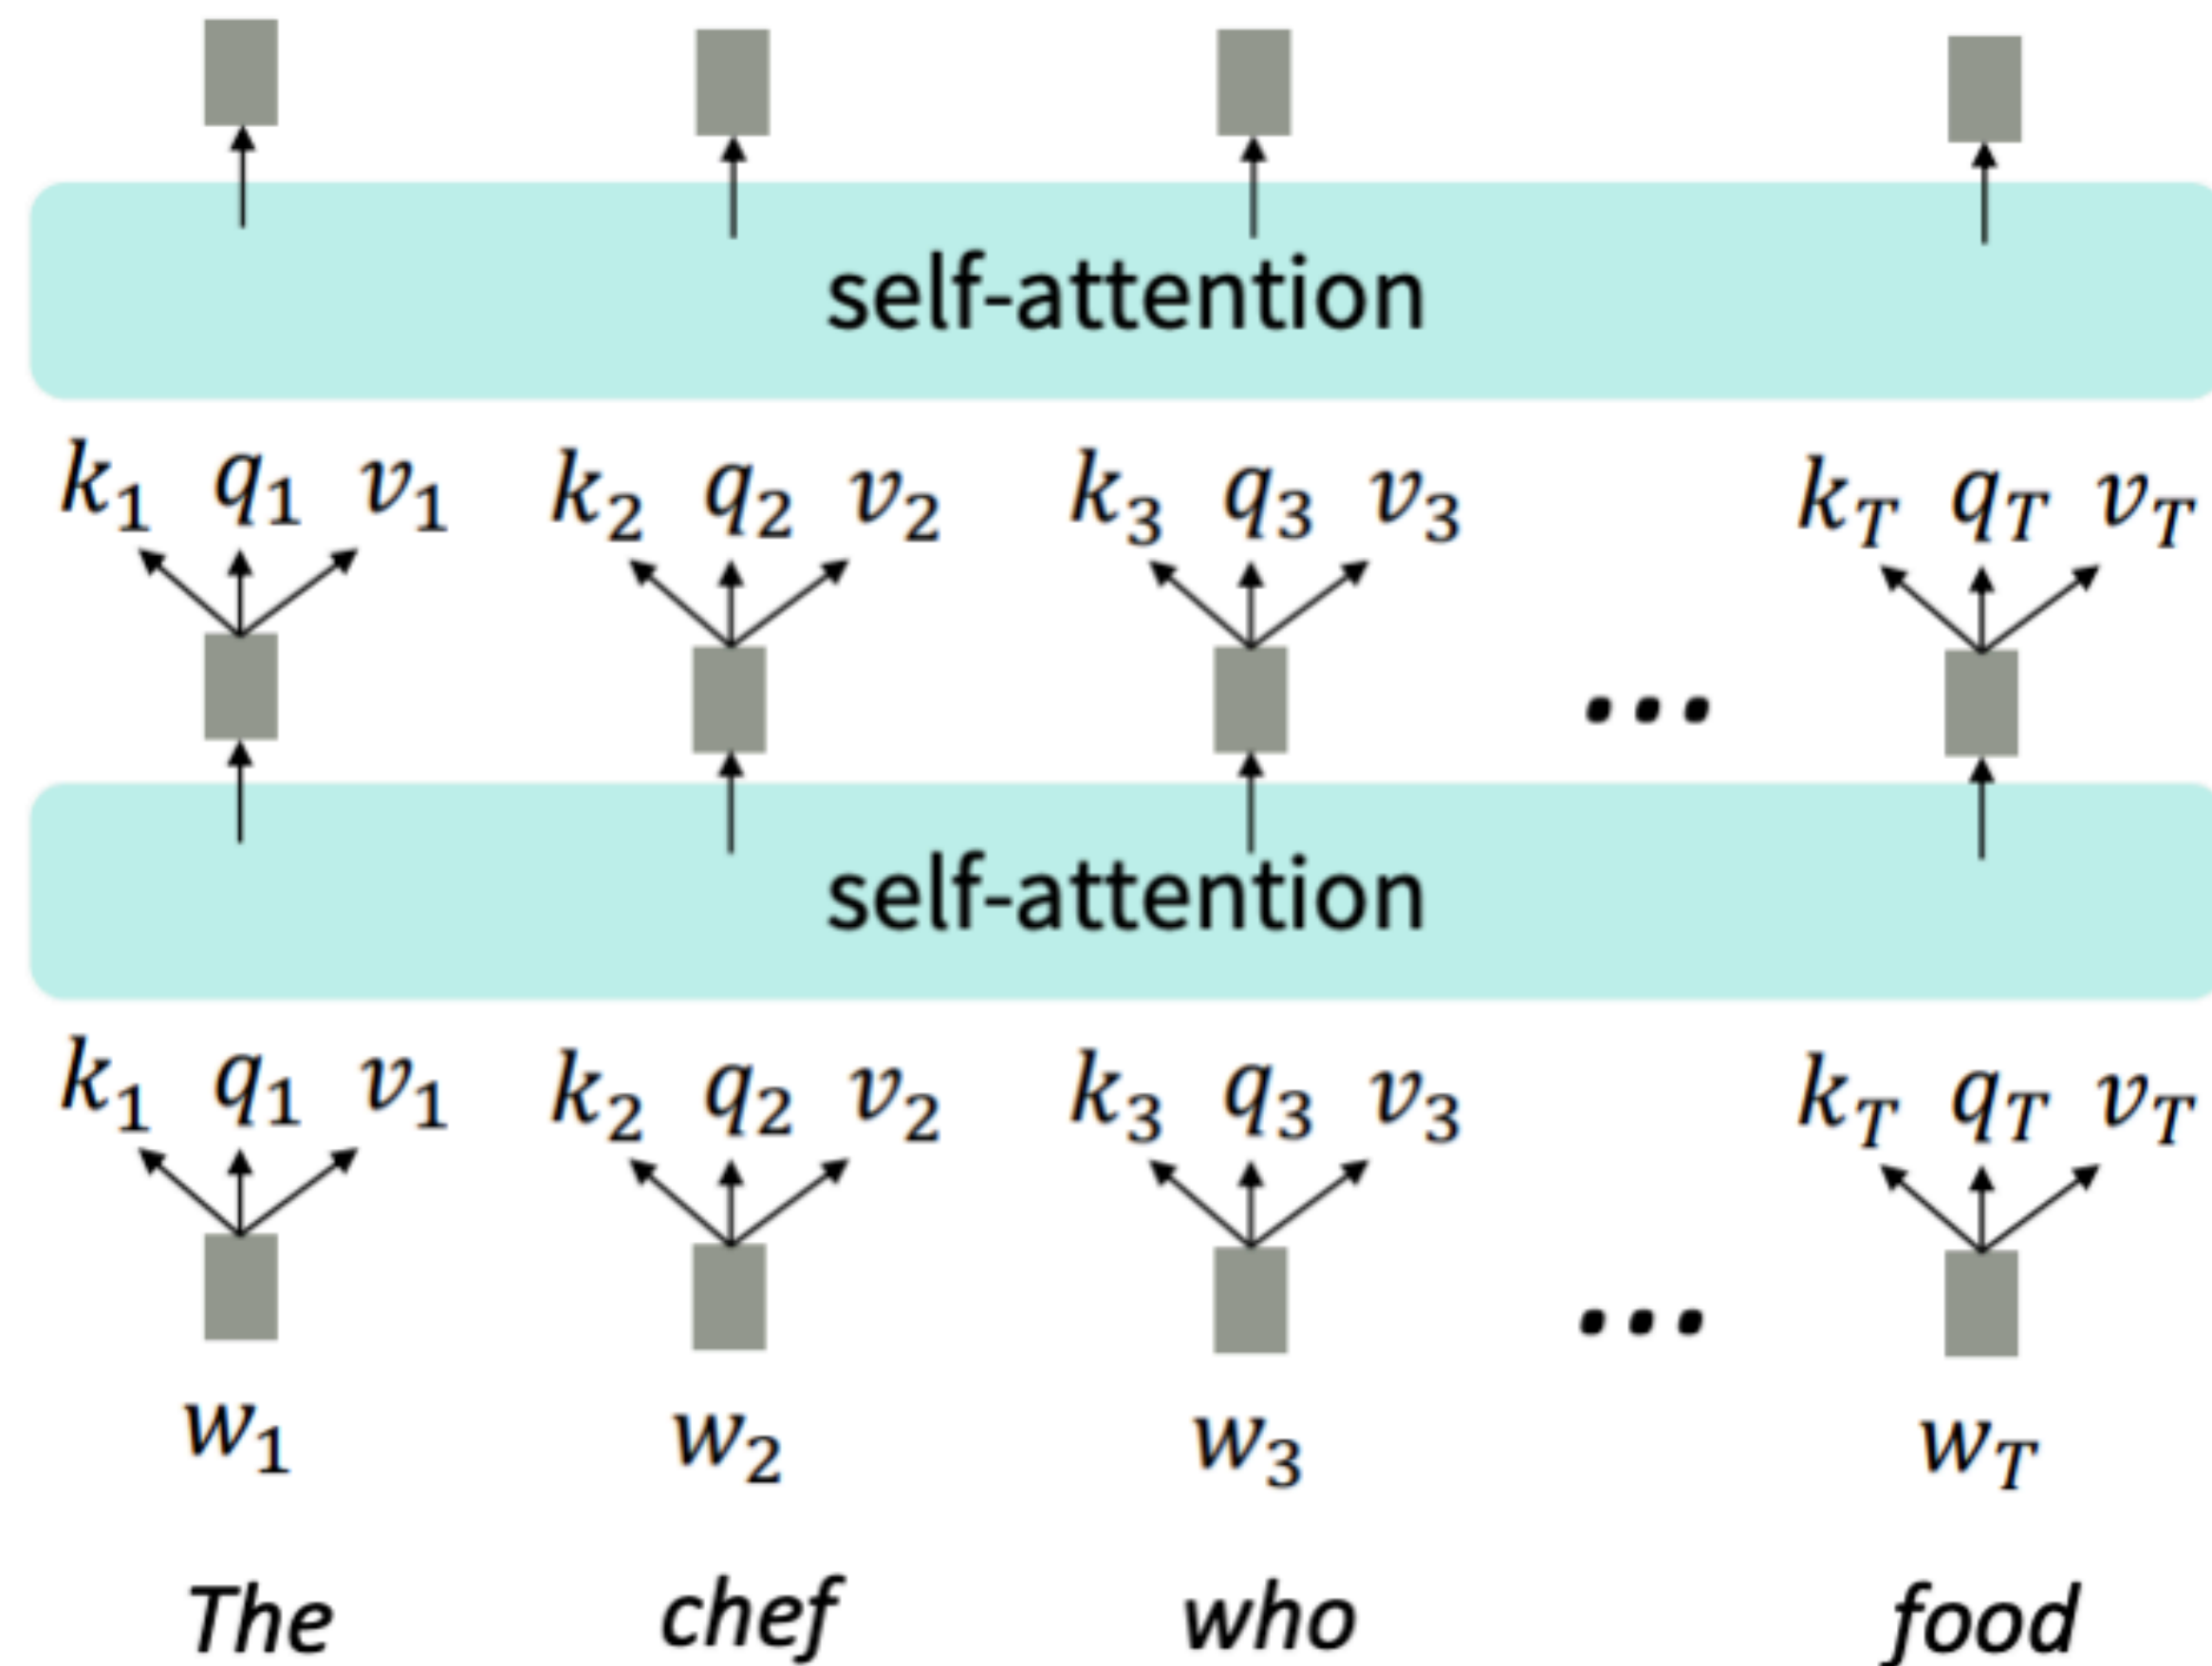
$$Y = \text{concat}(h_1, h_2) W_O$$

What does multi-head attention learn?



Transformer Encoder

- Replacing RNNs with multi-head self-attention




$$\text{MultiHead}(X) = \text{concat}(h_1, \dots, h_k)W_O$$
$$h_i = \text{attn}(Q_i, K_i, V_i)$$
$$Q_i = (XW_Q)^i, K_i = (XW_K)^i, V_i = (XW_V)^i$$

$$\text{attn}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right)V$$


Self-attention does not know the order of the inputs!

Missing Piece: Positional Information

RNN



Self-attention




Missing Piece: Positional Information

- Unlike RNNs, self-attention does **not** build in order information
 - Encode the order of the sentence into the input x_1, \dots, x_n
- Solution: add **positional encoding** to the input embeddings

$$x_i \leftarrow x_i + p_i$$

- Use sine and cosine functions of different frequencies (not learnable)

$$p_i = \begin{pmatrix} \sin(i/10000^{2*1/d}) \\ \cos(i/10000^{2*1/d}) \\ \vdots \\ \sin(i/10000^{2*\frac{d}{2}/d}) \\ \cos(i/10000^{2*\frac{d}{2}/d}) \end{pmatrix}$$



Adding Nonlinearities

- There is no elementwise nonlinearities in self-attention; stacking more self-attention layers just re-averages value vectors
- Simple fix: add a feed-forward network to post-process each output vector


$$\text{FFN}(\mathbf{x}_i) = W_2 \text{ReLU}(W_1 \mathbf{x}_i + \mathbf{b}_1) + \mathbf{b}_2$$

A large number
of parameters

$$W_1 \in \mathbb{R}^{d_{ff} \times d}, \mathbf{b}_1 \in \mathbb{R}^{d_{ff}}$$

$$W_2 \in \mathbb{R}^{d \times d_{ff}}, \mathbf{b}_2 \in \mathbb{R}^d$$

In practice, they use $d_{ff} = 4d$



Feed-Forward Layers

- Feed-forward layers constitute **two-thirds** of parameters
- Operates as memories of textual patterns (Gova et al., 2021)

Key	Pattern	Example trigger prefixes
k_{449}^1	Ends with “ <i>substitutes</i> ” (shallow)	<i>At the meeting, Elton said that “for artistic reasons there could be no substitutes</i> <i>In German service, they were used as substitutes</i> <i>Two weeks later, he came off the substitutes</i>
k_{2546}^6	Military, ends with “ <i>base</i> ”/“ <i>bases</i> ” (shallow + semantic)	<i>On 1 April the SRSg authorised the SADF to leave their bases</i> <i>Aircraft from all four carriers attacked the Australian base</i> <i>Bombers flying missions to Rabaul and other Japanese bases</i>
k_{2997}^{10}	a “part of” relation (semantic)	<i>In June 2012 she was named as one of the team that competed</i> <i>He was also a part of the Indian delegation</i> <i>Toy Story is also among the top ten in the BFI list of the 50 films you should</i>
k_{2989}^{13}	Ends with a time range (semantic)	<i>Worldwide, most tornadoes occur in the late afternoon, between 3 pm and 7</i> <i>Weekend tolls are in effect from 7:00 pm Friday until</i> <i>The building is open to the public seven days a week, from 11:00 am to</i>
k_{1935}^{16}	TV shows (semantic)	<i>Time shifting viewing added 57 percent to the episode’s</i> <i>The first season set that the episode was included in was as part of the</i> <i>From the original NBC daytime version , archived</i>

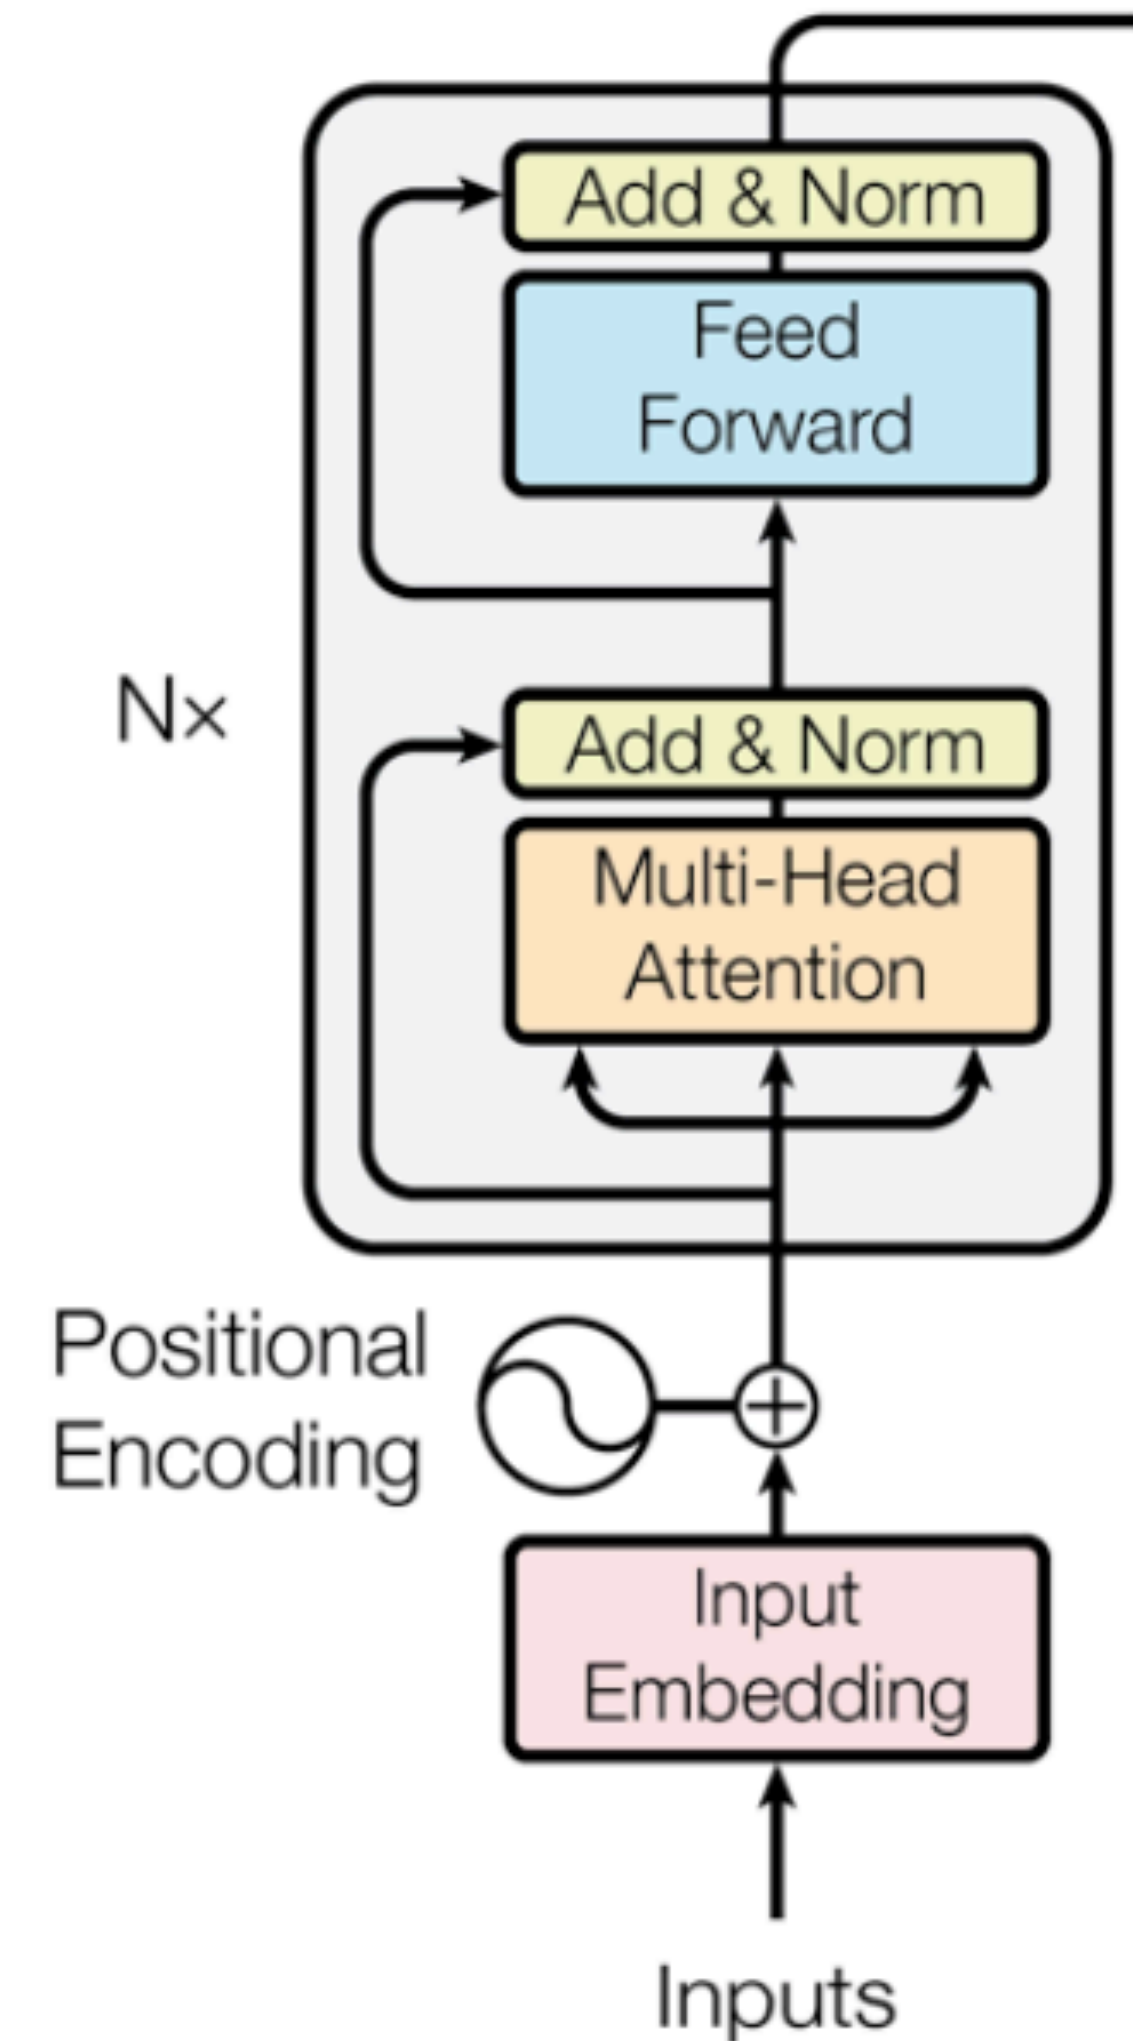
Transformer Encoder

- Each encoder layer has two sub-layers:
 - A multi-head self-attention layer
 - A feedforward layer
- Add & Norm:
 - Add: Residual connection (He et al., 2016)

$$Y \leftarrow Y + X$$

- Norm: Layer normalization (Ba et al., 2016)

$$Y = \frac{X - E[X]}{\sqrt{\text{Var}[X] + \epsilon}} * \gamma + \beta$$



In (Vaswani et al., 2017), $N = 6$

Question

Which of the following statements is correct?

- (a) Transformers run faster than LSTMs
- (b) Transformers are easier to parallelize compared to LSTMs
- (c) Transformers have less parameters compared to LSTMs
- (d) Transformers are better at capturing positional information than LSTMs

Transformer: Pros and Cons

- **Easier to capture dependencies:** we draw attention between every pair of words!
- **Easier to parallelize:**

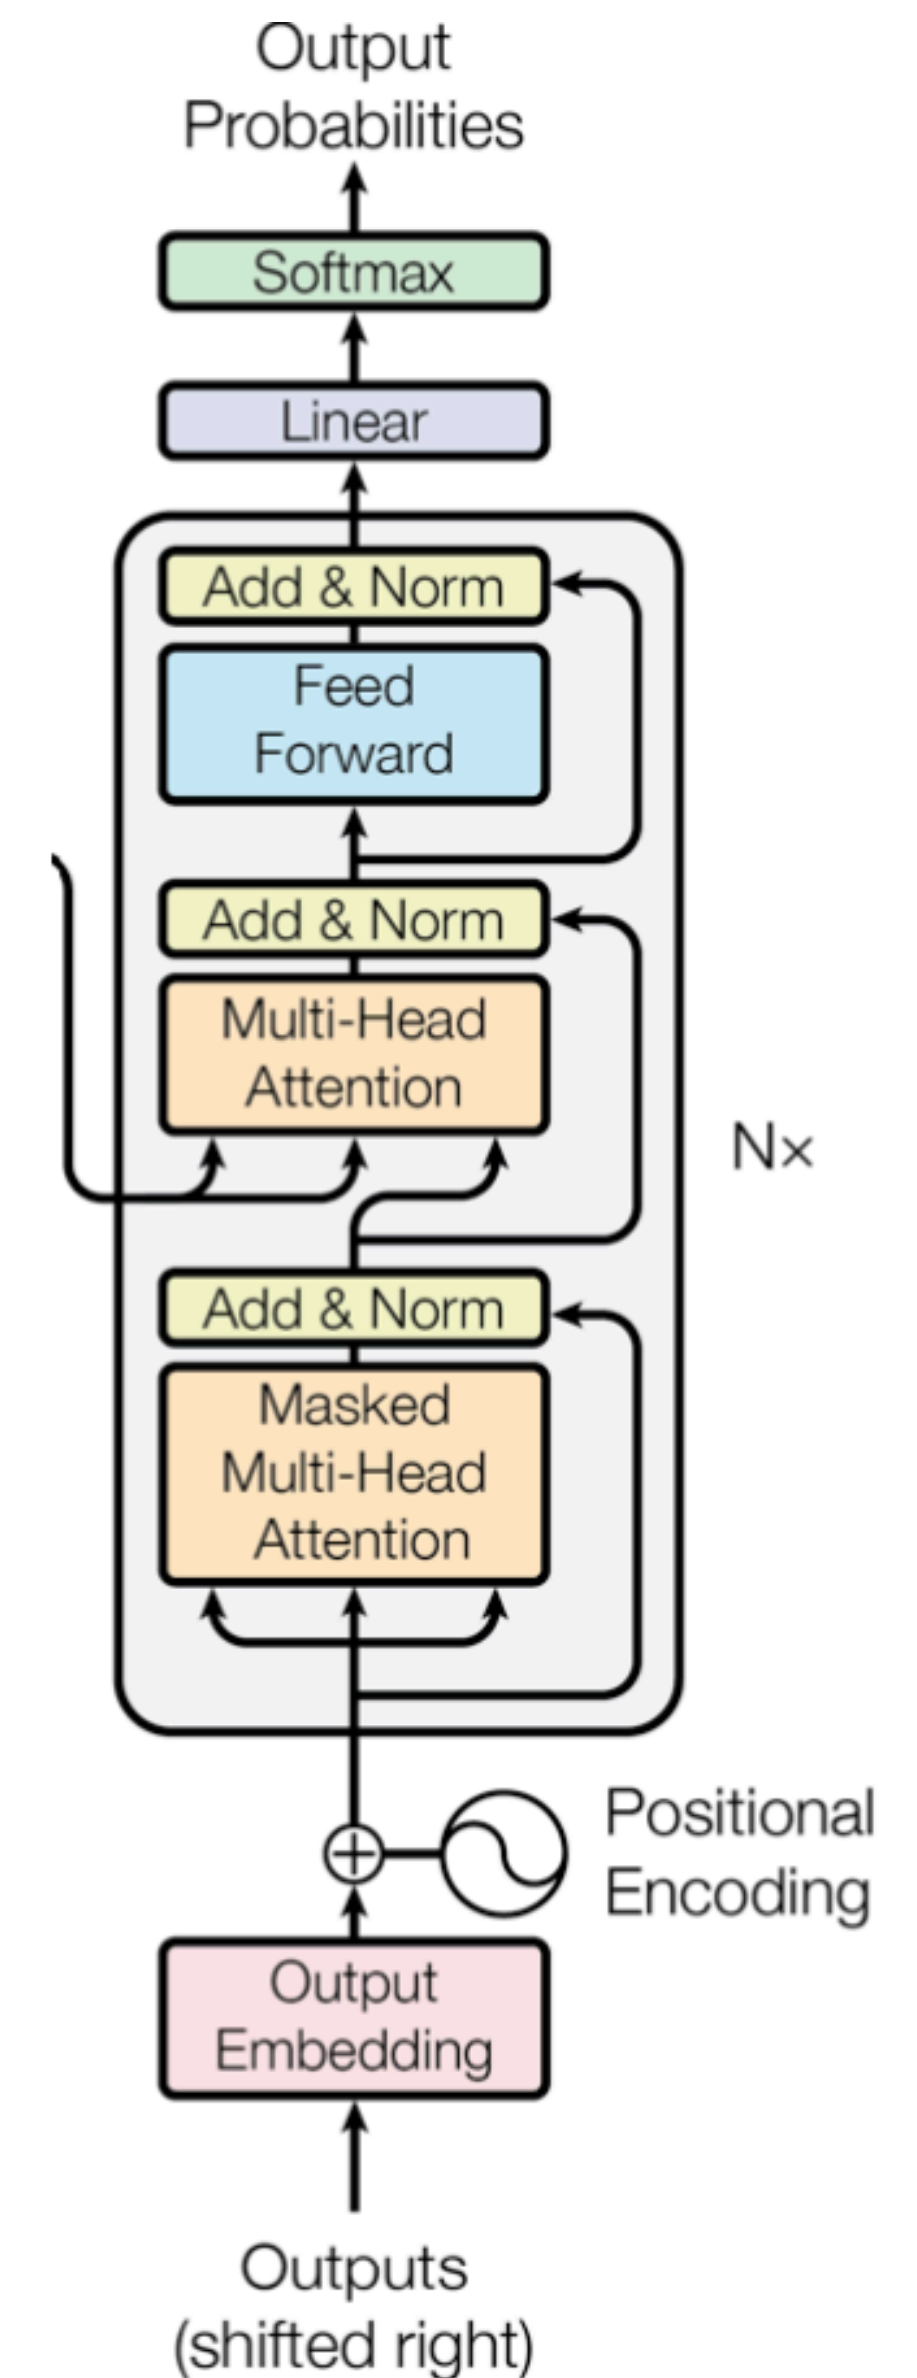
$$\begin{aligned}\text{MultiHead}(X) &= \text{concat}(h_1, \dots, h_k)W_O \\ h_i &= \text{attn}(Q_i, K_i, V_i) \\ Q_i &= (XW_Q)^i, K_i = (XW_K)^i, V_i = (XW_V)^i\end{aligned}$$

$$\text{attn}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right)V$$

- **Quadratic computation in self-attention:**
 - Can be very expensive when the sequence is very long
- **Harder to model positional information**

Transformer Decoder

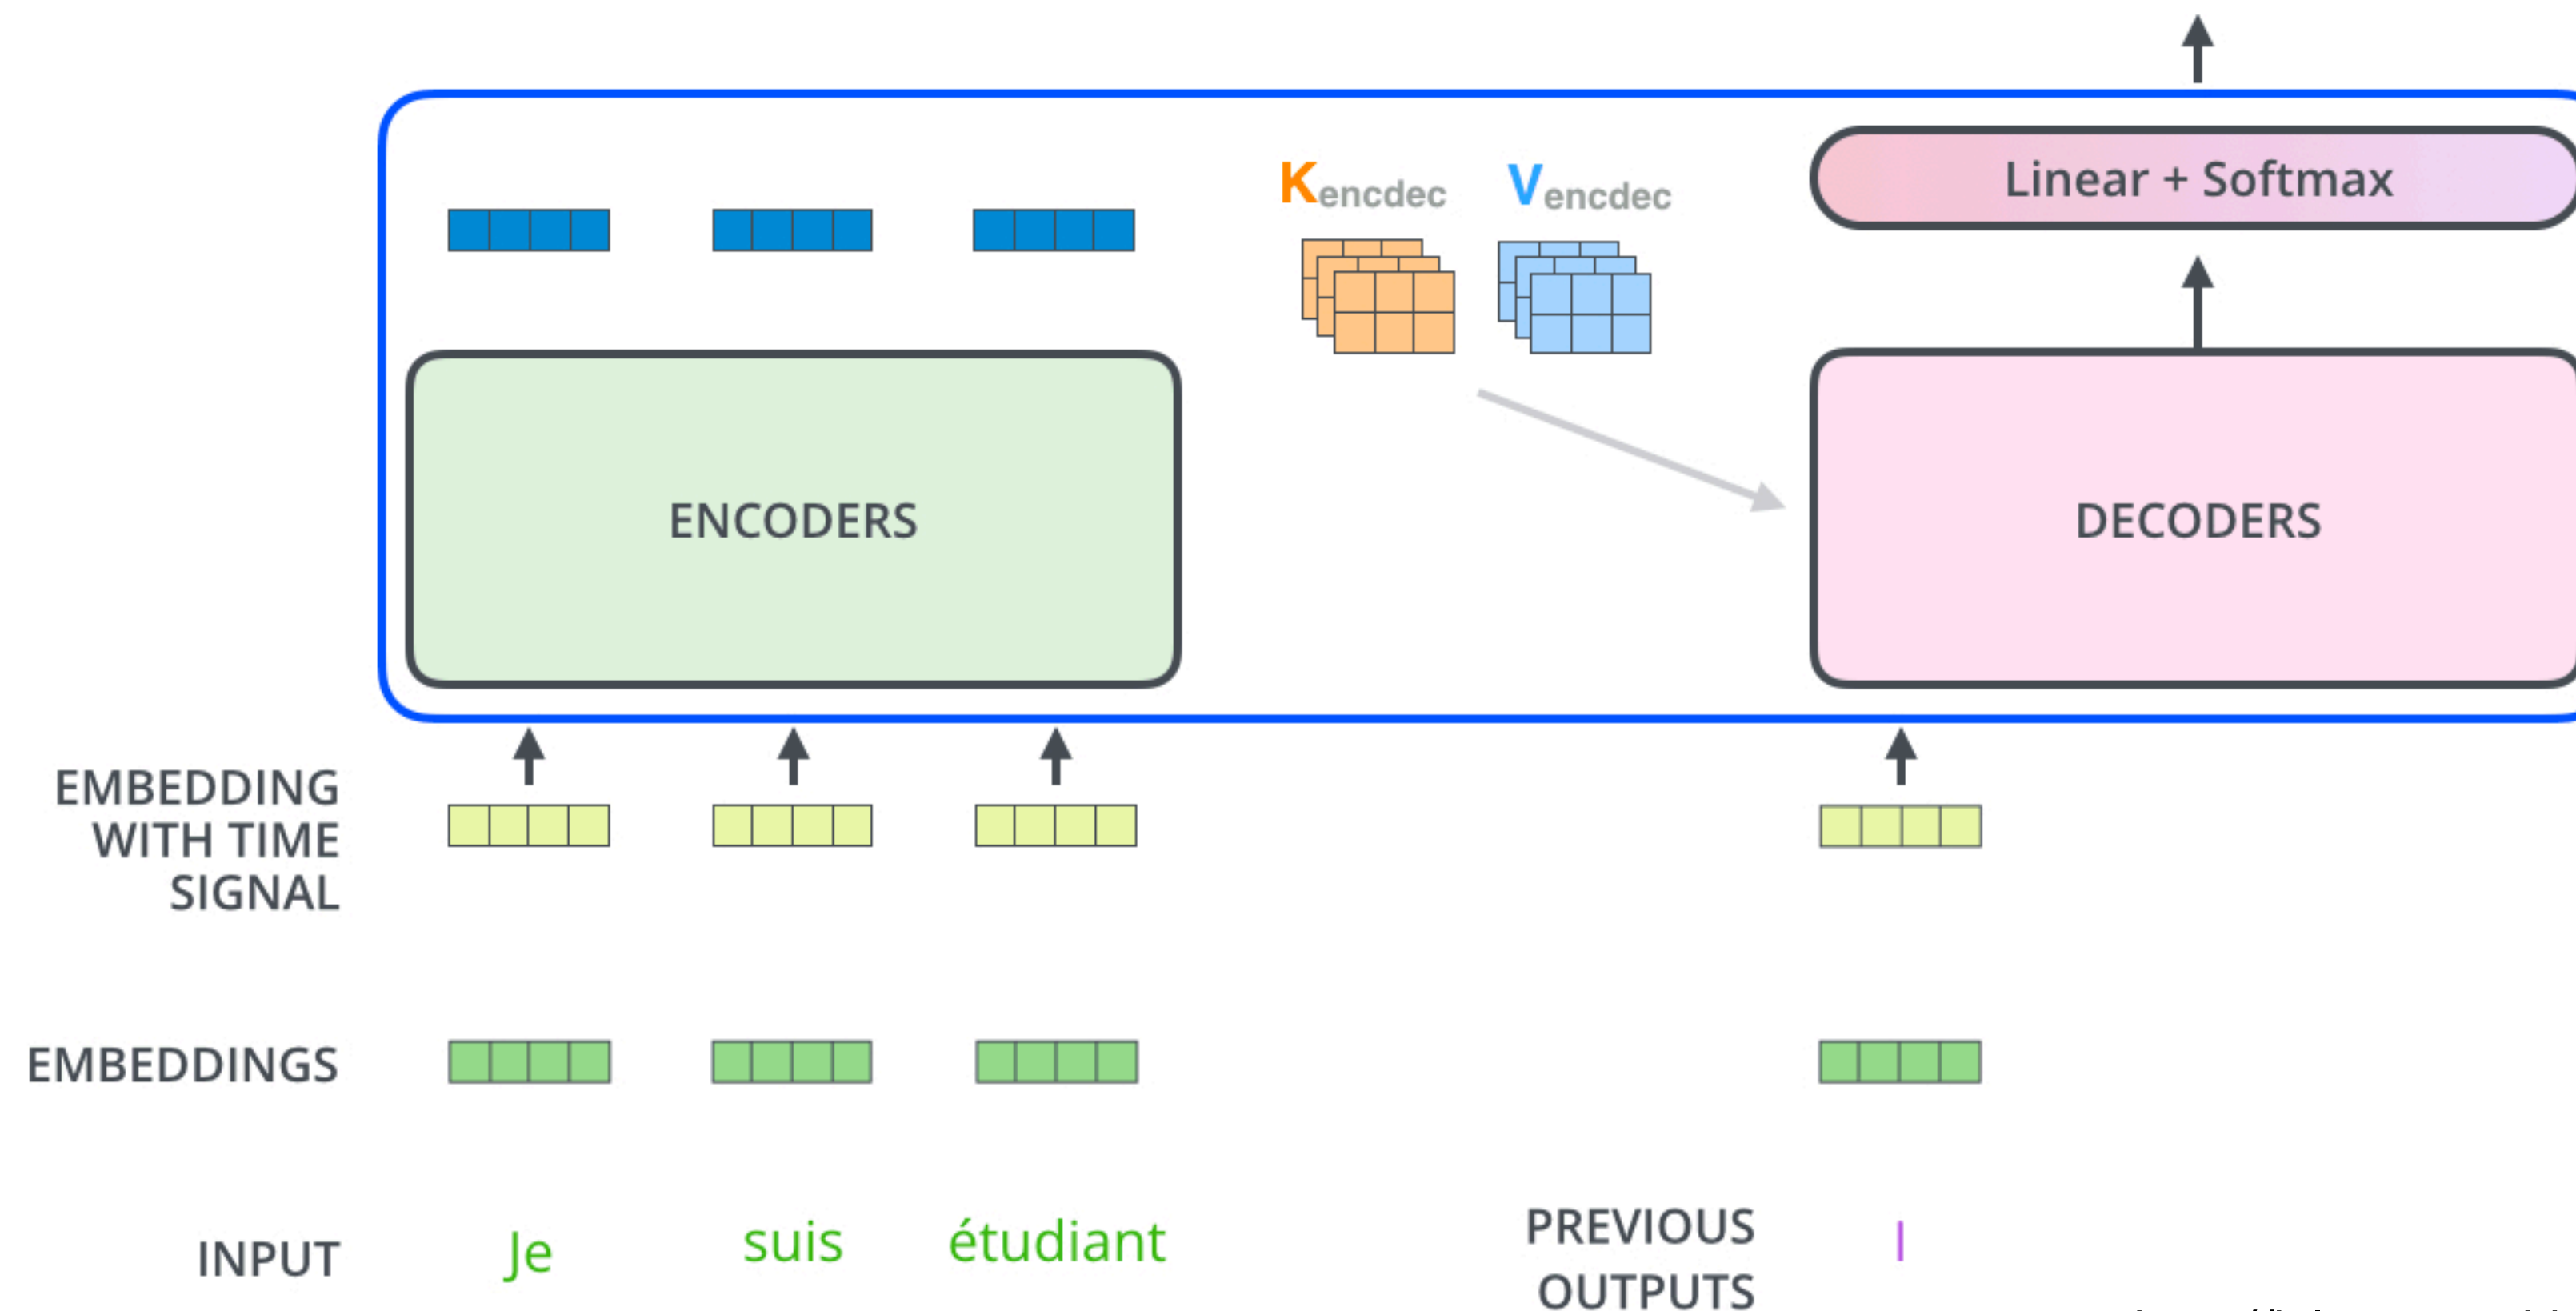
- Each decoder layer has three sub-layers:
 - A **masked multi-head attention** layer
 - A **multi-head cross attention** layer
 - A feedforward layer
- **Masked multi-head attention**
 - self-attention on the decoder states
- **Multi-head cross attention**
 - Decoder attends to encoder states
 - Encoder: **keys/values**
 - Decoder: **queries**



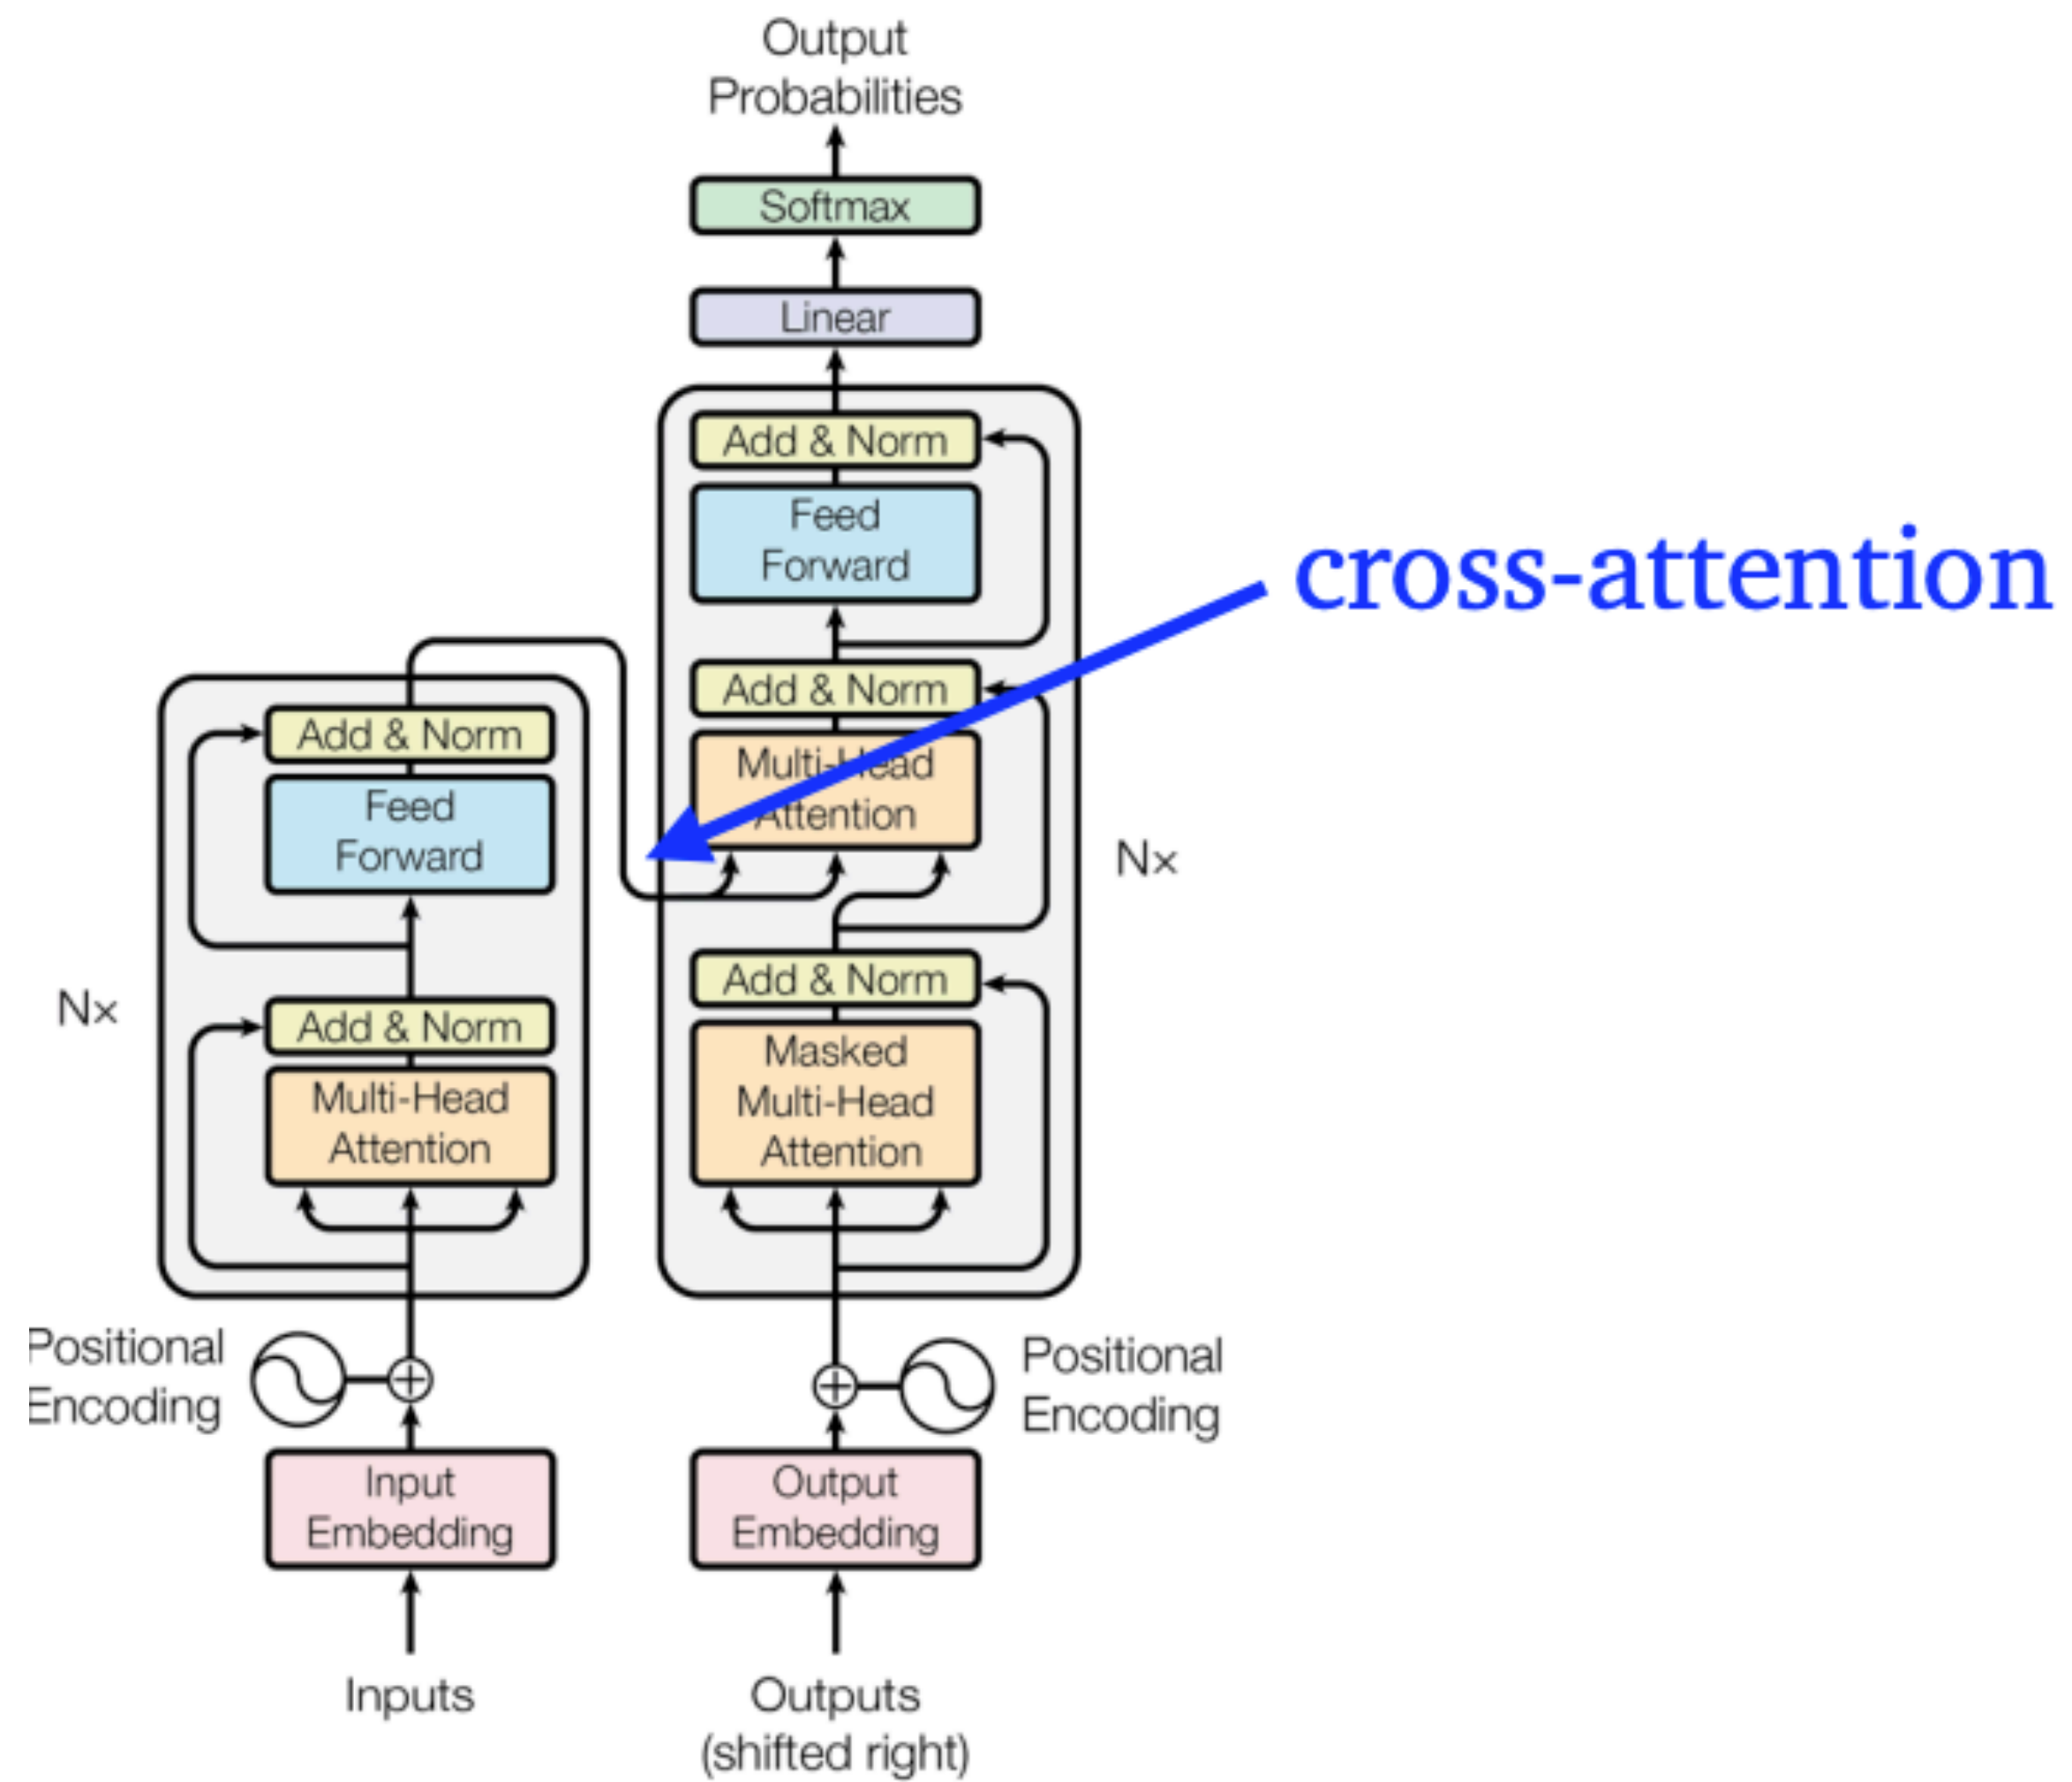
In (Vaswani et al., 2017), $N = 6$

Multi-head Cross Attention

- Decoder attends to encoder states
 - Encoder: **keys**/**values**
 - Decoder: **queries**



Multi-head Cross Attention




$$q_i = W_Q y_i, \quad k_i = W_K x_i, \quad v_i = W_V x_i$$

$$a_{i,j} = \text{softmax}\left(\frac{q_i^T k_j}{\sqrt{d}}\right)$$

$$y'_i = \sum_{j=1}^n a_{i,j} v_j$$


Masked Multi-Head Attention

- **Key point:** we cannot see the future words in decoder
- **Solution:** for every q_i , only attend to $\{(k_j, v_j)\}, j \leq i$




Masked Multi-Head Attention

- **Key point:** we cannot see the future words in decoder
- **Solution:** for every q_i , only attend to $\{(k_j, v_j)\}, j \leq i$



Masked Multi-Head Attention

- **Key point:** we cannot see the future words in decoder
- **Solution:** for every q_i , only attend to $\{(k_j, v_j)\}, j \leq i$




cannot be parallel!

Masked Multi-Head Attention

$$q_i = W_Q x_i, k_i = W_K x_i, v_i = W_V x_i$$

$$a_{i,j} = \text{softmax}\left(\frac{q_i^T k_j}{\sqrt{d}}\right)$$



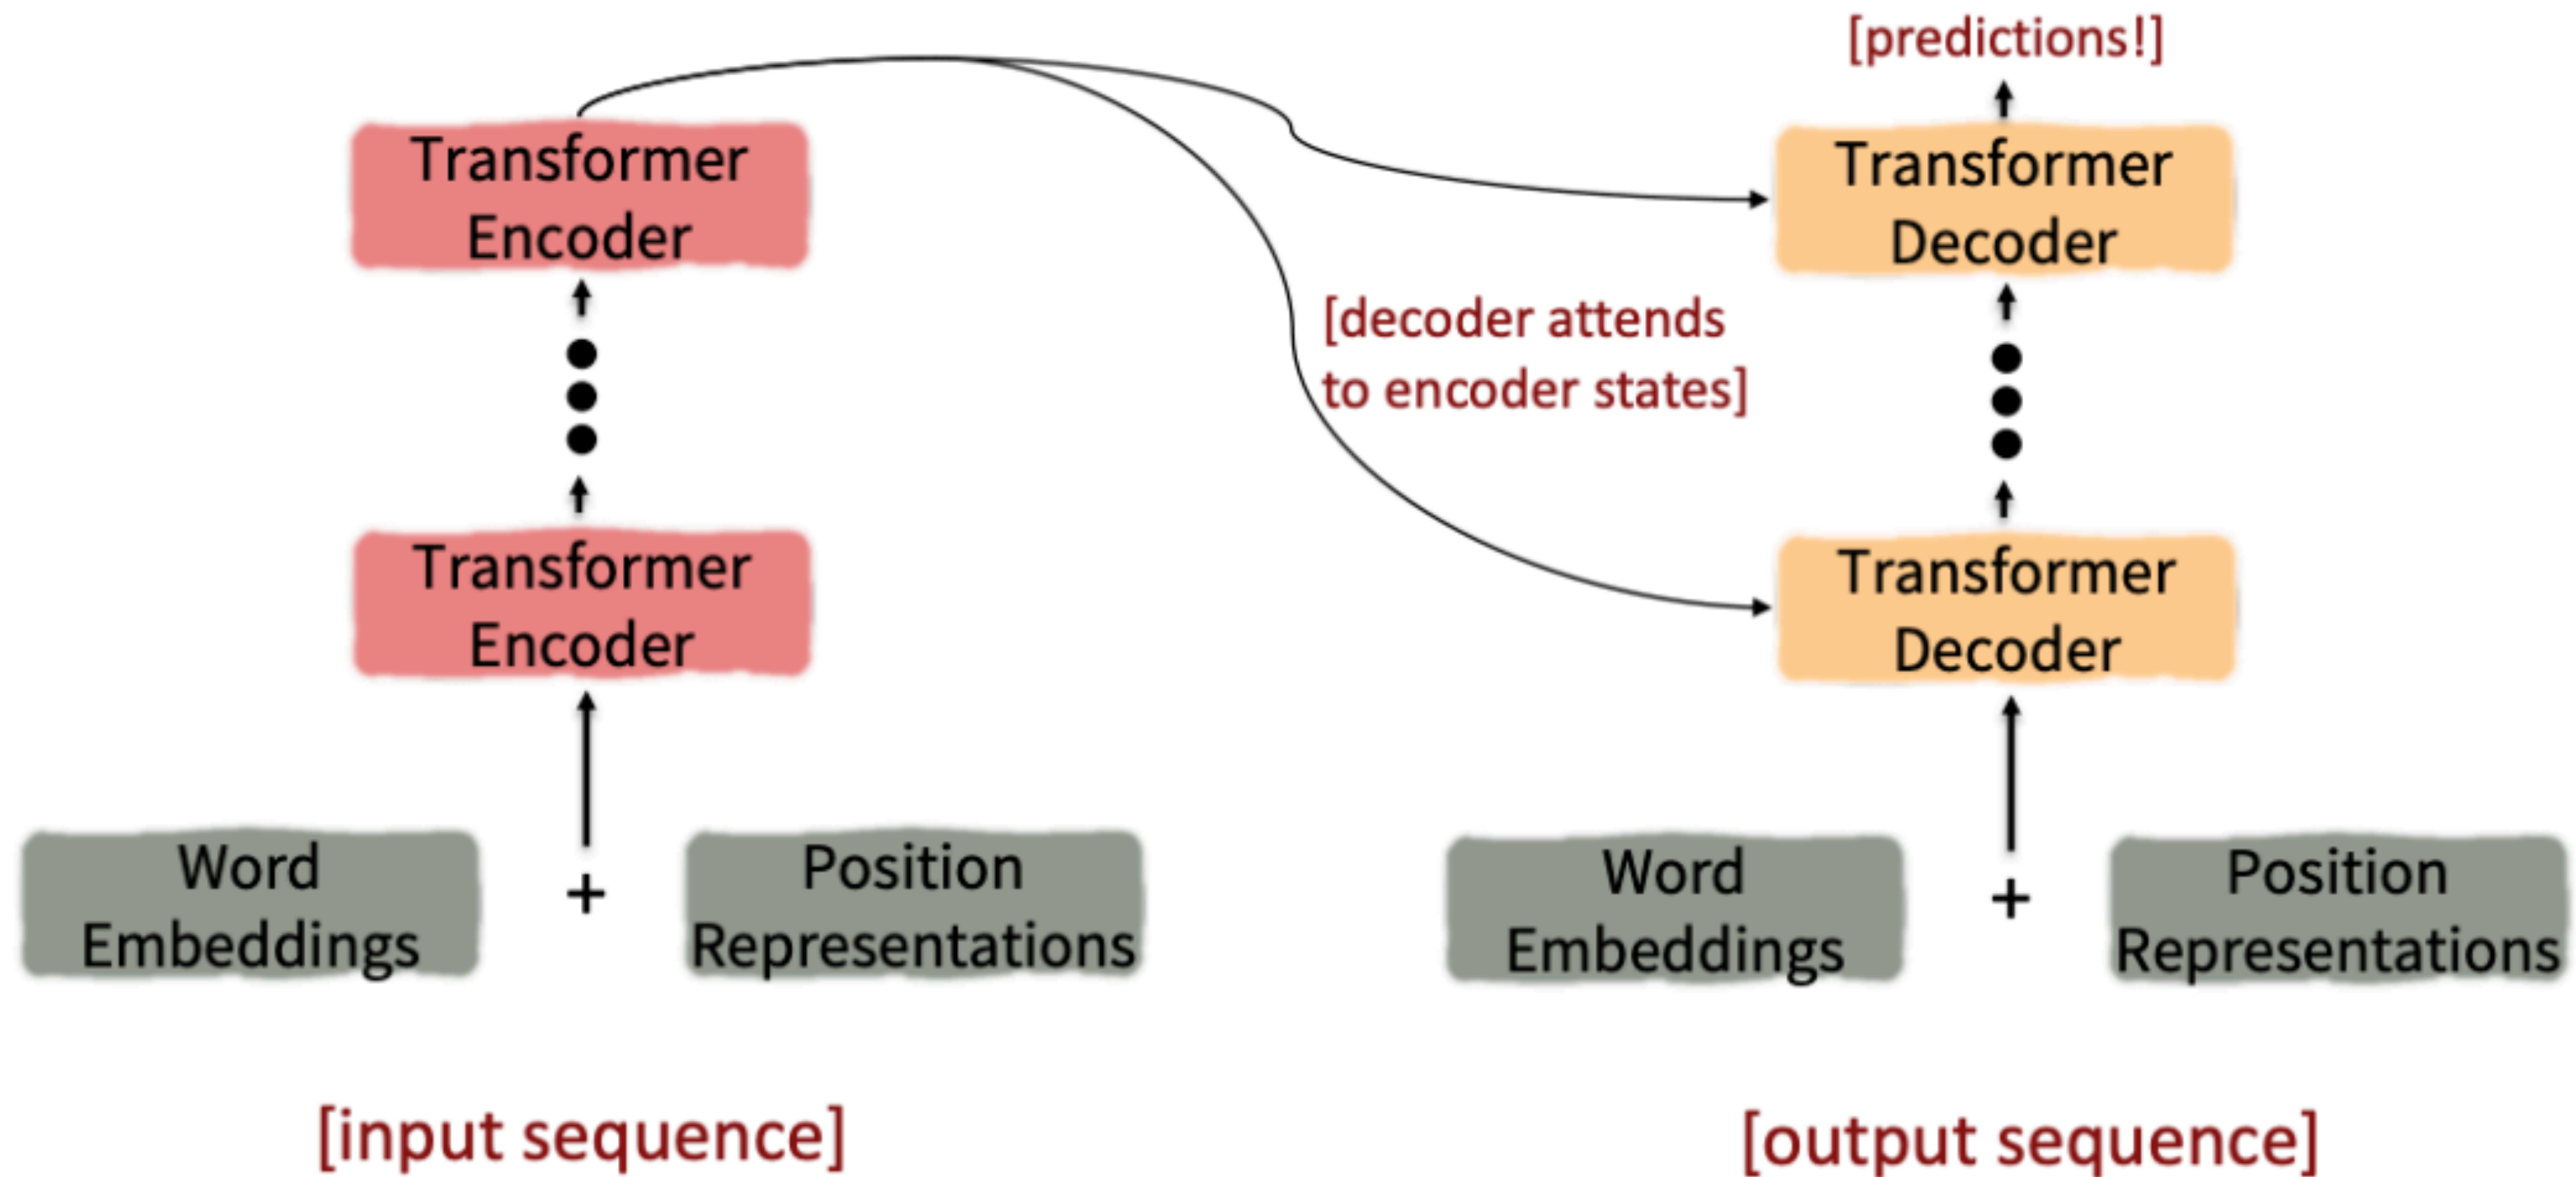
Efficient implementation: compute attention as we normally do, mask out attention to future words by setting attention scores to $-\infty$

```
dot = torch.bmm(queries, keys.transpose(1, 2))

indices = torch.triu_indices(t, t, offset=1)
dot[:, indices[0], indices[1]] = float('-inf')

dot = F.softmax(dot, dim=2)
```

Putting the pieces together



Transformer: Machine Translation

Model	BLEU		Training Cost (FLOPs)	
	EN-DE	EN-FR	EN-DE	EN-FR
ByteNet [18]	23.75			
Deep-Att + PosUnk [39]		39.2		$1.0 \cdot 10^{20}$
GNMT + RL [38]	24.6	39.92	$2.3 \cdot 10^{19}$	$1.4 \cdot 10^{20}$
ConvS2S [9]	25.16	40.46	$9.6 \cdot 10^{18}$	$1.5 \cdot 10^{20}$
MoE [32]	26.03	40.56	$2.0 \cdot 10^{19}$	$1.2 \cdot 10^{20}$
Deep-Att + PosUnk Ensemble [39]		40.4		$8.0 \cdot 10^{20}$
GNMT + RL Ensemble [38]	26.30	41.16	$1.8 \cdot 10^{20}$	$1.1 \cdot 10^{21}$
ConvS2S Ensemble [9]	26.36	41.29	$7.7 \cdot 10^{19}$	$1.2 \cdot 10^{21}$
Transformer (base model)	27.3	38.1	$3.3 \cdot 10^{18}$	
Transformer (big)	28.4	41.8	$2.3 \cdot 10^{19}$	

Transformer: Document Generation

Model	Test perplexity	ROUGE-L
<i>seq2seq-attention, $L = 500$</i>	5.04952	12.7
<i>Transformer-ED, $L = 500$</i>	2.46645	34.2
<i>Transformer-D, $L = 4000$</i>	2.22216	33.6
<i>Transformer-DMCA, no MoE-layer, $L = 11000$</i>	2.05159	36.2
<i>Transformer-DMCA, MoE-128, $L = 11000$</i>	1.92871	37.9
<i>Transformer-DMCA, MoE-256, $L = 7500$</i>	1.90325	38.8

Significant gains compared to
seq2seq-attention with LSTMs