Computer Science 572 Exam
Prof. Horowitz
Monday, November 26, 2018, 8:00am – 8:50am

Name: **Student Id Number:**

1. This is a closed book exam.
2. Please answer all questions.
3. There are a total of 30 questions. Question points may vary.
4. **Place your answer immediately below the question. Limit answers to ONE SENTENCE unless more is requested.**

1. [3 pts] Given the following two threads and initial values for x = 6 and y = 0, what are two different possible ending values for x and y? For each possible ending values provide the order of execution that supports your results.

Thread 1
```
void foo( ) {
  x++;
  y = x;
}
```

Thread 2
```
void bar( ) {
  y++;
  x++;
}
```

Thread 2,1 = 8,8 &
Thread 1,2,1,2 = 8,7

2. [3 pts] Does correlation necessarily imply causation. Yes or No?

NO

3. [3 pts] In the Google cloud, once your cluster has been set up what command is used to connect to the master machine?

SSH

4. [3 pts] Suppose one advertiser bids $1.50 for his ad to be displayed and a second advertiser bids $1.25 for his ad to be displayed and all other factors affecting ads are identical. If the first advertiser's ad is clicked on how much does the advertiser pay Google?

$1.26

5. [3 pts] Given two documents below, doc1 and doc2, provide the mapper output if an invertedIndex is run on the documents in a Hadoop cluster.

   doc1 – seven years ago our fathers did
   doc2 – three years ago our sisters did

   seven doc1
   years doc1
   ago doc1
   our doc1
   fathers doc1
   did doc1
   three doc2
   years doc2
   ago doc2
   our doc2
   sisters doc2
   did doc2

6. [3 pts] Given the two documents above, doc1 and doc2, provide the reducer output if an invertedIndex is run on the documents in a Hadoop cluster.

   seven doc1:1
   years doc1:1 doc2:1
   ago doc1:1 doc2:1
   our doc1:1 doc2:1
   fathers doc1:1
   did doc1:1 doc2:1
   three doc2:1
   sisters doc2:1

7. [3 pts] The class notes describe long-tailed keywords as search queries made up of three, four or more word phrases. Is the conversion rate for long-tailed keywords better or worse that it is for shorter keywords?

   better, 2.5 times than shorter

8. [3 pts] What is the name of Google's knowledgebase?

   Google's Knowledge Grpah

9. [3 pts] When viewed as a graph, an ontology is what sort of graph? Use conventional graph terms. We are expecting at least three graph properties.

labeled, directed & cyclic

10. [4 pts] The discussion of search engine optimization (SEO) identified many factors which correlate strongly with attaining a high ranking on the search engine result page. Mention four of them.

1. Content factors • Content relevance • Word count
2. User signals • Click Through Rate (CTR) • Bounce rate
3. Technical factors • Presence of H1/H2 • Use of HTTPS
4. User experience • Number of internal/external links
5. Social signals • Facebook total • Tweets
6. Backlinks- No. of backlinks, No.of DoFollow backlinks

11. [3 pts] Define Mean Reciprocal Rank scoring

Repeated from Spring '19.

Below is the Norvig spelling corrector program written in Python and presented in class. Please answer the questions that follow the program.

```
import re, collections
def words(text): return re.findall('[a-z]+', text.lower())
def train(features):
    model = collections.defaultdict(lambda: 1)
    for f in features:
        model[f] += 1
    return model
NWORDS = train(words(file('big.txt').read()))
alphabet = 'abcdefghijklmnopqrstuvwxyz'
def edits1(word):
   splits     = [(word[:i], word[i:]) for i in range(len(word) + 1)]
   deletes    = [a + b[1:] for a, b in splits if b]
   transposes = [a + b[1] + b[0] + b[2:] for a, b in splits if
len(b)>1]
   replaces   = [a + c + b[1:] for a, b in splits for c in alphabet if
b]
   inserts    = [a + c + b     for a, b in splits for c in alphabet]
   return set(deletes + transposes + replaces + inserts)
def known_edits2(word):
    return set(e2 for e1 in edits1(word) for e2 in edits1(e1) if e2 in
NWORDS)
def known(words): return set(w for w in words if w in NWORDS)
def correct(word):
    candidates = known([word]) or known(edits1(word)) or
known_edits2(word) or [word]
    return max(candidates, key=NWORDS.get)
```

12. [3 pts] What functions are defined?

words(text)
train(features)
edits1(word)
known_edits2
known
correct

13. [3 pts] What function is used to invoke (start) the program

words

14. [3 pts] What edit operations are included in the program?

inserts
deletes
transposes
replaces

15. [3 pts] How many levels of edits does the program investigate?

2

16. [3 pts] For a clustering to be considered good (or successful) what is the similarity property to be satisfied for its intra-class elements and inter-class elements?

intra class-high
inter class-low

17. [3 pts] Define soft clustering

Repeated from Spring '19.

18. [4 pts] In the k-means clustering algorithm there are several possible criteria for termination. mention two.

Repeated from Spring '19.

19. [4 pts] Given n documents each expressed as an m-element vector, what is the computing time for the Agglomerative Clustering algorithm assuming priority queues are used?

O(mnlogn)

20. [4 pts] Given the two strings: “information” and “interrogation”, what is their minimum edit distance assuming the operations (replace, delete, insert) all have a count of *1*?

5

21. [3 pts] In class we discussed that there are three distinct phases for a question/answering system. Name them.

Question Processing
Passage Retrieval
Answer Processing

22. [4 pts] The terms hyperonymy and hyponymy are used with respect to WordNet. Define one of the two terms and give an example

Repeated from Spring '19.

Hyponym: plant -> tree (specialization)
Hypernym: apple -> fruit (generalization)

23. [3 pts] Given 75 documents divided into three clusters where cluster 1 has 10 related documents, cluster 2 has 5 related documents and cluster 3 has 10 related documents, what is the Purity Index of this clustering?

1/3, NOT IN PORTION

24. [3 pts] For the k-means algorithm, is the centroid necessarily a document in the set of documents?

NO

25. [3 pts] Is the graph provided to NetworkX directed or undirected?

directed

26. [4 pts] This semester we examined two algorithms for clustering and two algorithms for classification. Name all four.

Clustering-
K-means
Heirarchical

Classification-
KNN
Rocchio

27. [3 pts] There are three types of spelling errors. Non-word errors, Typographical errors and Cognitive errors. Define cognitive errors and give an example.

Cognitive errors (homophones, sounds alike)
– piece - peace,
– too - two,
– your - you're

28. [4 pts] There are six different strategies to speed up indexed retrieval mentioned in class. Mention any three of them.

Repeated from Spring '19.

29. [4 pts] When viewed as a graph, a knowledge graph is what sort of graph? Use conventional graph terms. We are expecting at least two graph properties.

Repeated from Spring '19.

30. [5 pts] Below is the equation for maximum likelihood estimation for bigram probabilities? Explain the terms in the equation.

$$P(w_i \mid w_{i-1}) = \frac{count(w_{i-1},\ w_i)}{count(w_{i-1})}$$

w(i-1): (i-1)th word in the document
w(i): ith word in the document

P(w(i)|w(i-1)): Probability of (i)th word occuring given (i-1)th word has already occured
count(w(i-1),w(i)): number oftimes that (i-1)th word and (i)th word occur together in given order
count(w(i-1)): frequency of (i-1)th word in the document